



Extensible Autonomous Transactions in the world of Microservices using Apache Kafka


Nice to meet you



NICOLA GIACCHETTA



[nicolagiacchetta](#)




DIVYA NAGAR



[divya2661](#)

We work in a team of 3 people along with **NARAIN RAMASWAMY** at **nexmo**[®]
The **Vonage** API Platform

What is Nexmo, The Vonage API Platform?



Full Product Suite



Voice



Messages



Video



Numbers Insight



Verify



Dispatch



SMS



Chat



SIP Trunking

Enterprise DNA



Loved by developers.
Built for business.

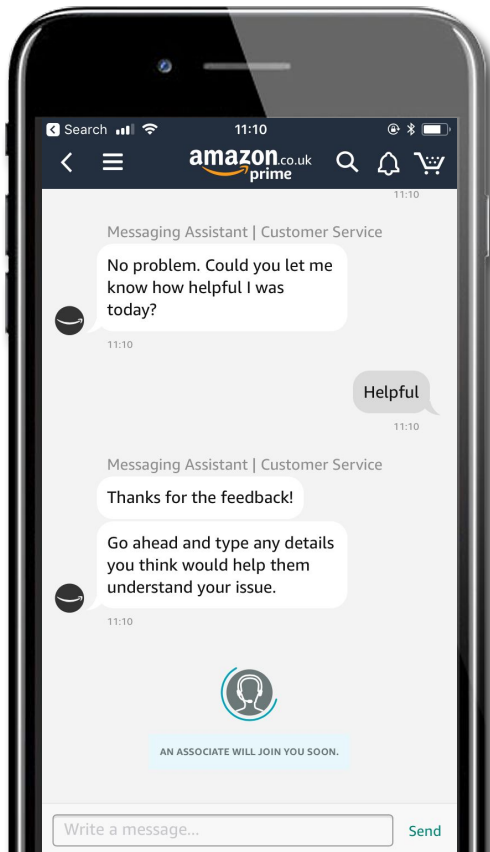
Global Coverage



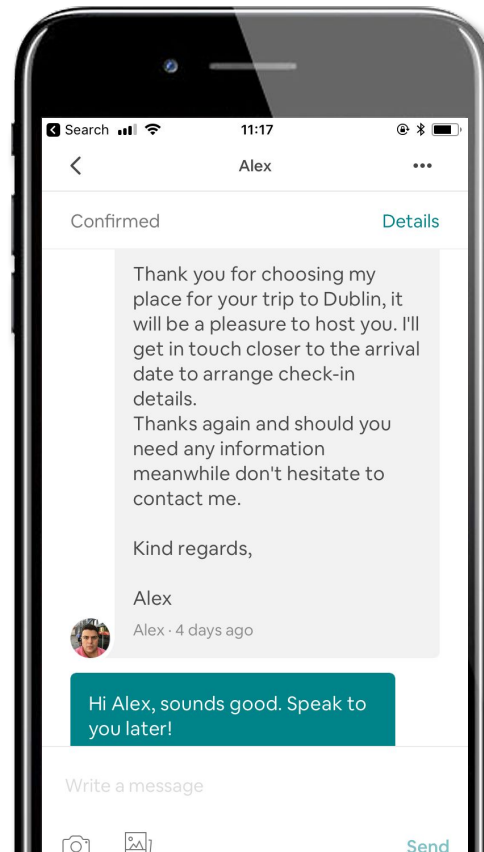
225 countries and over 1,600
telecommunications networks.

Digital native companies have embraced programmable communications

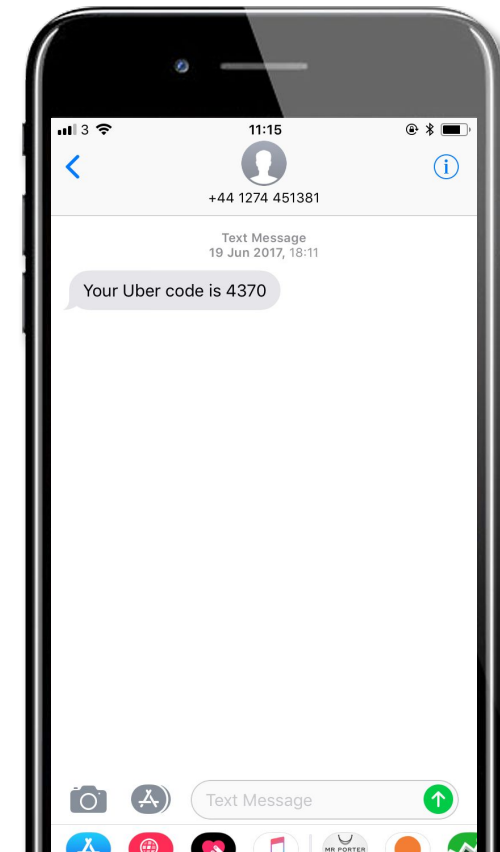
amazon



airbnb



UBER





- Small Backend Application → Monolith
- Scaling Read Workload with Volga
- Monolith → Microservices
- Transactions in the World of Microservices


Once upon a time...



Americas

EMEA

APAC



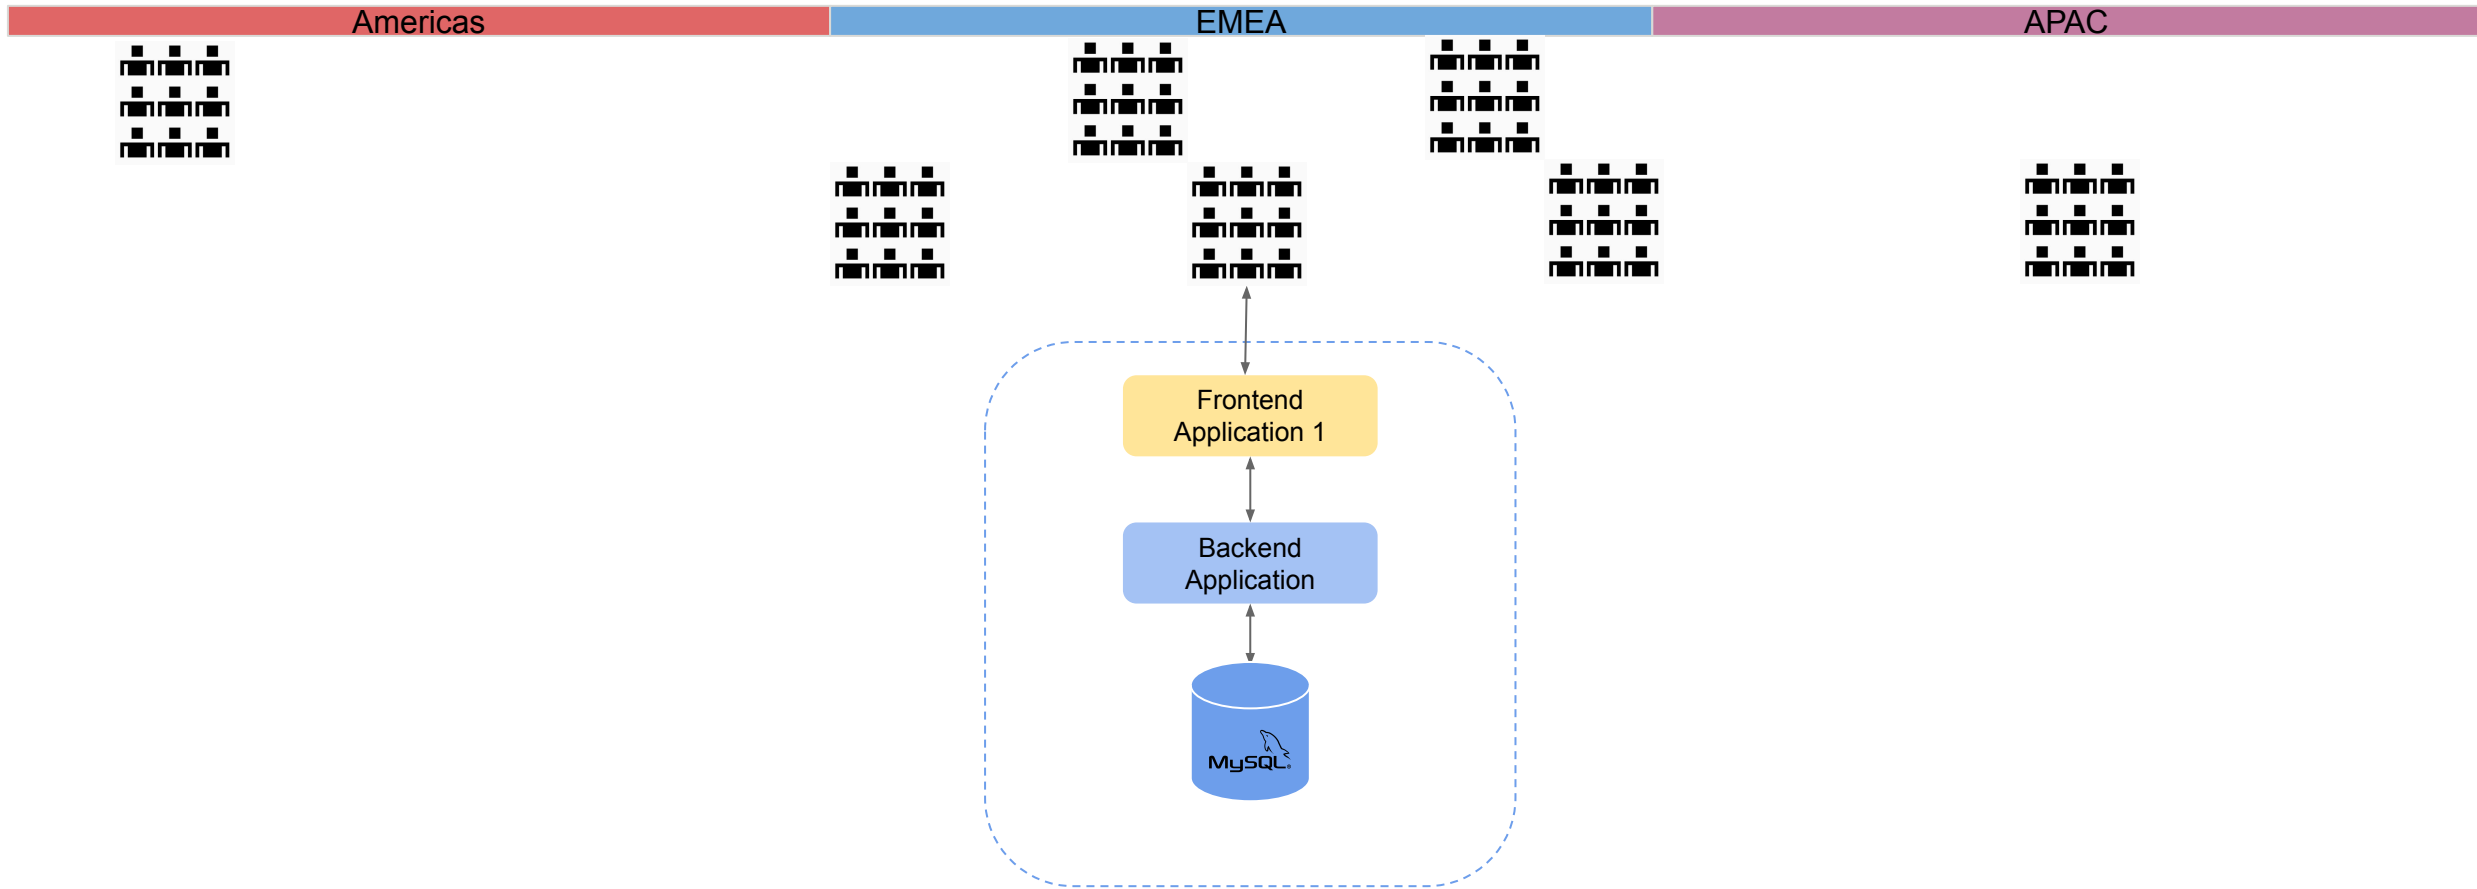
1 application in **1 DC** backed by **1 DB** handling **all reads** and **all writes**

Something more about our Backend Application



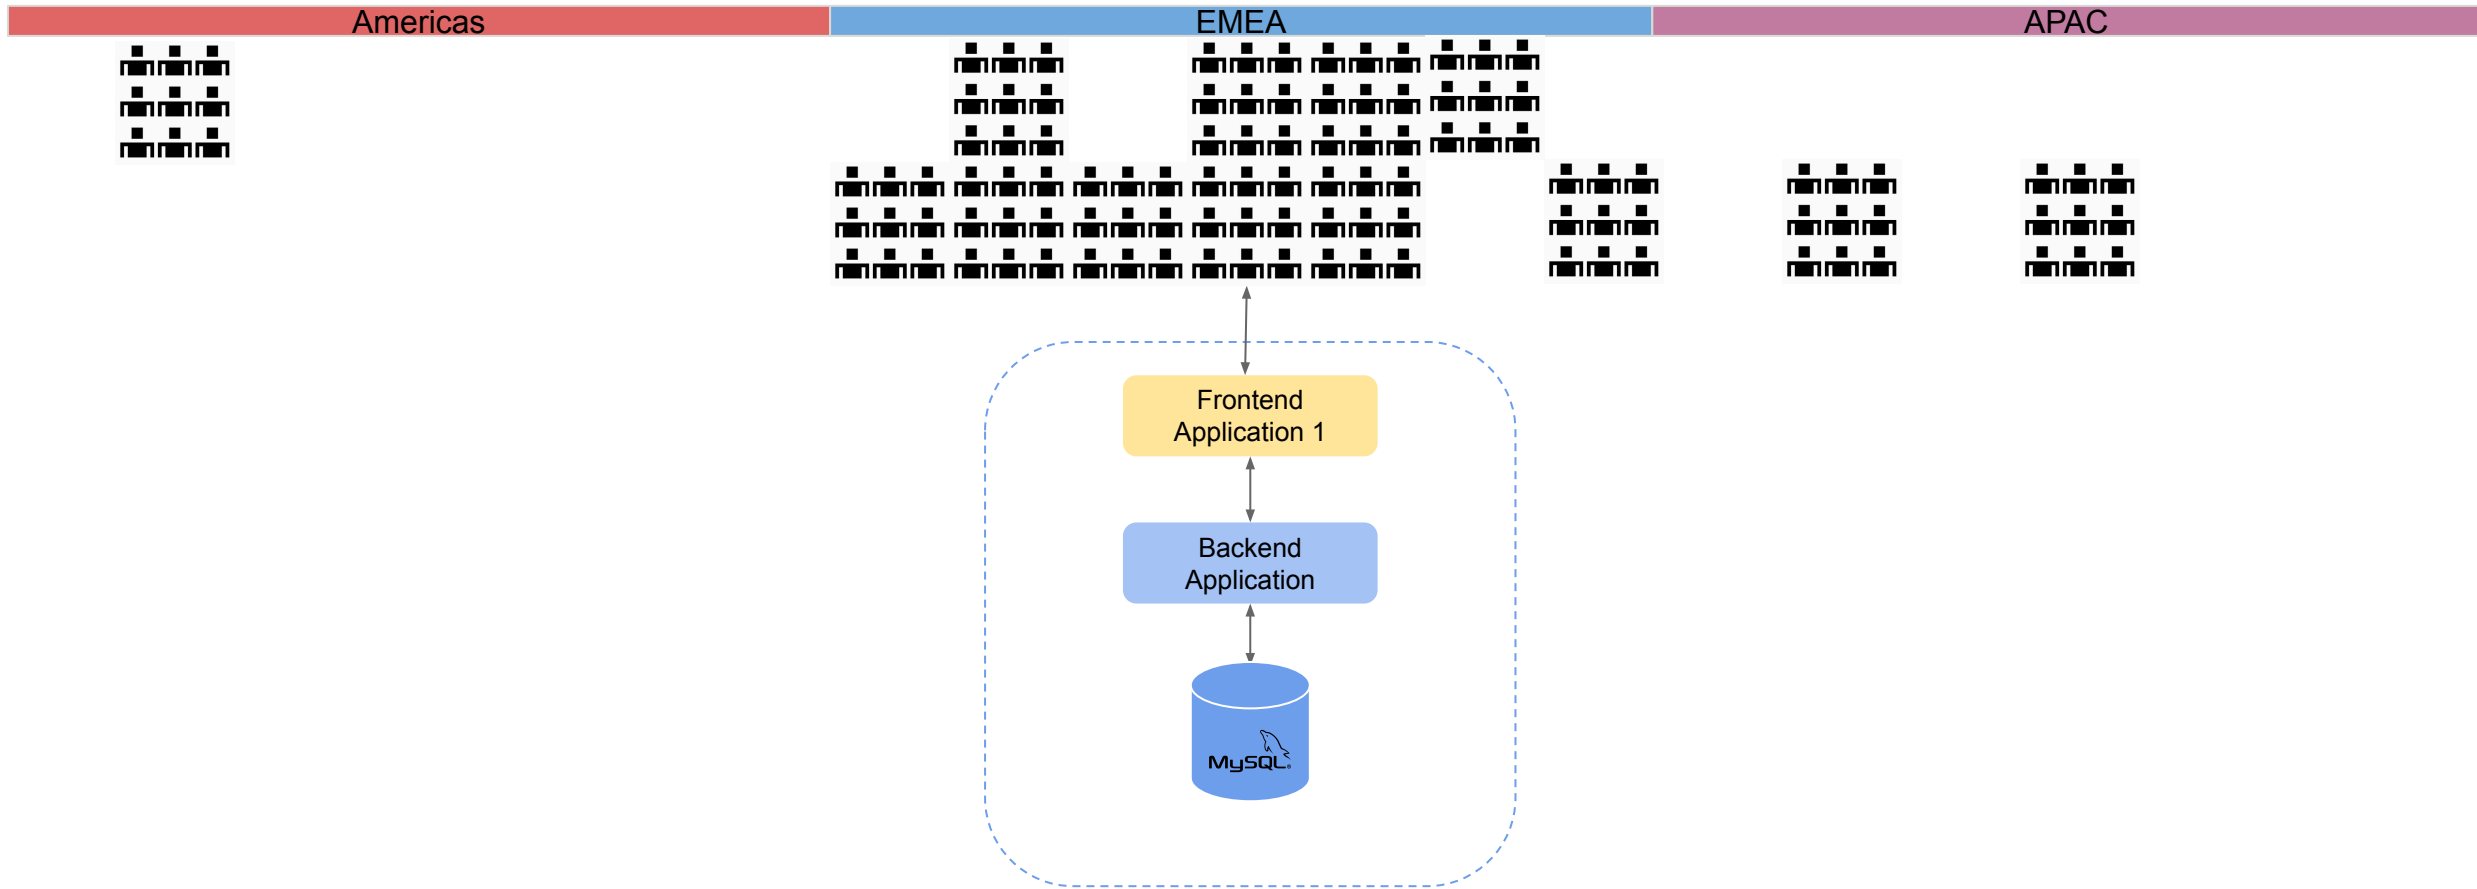
- Backed by a **MySQL** Database
- Reads >>> Writes

We were starting to grow



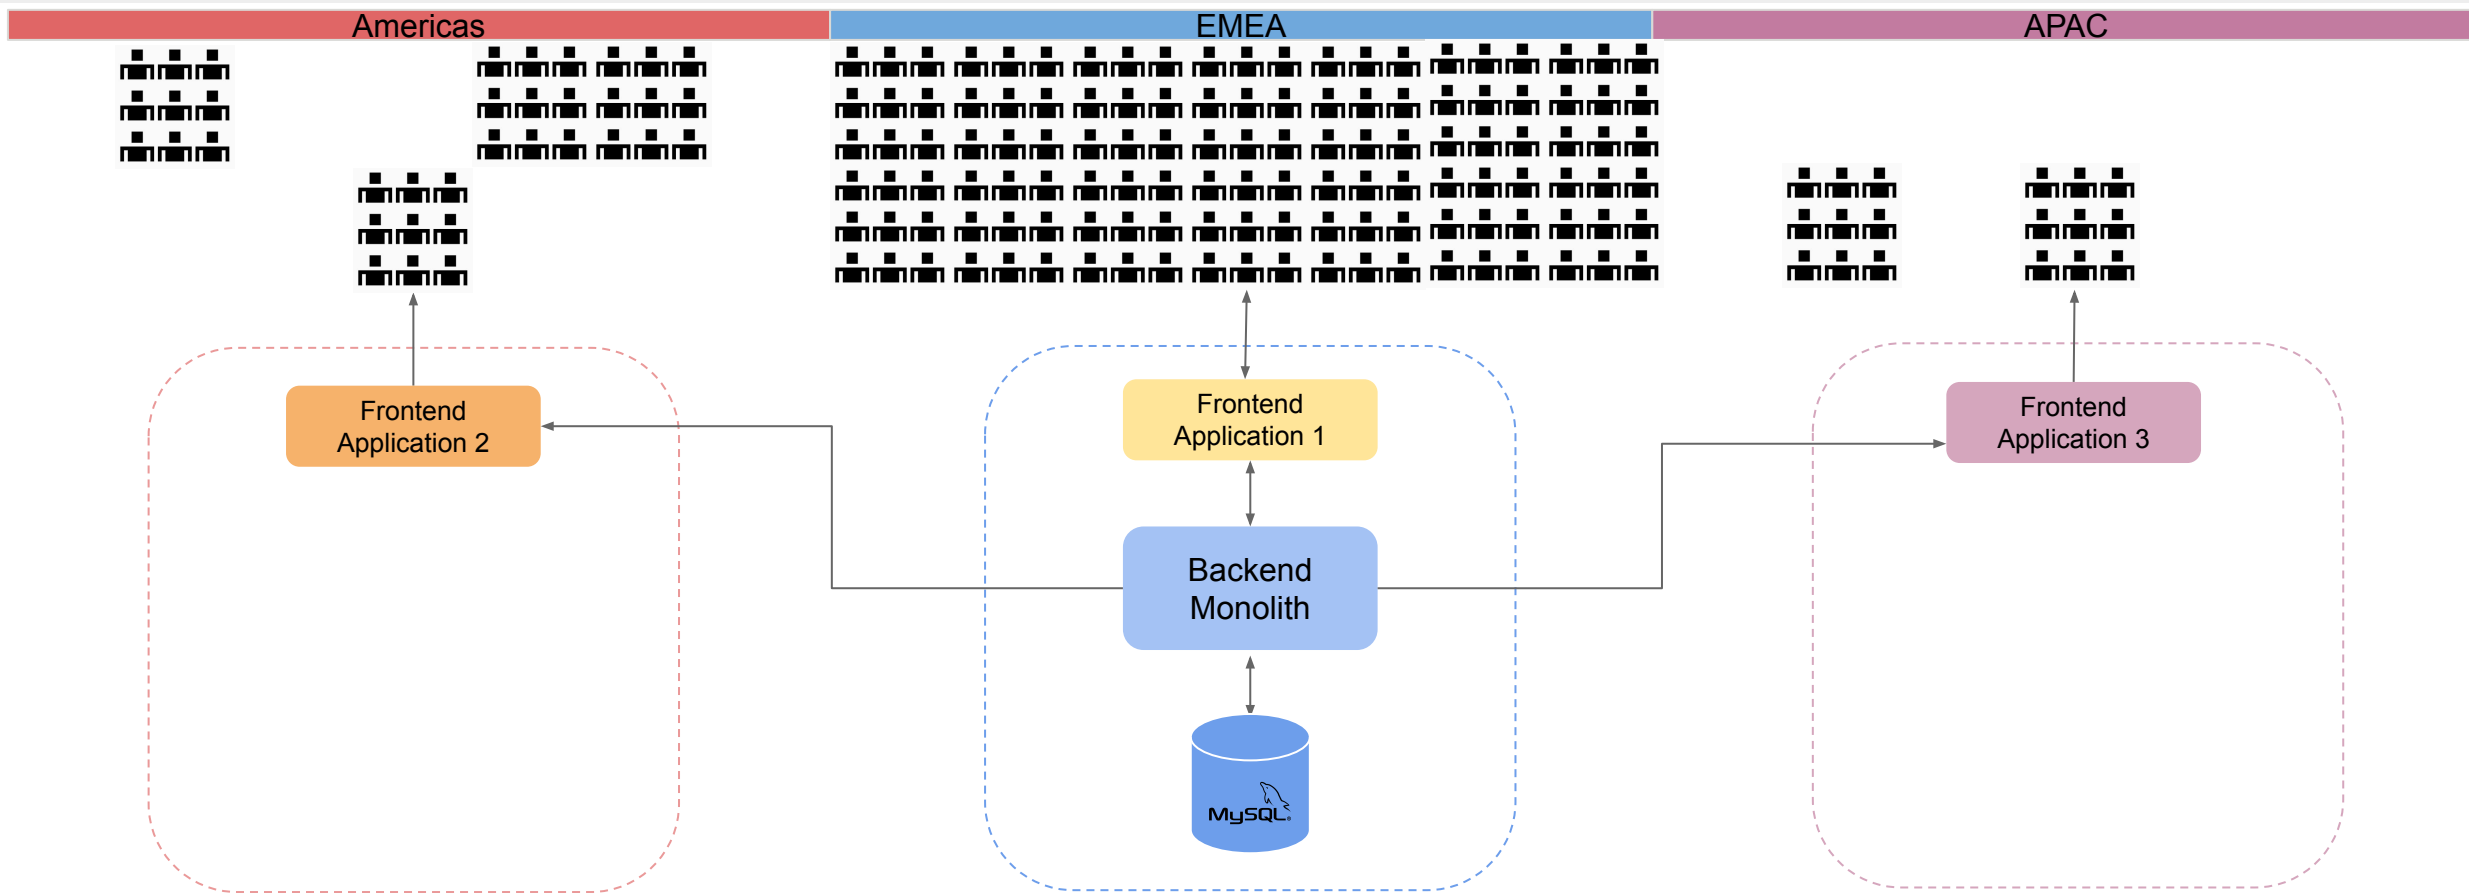
1 application in **1 DC** backed by **1 DB** handling **all reads** and **all writes**

We were starting to grow and grow



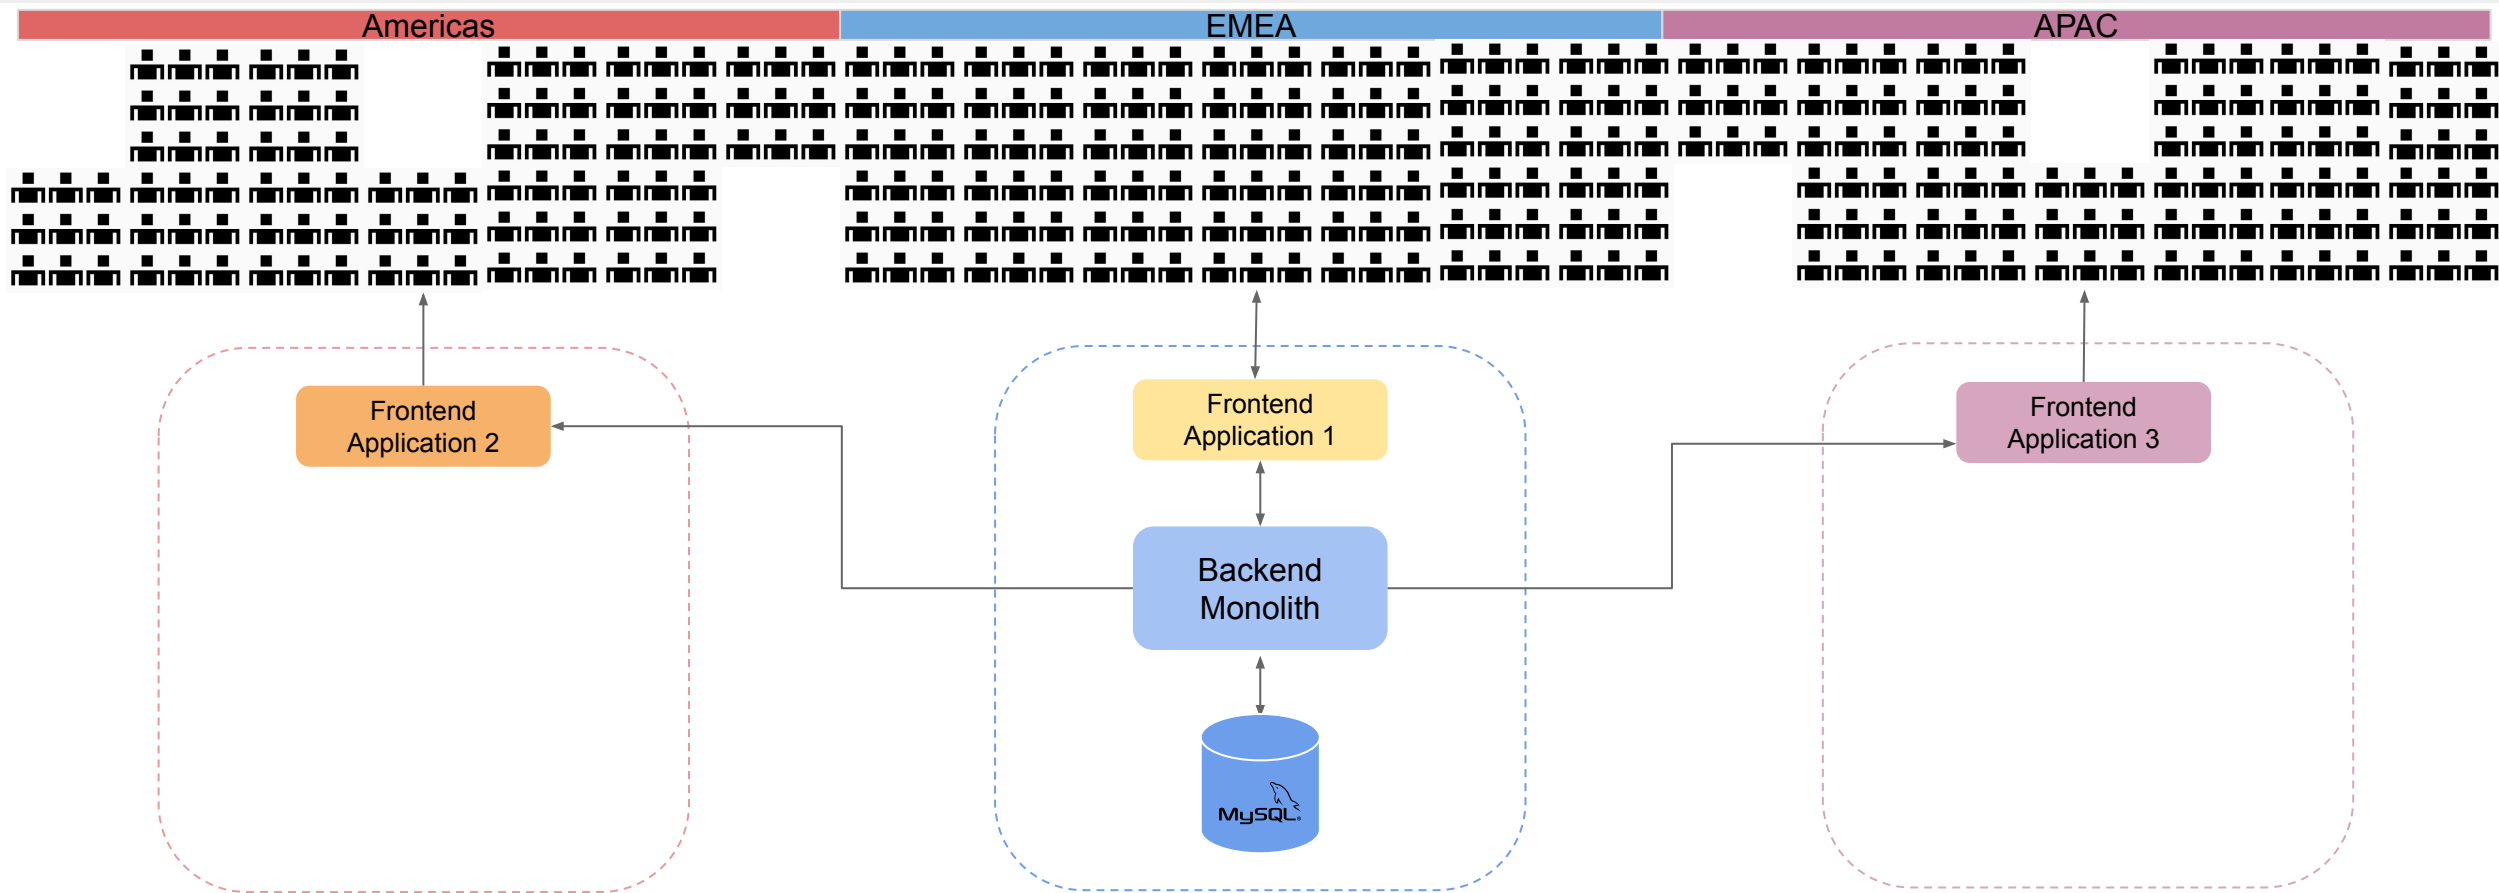
1 application in **1 DC** backed by **1 DB** handling **all reads** and **all writes**

We were starting to grow and grow and grow



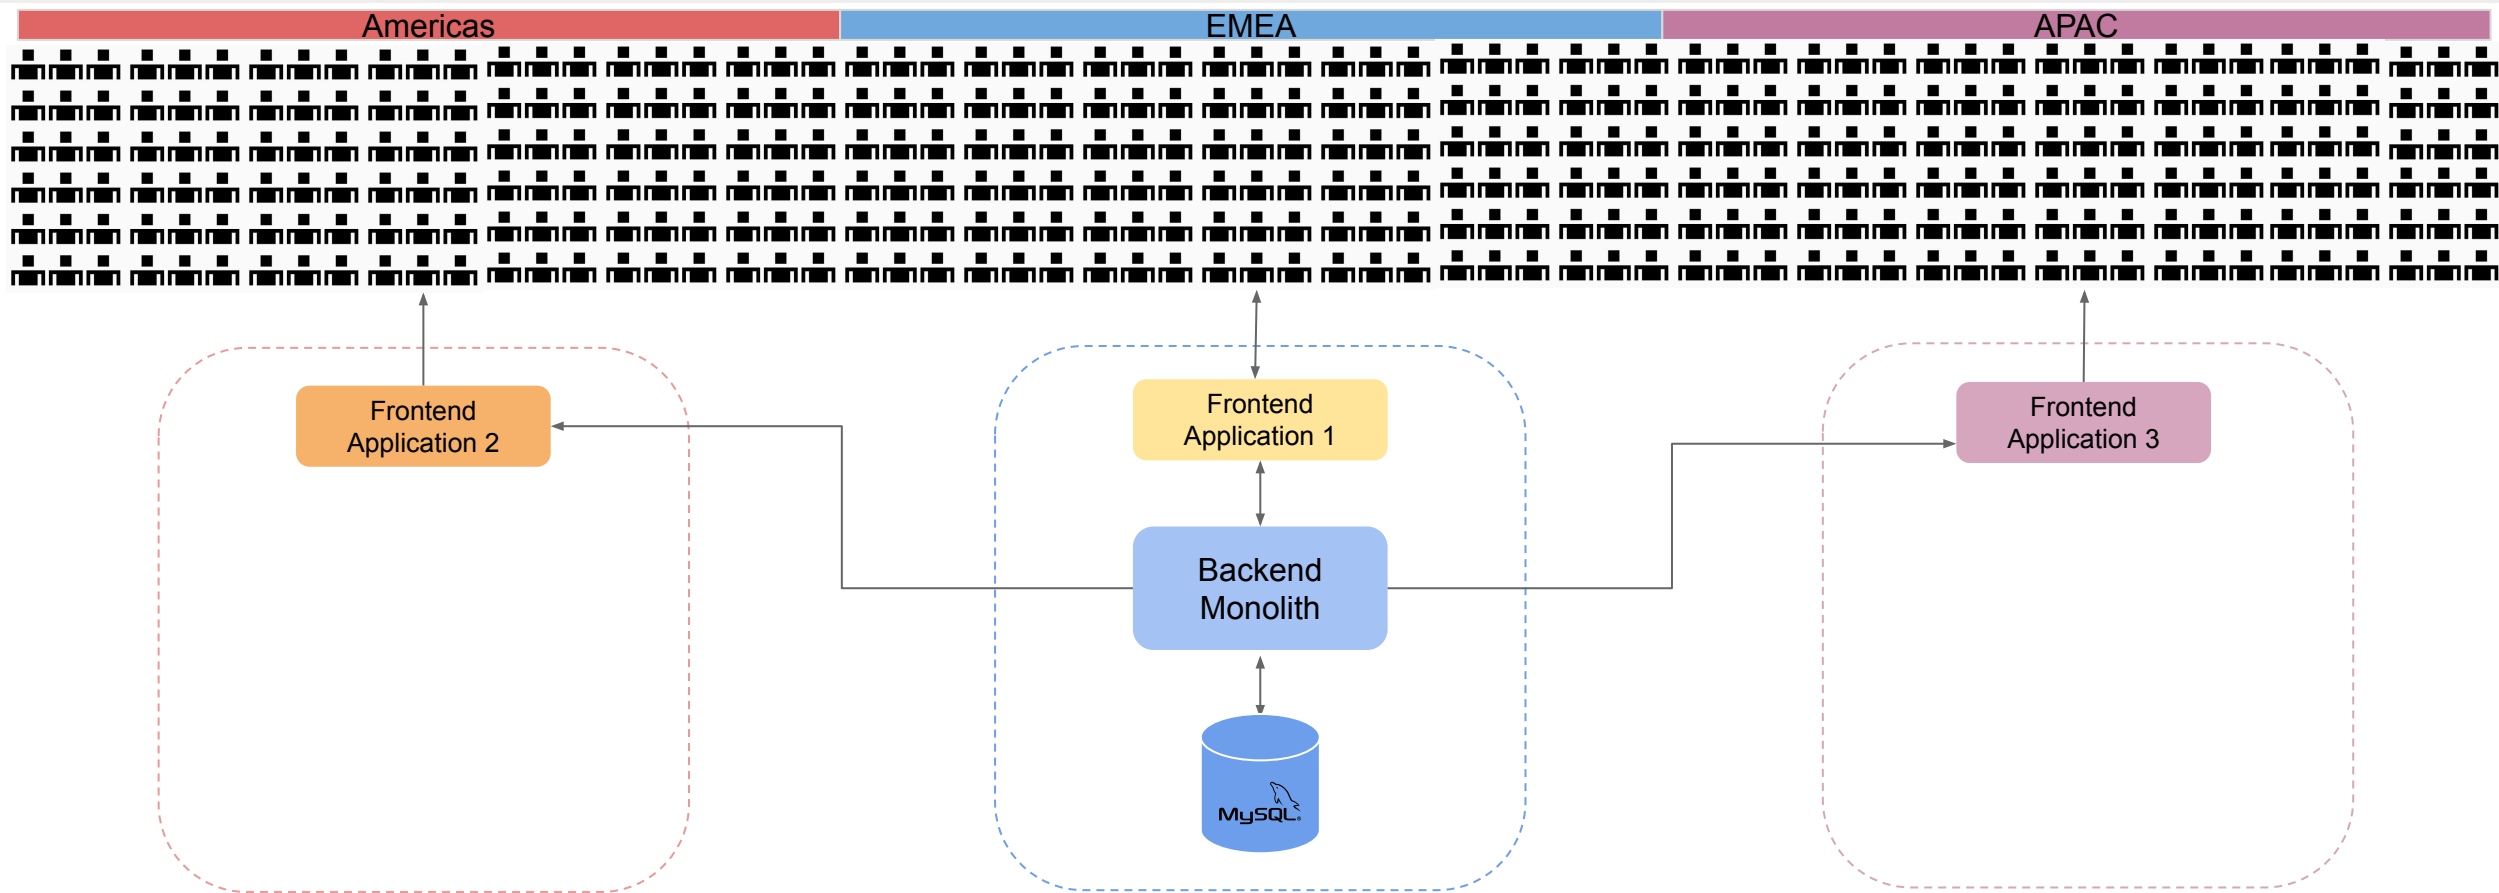
1 application in **1 DC** backed by **1 DB** handling **all reads** and **all writes**

We were starting to grow and grow and grow and grow



1 application in **1 DC** backed by **1 DB** handling **all reads** and **all writes**

We were starting to grow and grow and grow and grow and grow



1 application in **1 DC** backed by **1 DB** handling **all reads** and **all writes**



1. Read workload

2. Monolithic Architecture



1. Read workload

2. Monolithic Architecture



Problems:

- Too much Load
- Too high Latency




Problems:

- Too much Load
- Too high Latency


Solutions:

- Adding **Caches** closer to the Applications
- Add **Read Replicas** in other regions

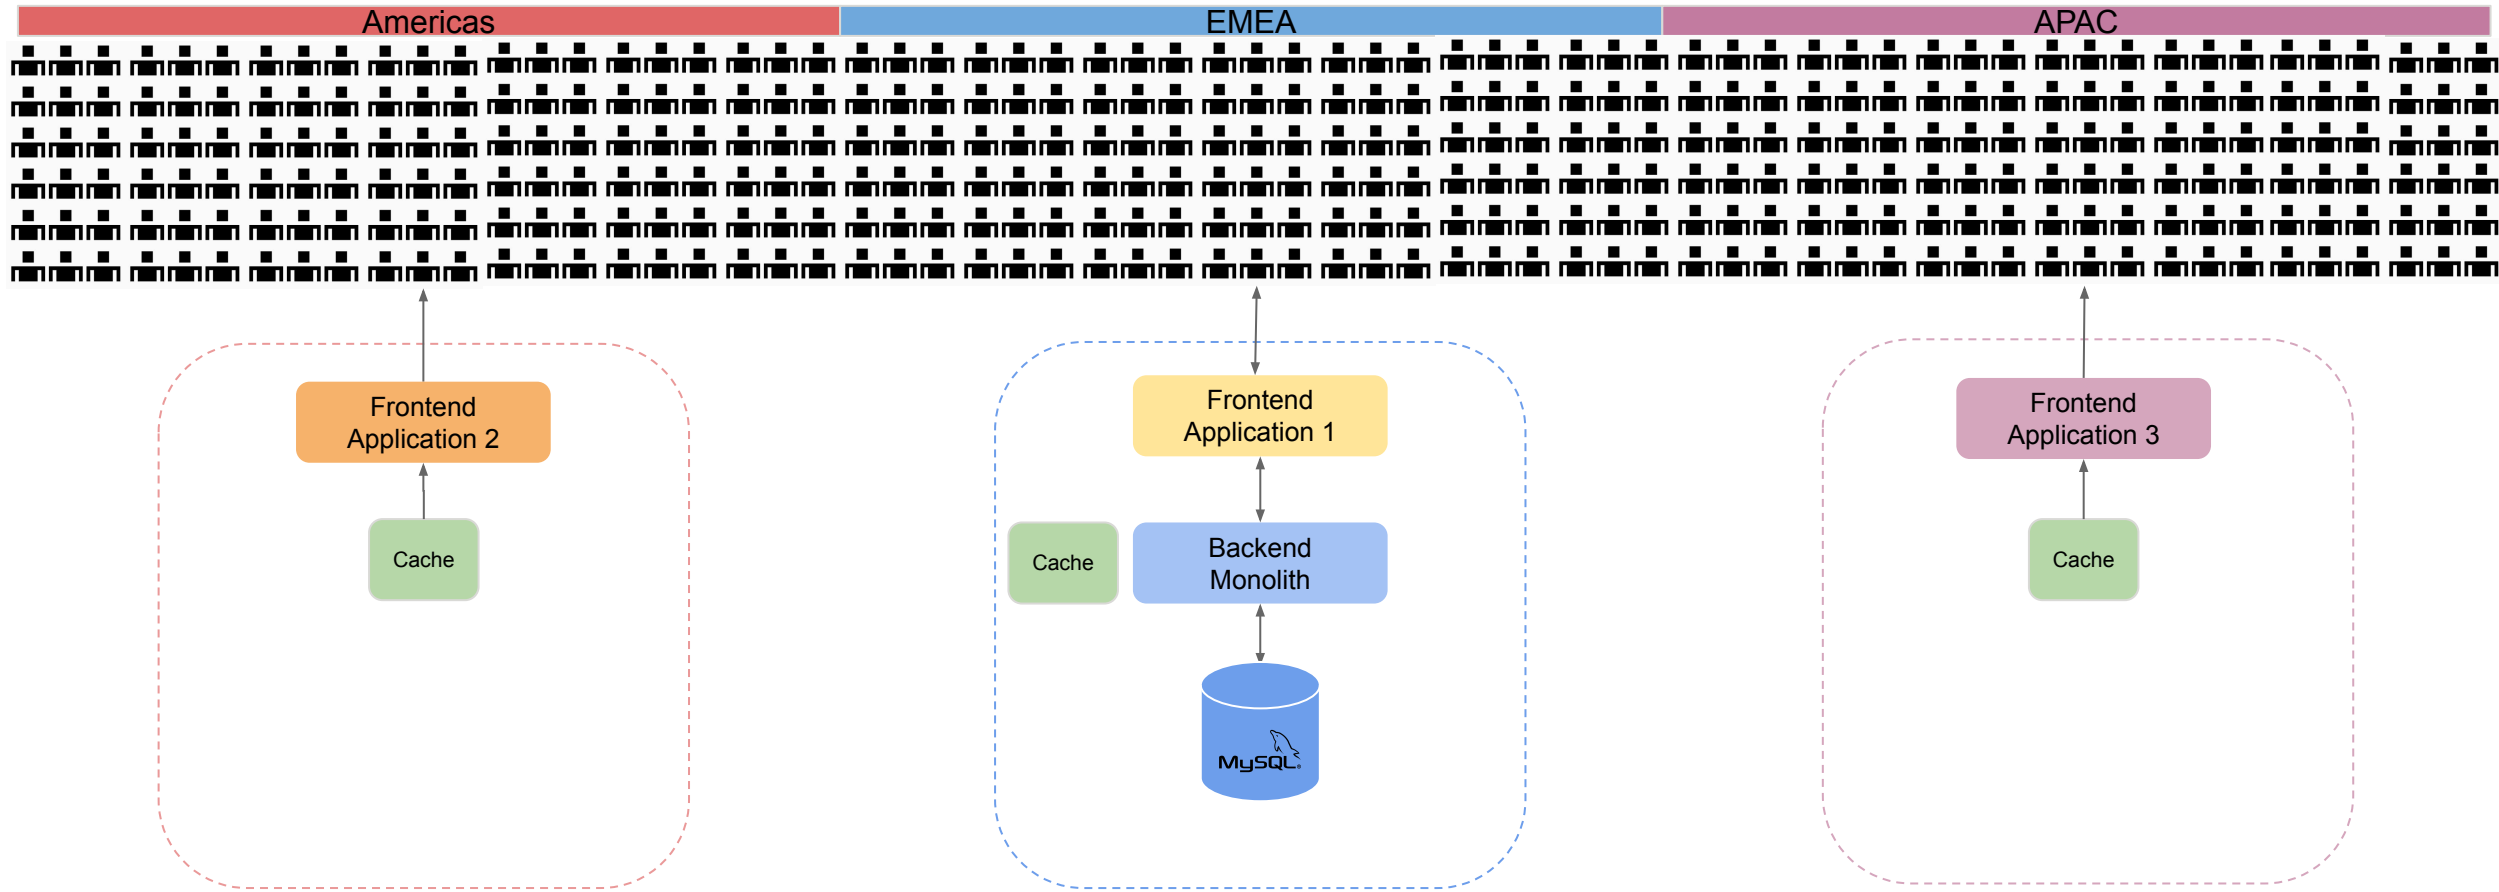
High Consistency, High-Availability, Partition Tolerance: choose **2**.



High Consistency, High-Availability, ~~Partition Tolerance~~: choose **1**.

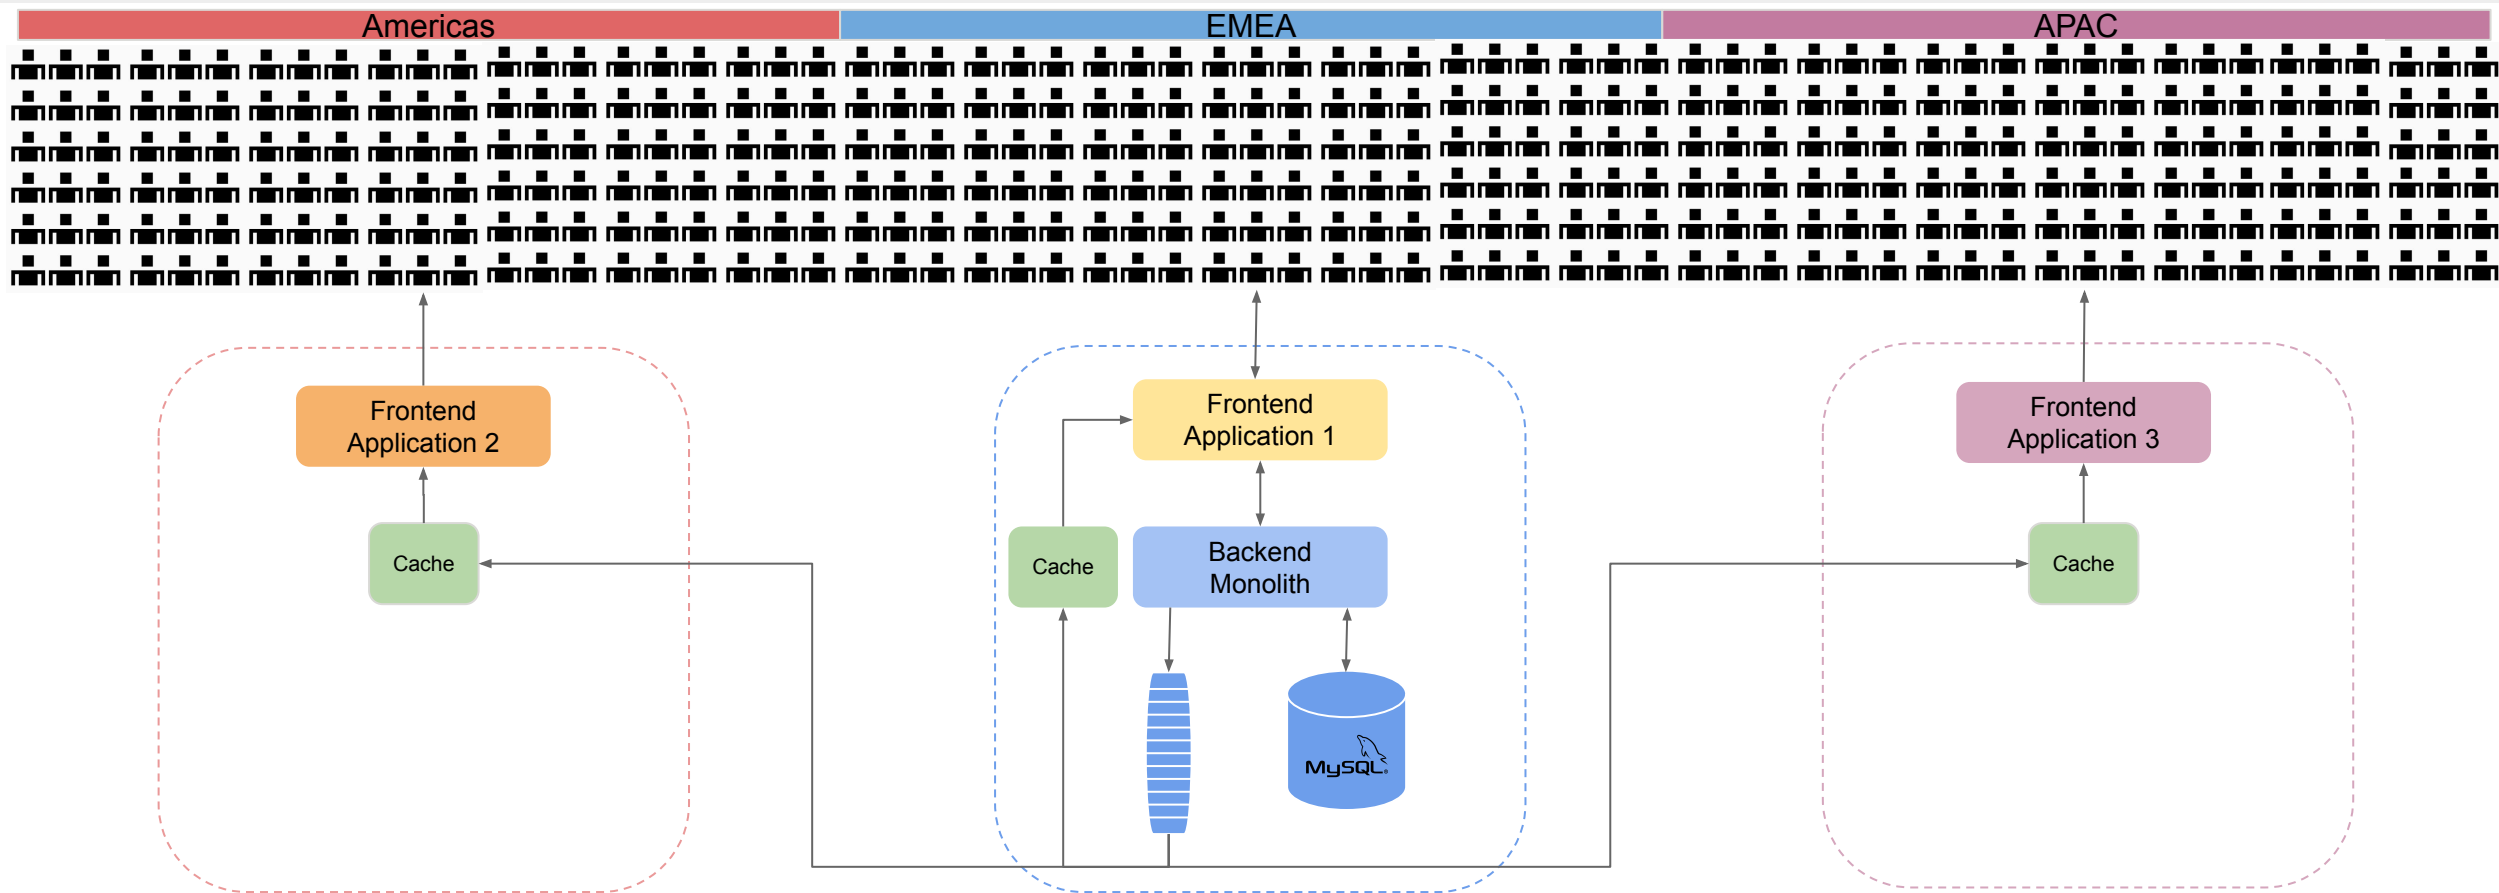


How do we notify the caches?



«There are only two hard things in Computer Science: cache invalidation and naming things.» -- **Phil Karlton**

Dual Writes




The application writes to the database and to another messaging system **in parallel**

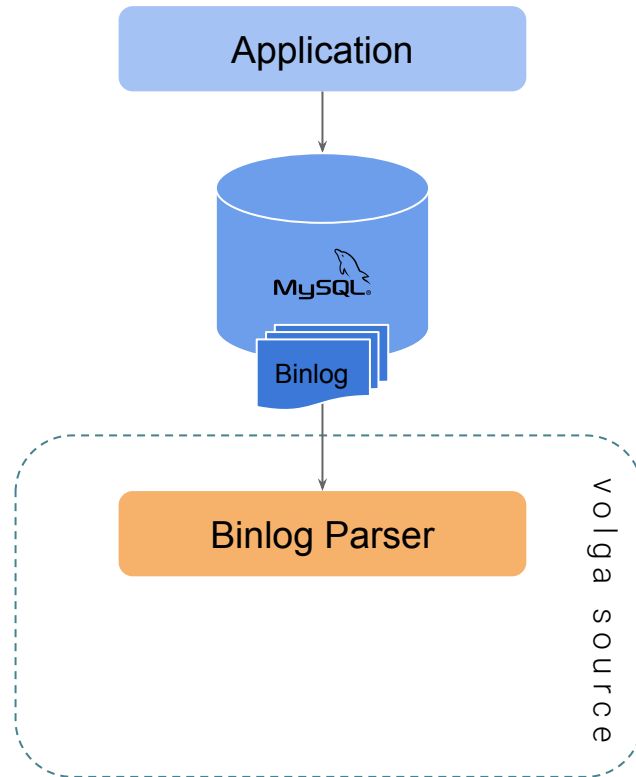


The application writes to the database and **in parallel** to another messaging system creating two main issues:

1. Without a complex **coordination protocol** (e.g. 2-Phase Commit or Consensus) it is hard to achieve sequential consistency between the database and the messaging system
2. The logic to write to the messaging system is **strictly-coupled** with the application



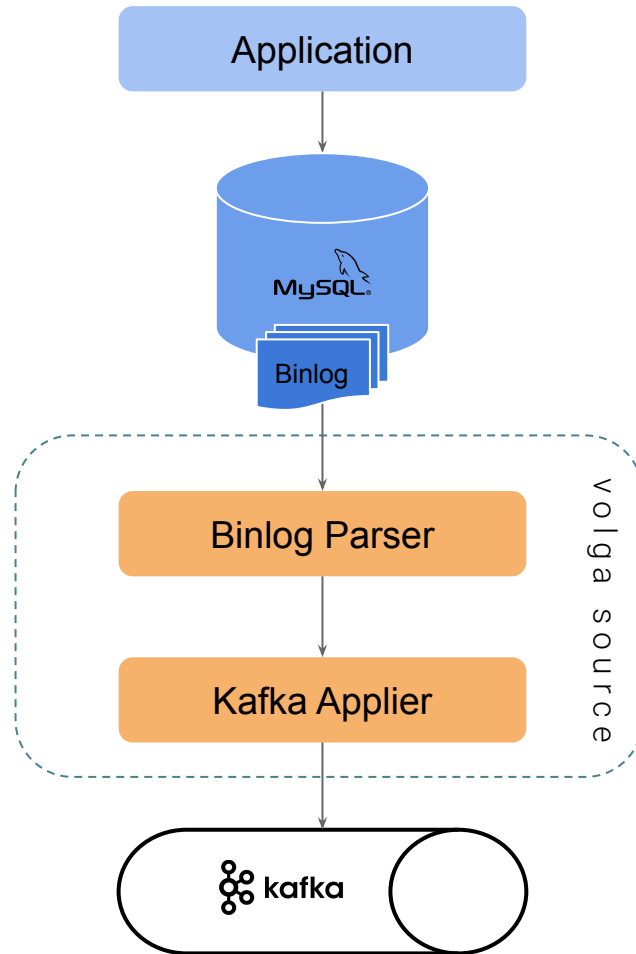
The database is the **single source-of-truth** and changes are extracted from its transaction or commit log.



The database is the **single source-of-truth** and changes are extracted from its transaction or commit log.

Volga:

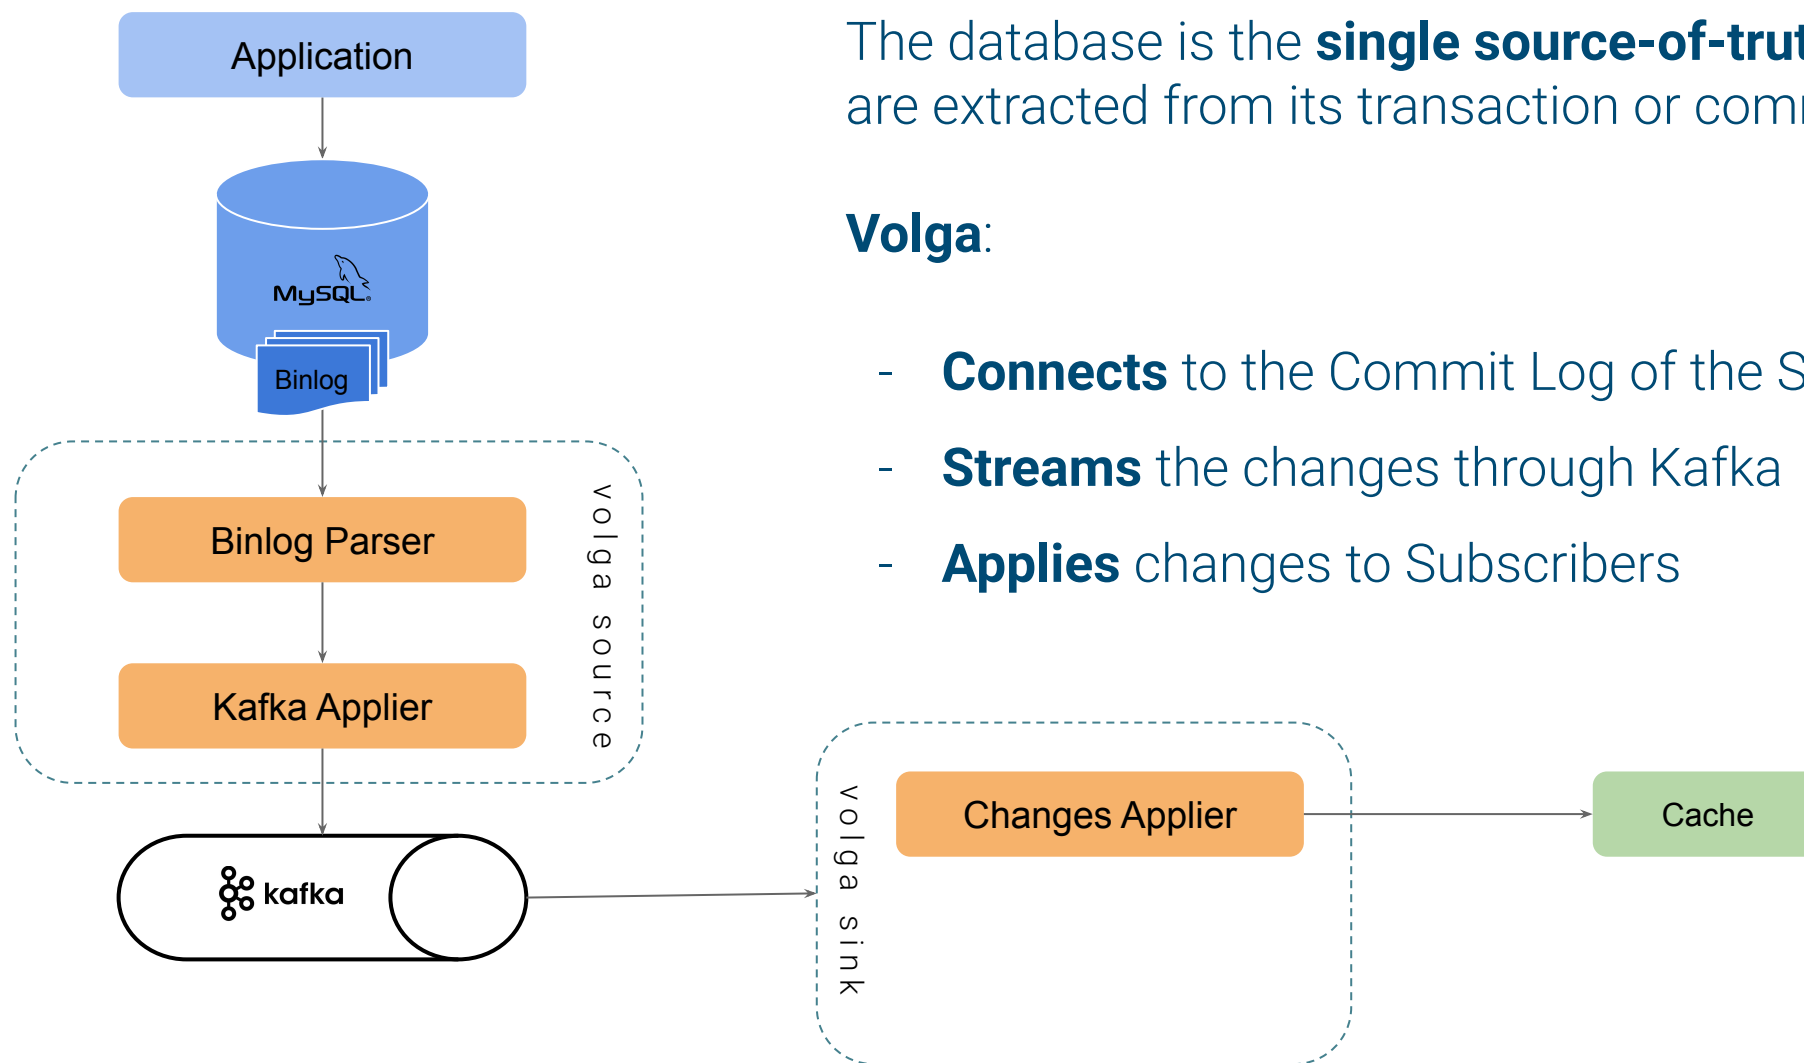
- **Connects** to the Commit Log of the Source DB



The database is the **single source-of-truth** and changes are extracted from its transaction or commit log.

Volga:

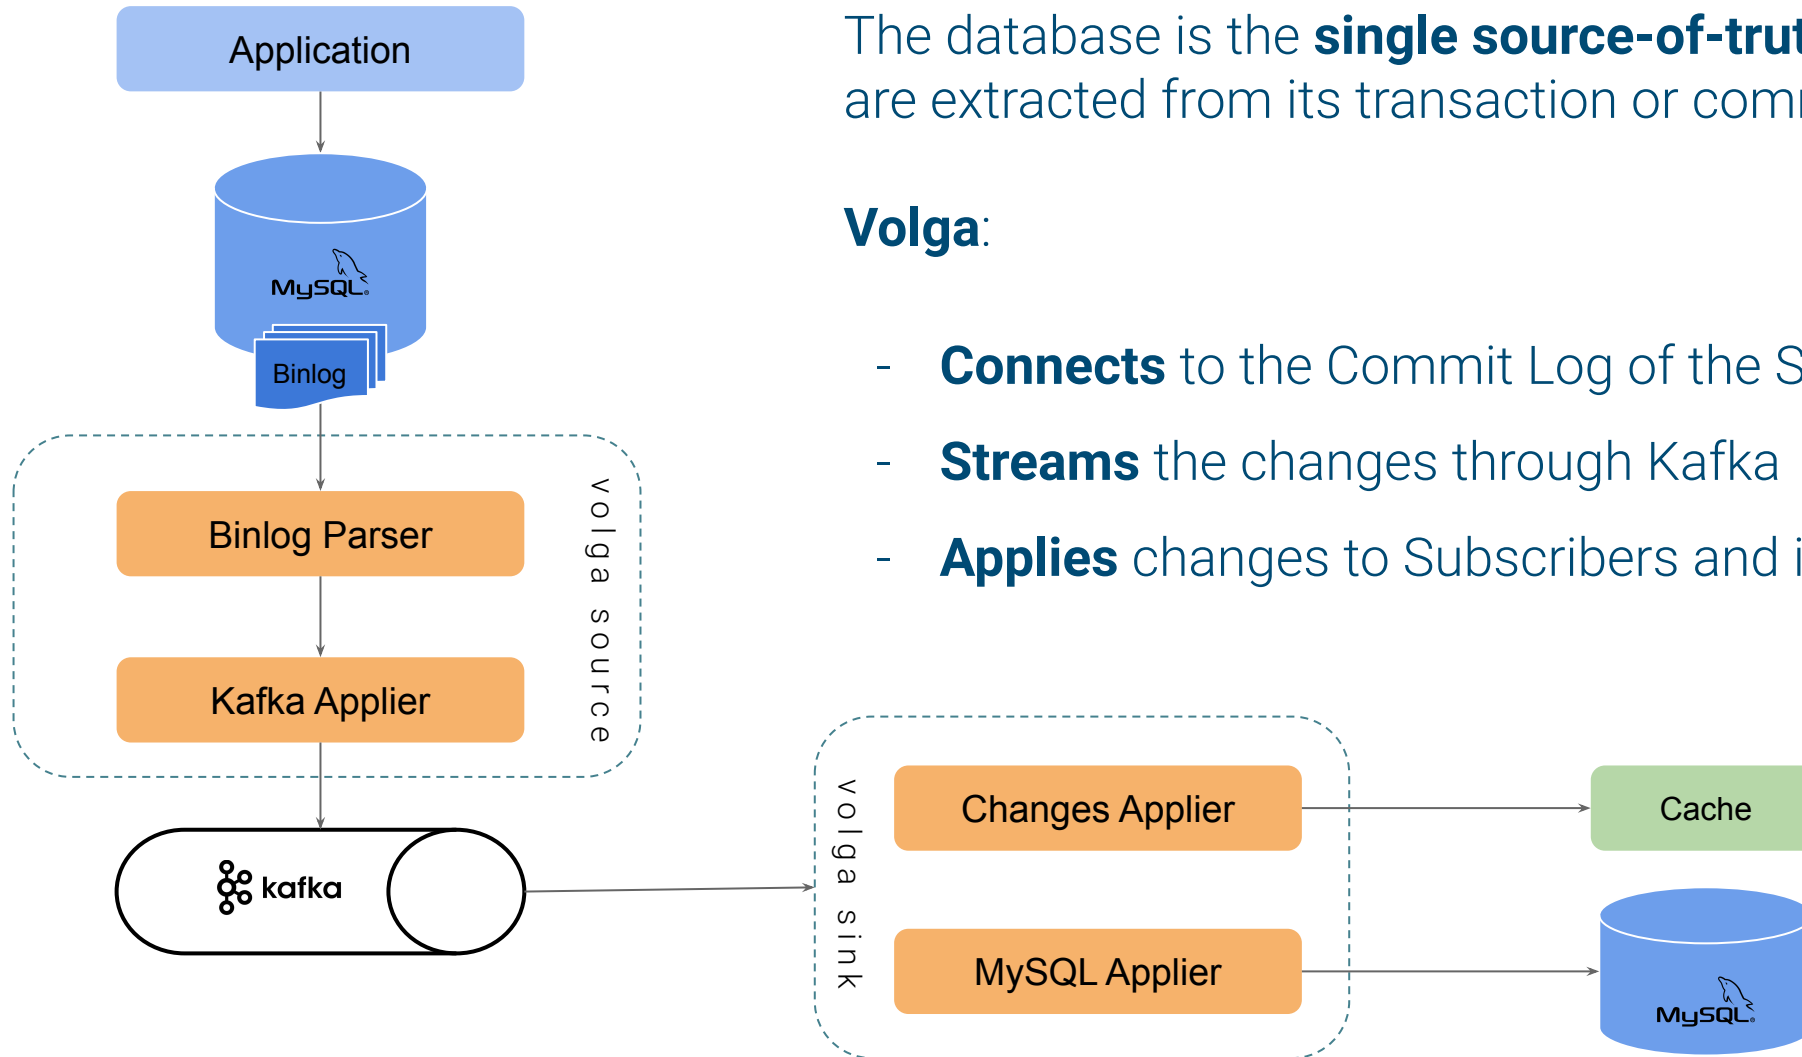
- **Connects** to the Commit Log of the Source DB
- **Streams** the changes through Kafka



The database is the **single source-of-truth** and changes are extracted from its transaction or commit log.

Volga:

- **Connects** to the Commit Log of the Source DB
- **Streams** the changes through Kafka
- **Applies** changes to Subscribers

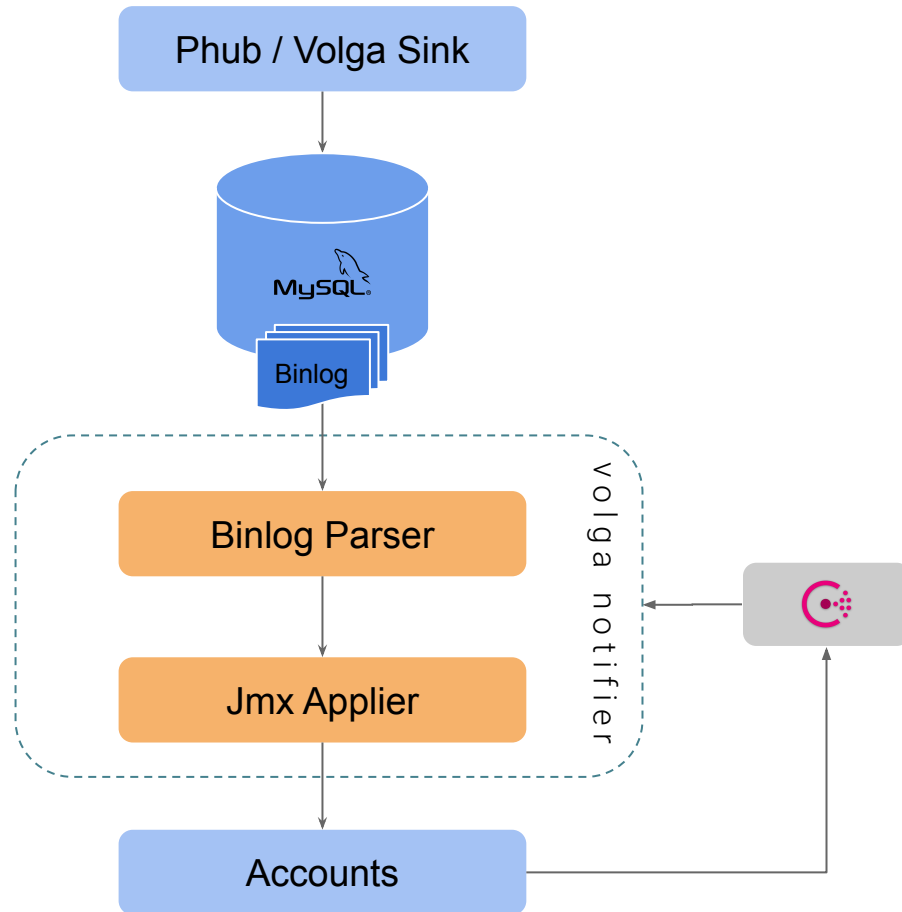


The database is the **single source-of-truth** and changes are extracted from its transaction or commit log.

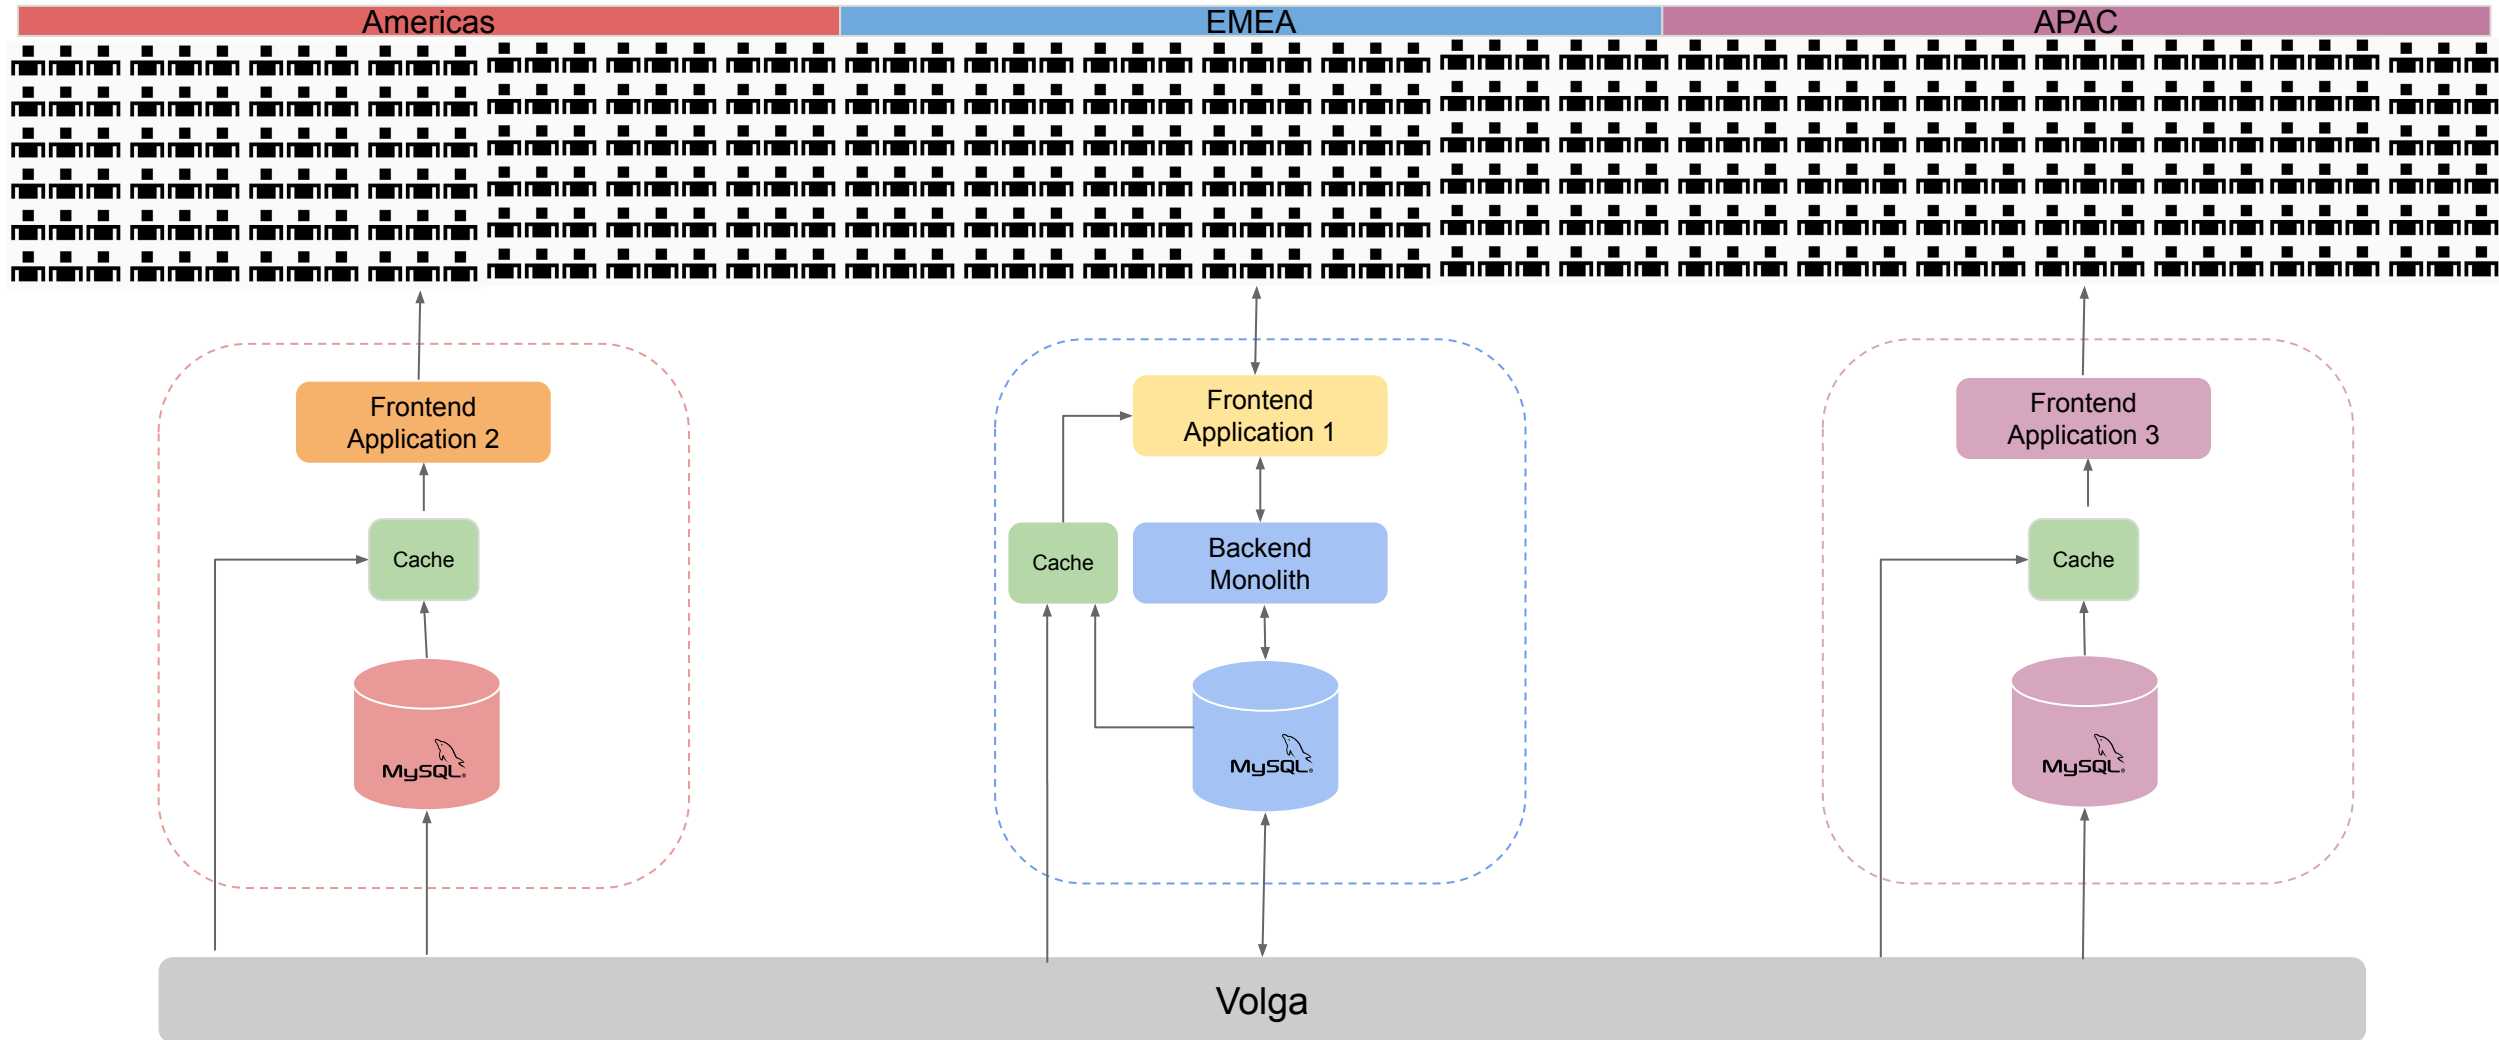
Volga:

- **Connects** to the Commit Log of the Source DB
- **Streams** the changes through Kafka
- **Applies** changes to Subscribers and in-sync Replicas

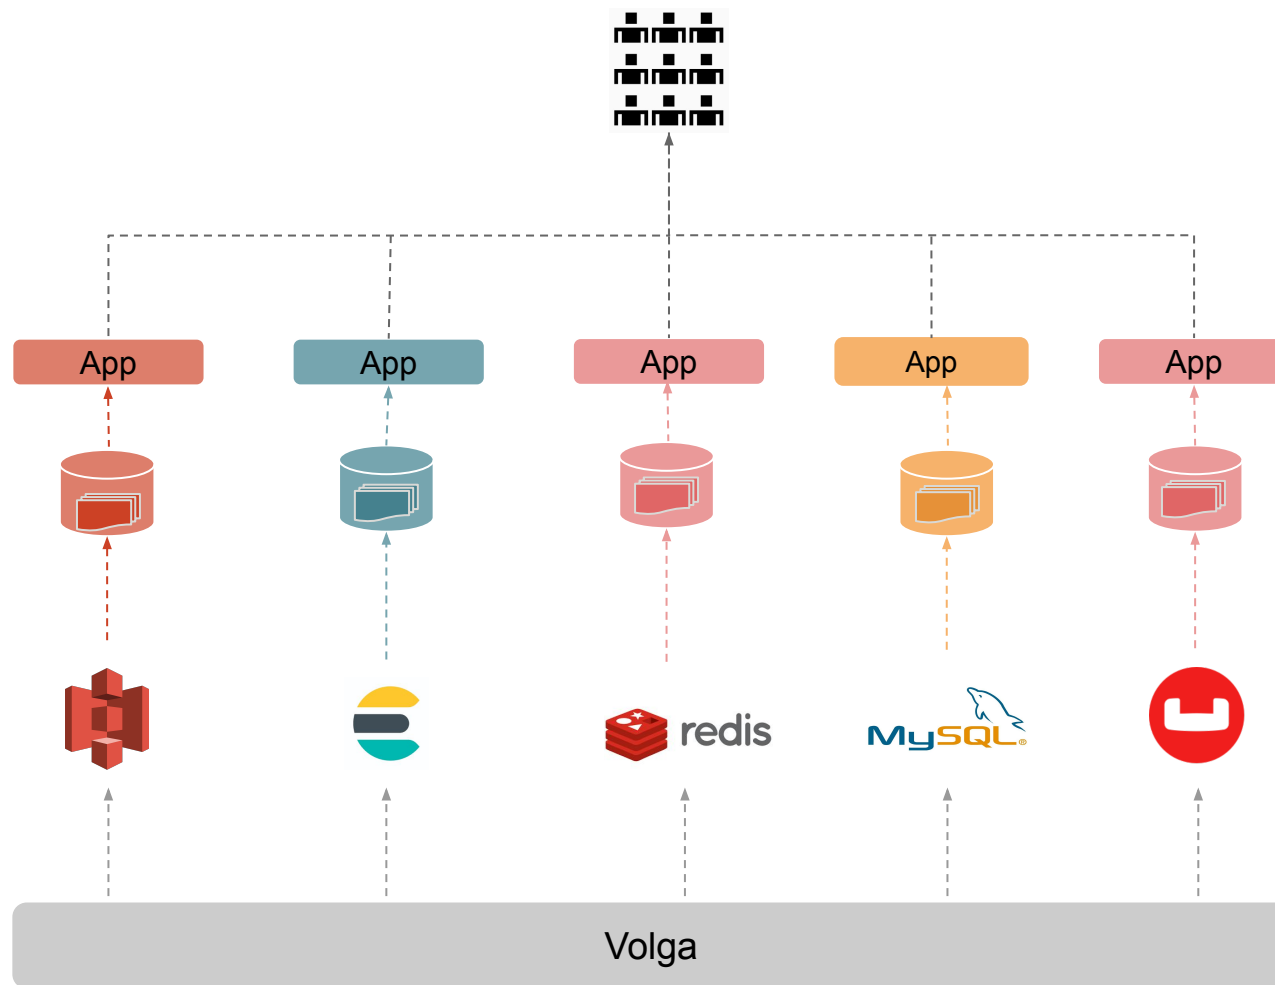
Log Mining with Volga



No Dual Writes with Volga



Make you data available for everyone, as they want



Via Volga, we can send data to:

- secondary indexes
- analytic platforms
- caches
- homogeneous in-sync replicas
- heterogeneous in-sync replicas



- **At-least-once** semantic
- **Sequential Consistency** (Binlog ordering + Kafka in-order delivery)

Additional Features:

- Snapshotting for Bootstrap, Correctness Check and DR
- Database Write Access Patterns Metrics
- Different level of partitioning (key-based, table, transactions)
- Homomorphic hashing



1.Scaling Read workload across regions ✓

2.**Monolithic Architecture**



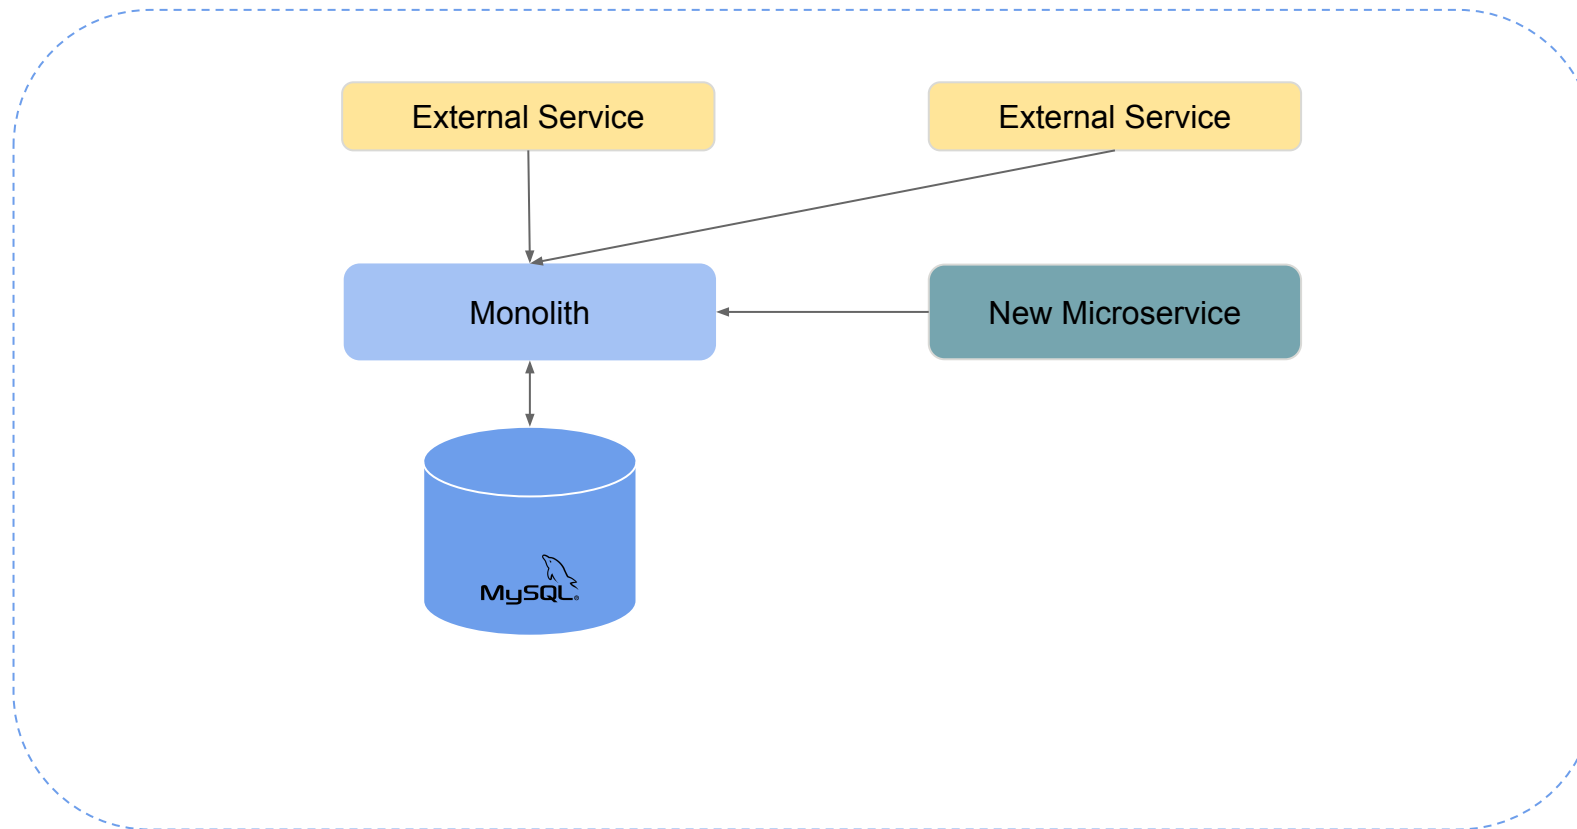
1. Incapacity to independently scale business domains
2. Barrier to innovation
3. Difficult coordination
4. Developers bad mood

The migration can be decomposed in 3 steps:

- **Transform** - Creation of the new Microservice extracting the new generation of a subset of functionalities from the Monolith
- **Coexist** - The functionalities are exposed from both the Monolith and the Microservice at the same time
- **Eliminate** - Remove the functionalities from the monolith




Strangler Pattern: Transform



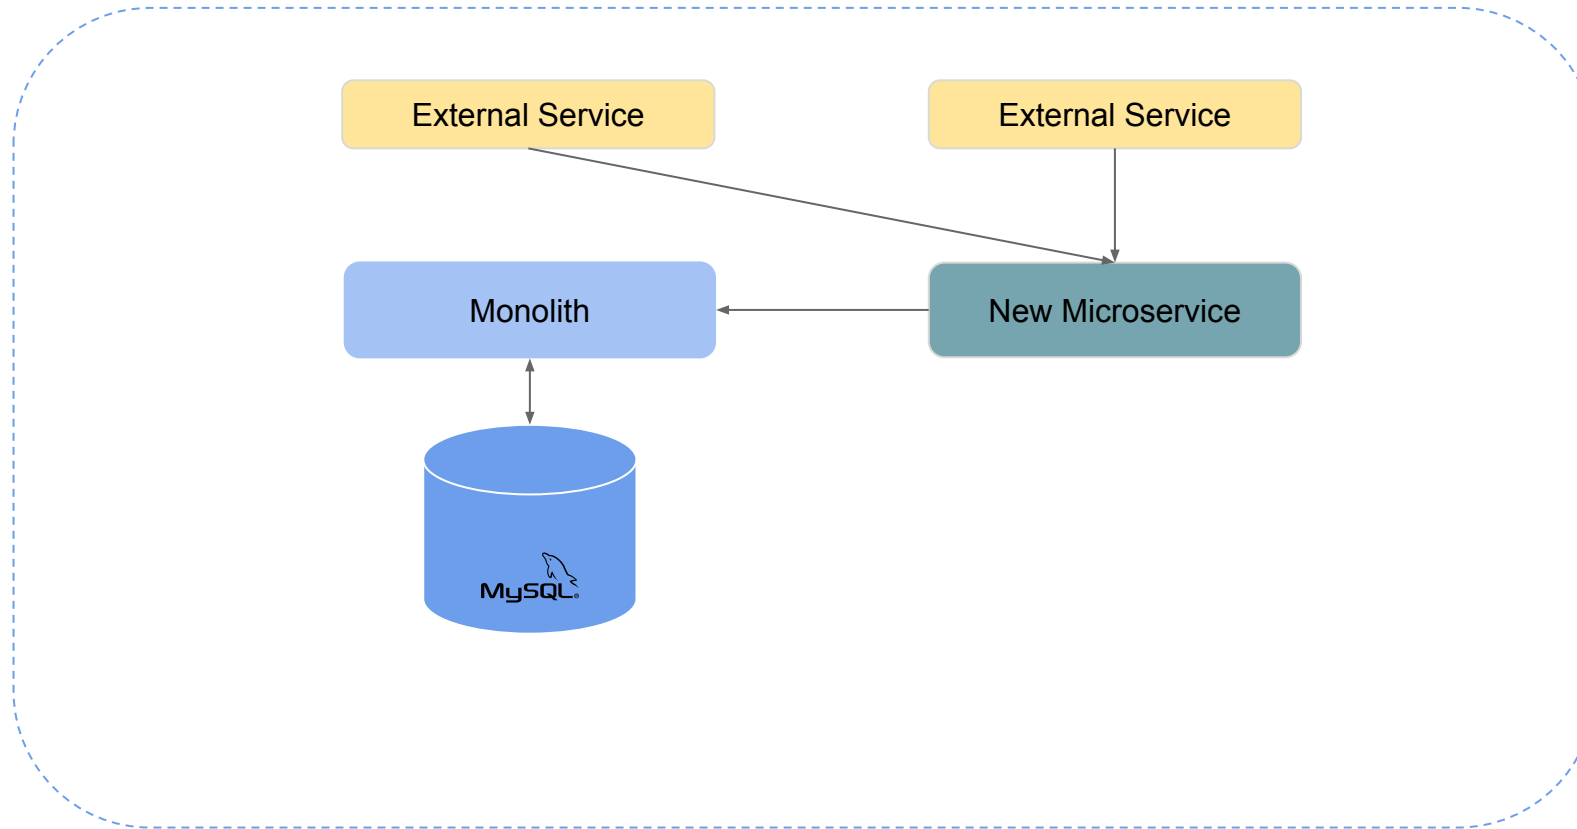
Build the new microservice still dependent on the Monolith

Strangler Pattern: Transform



Migrate the dependencies


Strangler Pattern: Transform



How do we migrate the dependency to the data store?

CQRS stands for Command Query Responsibility Segregation

Reads and Writes are segregated into separate models

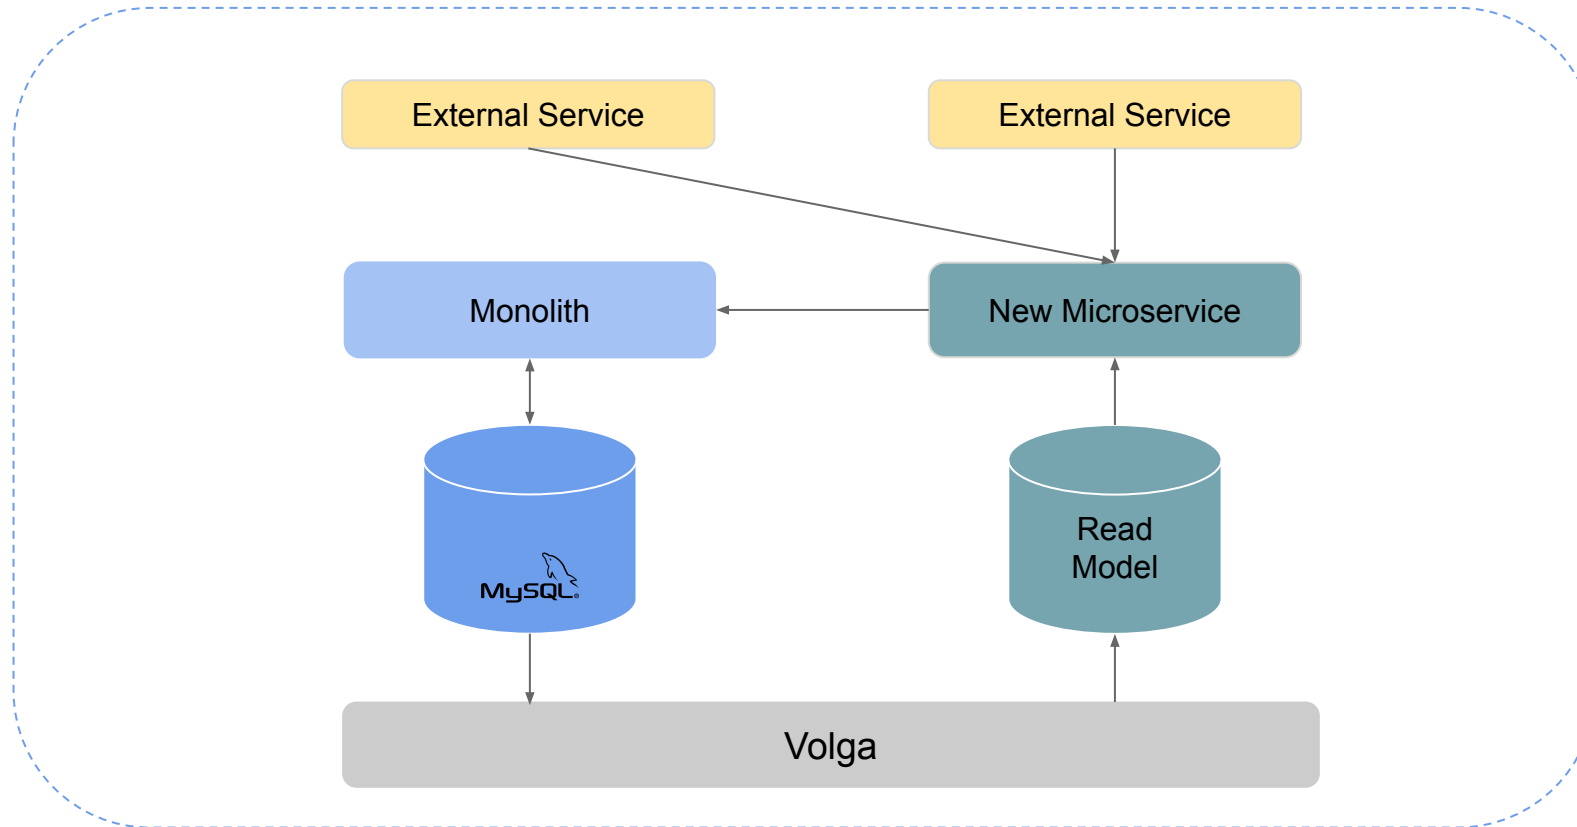


Benefits of Command Query Responsibility Segregation



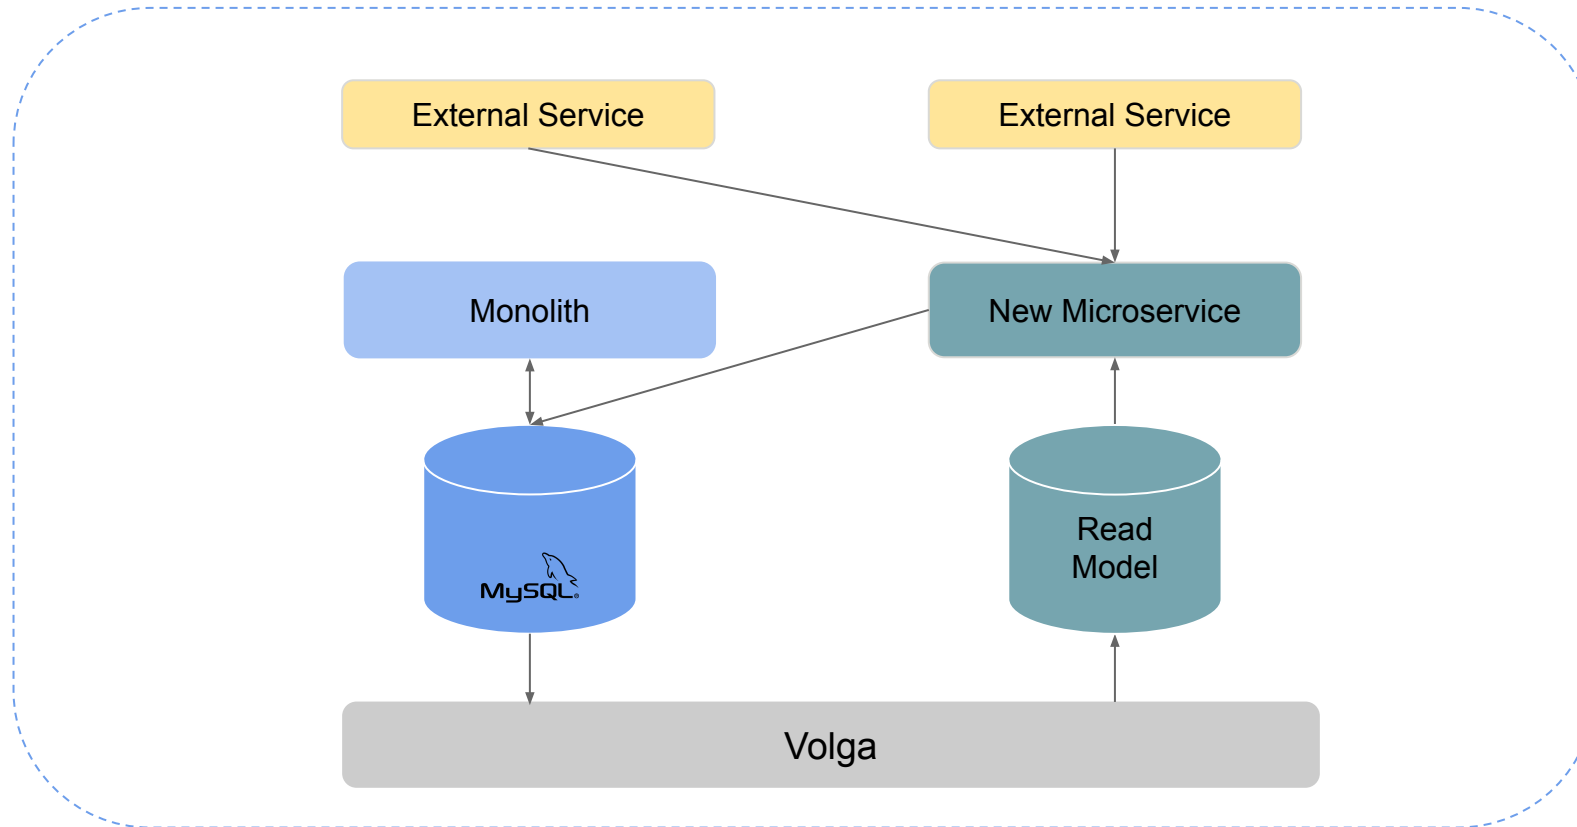
- Independent scaling of write and read workloads
- Independent Optimized data schemas
- Separation of concerns
- Improved security
- Loose coupling
- Support for incremental Migration

Strangler Pattern: Coexist



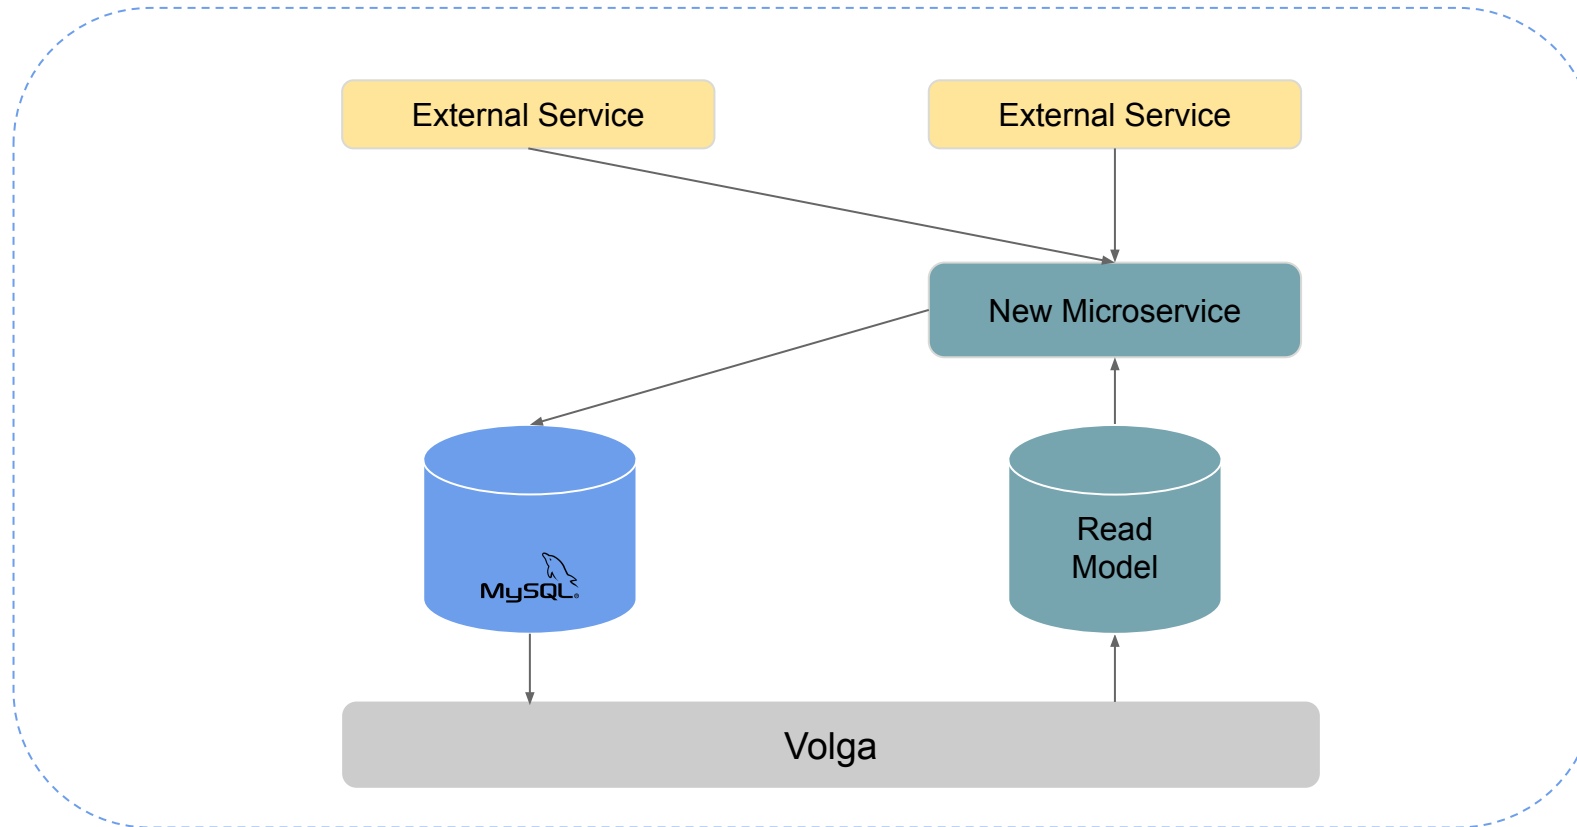
Writes are served by the old model, Reads by the new Read Model

Strangler Pattern: Eliminate

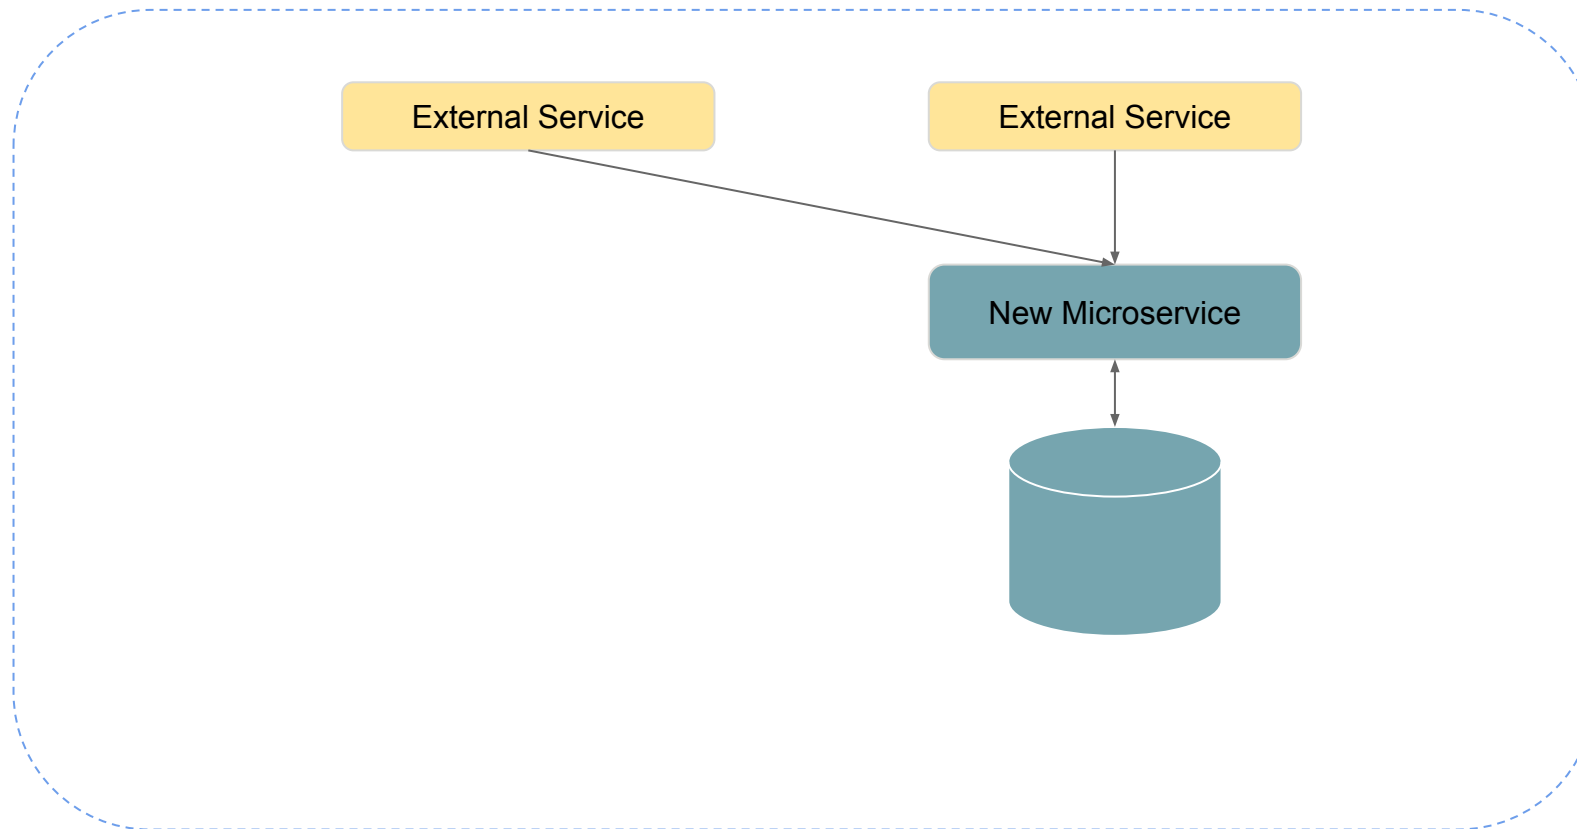


Remove the dependency to the Monolith

Strangler Pattern: Eliminate



Remove the dependency to the Monolith



Migrate the Write Model

Problem 2: The monolith




1. Read workload ✓

2. Drawbacks of the Monolithic Architecture ✓

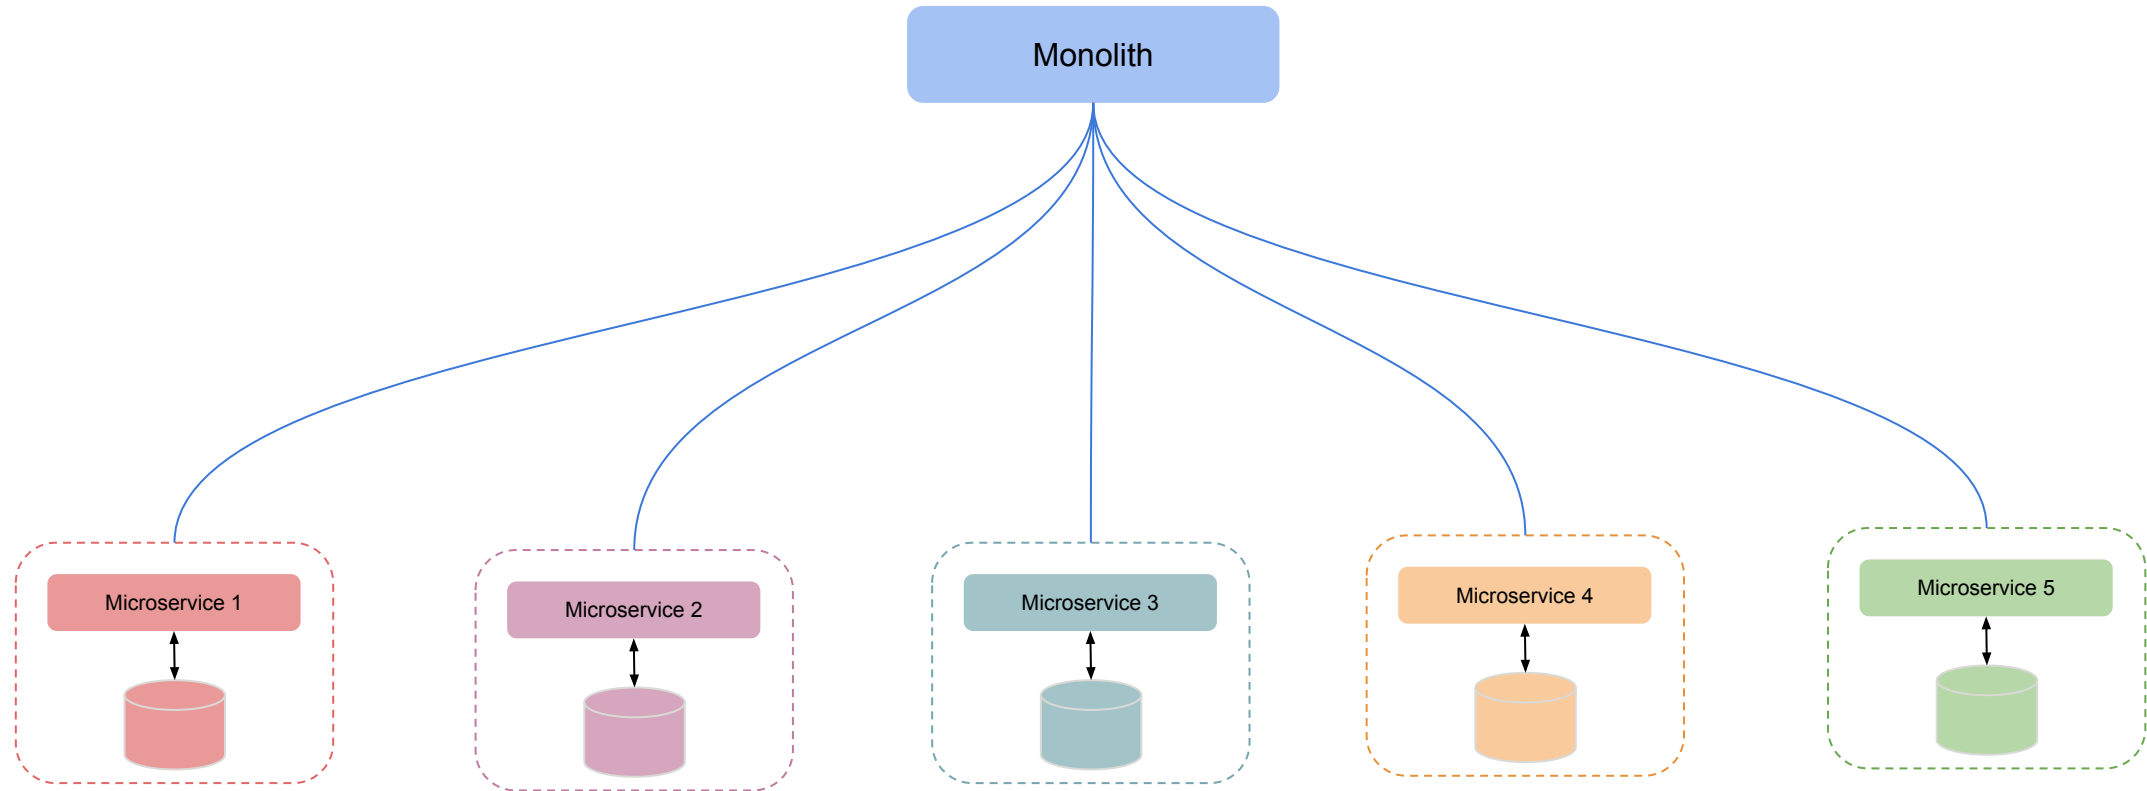
The Delight of Transactions



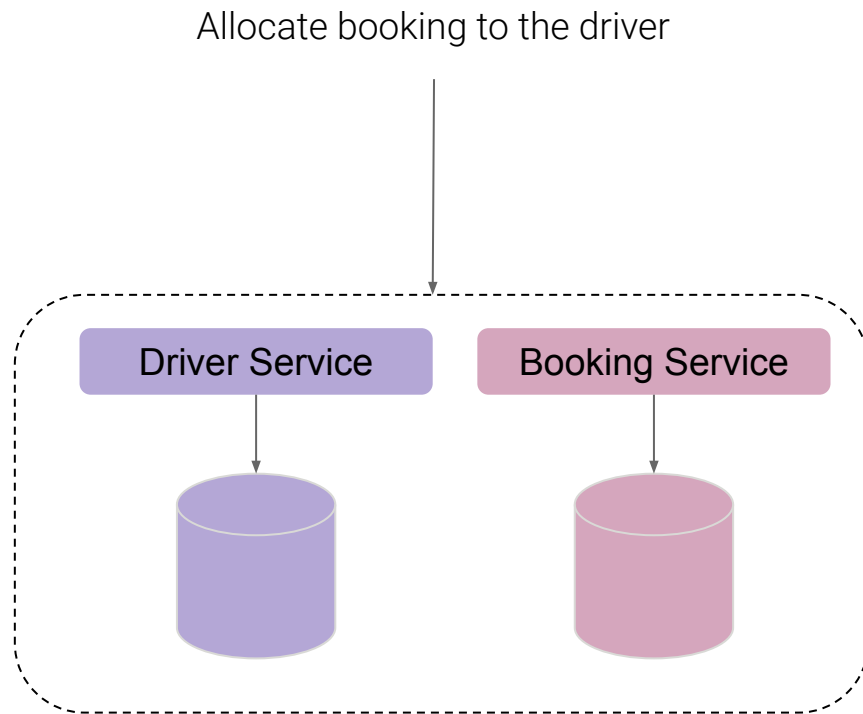
Allocate booking to the driver



```
Start Transaction;  
set driver_status=busy;  
set booking_status=allotted;  
Commit Transaction;
```



The Horror of transactions



```
Start Transaction;  
set driver_status=busy  
set booking_status=allotted  
Commit Transaction;
```

How to implement Transactions with Microservices?

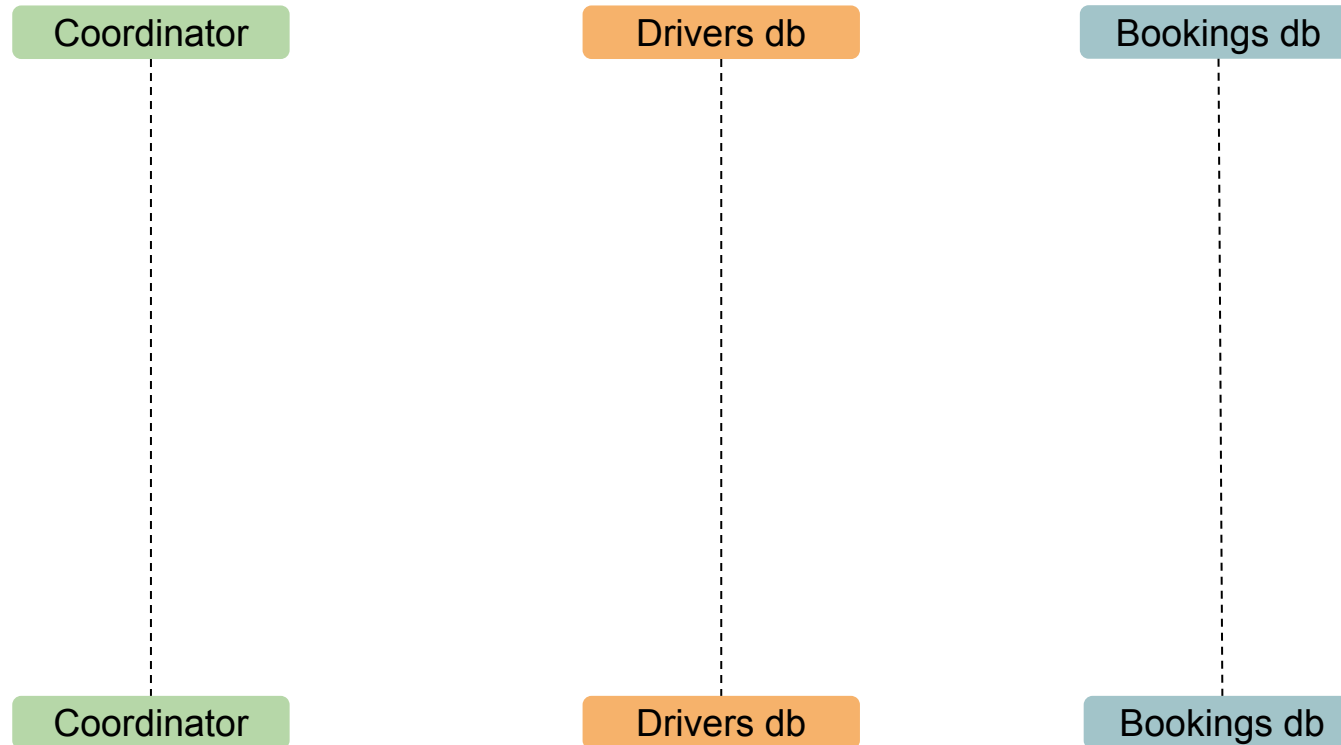


2-Phase Commit

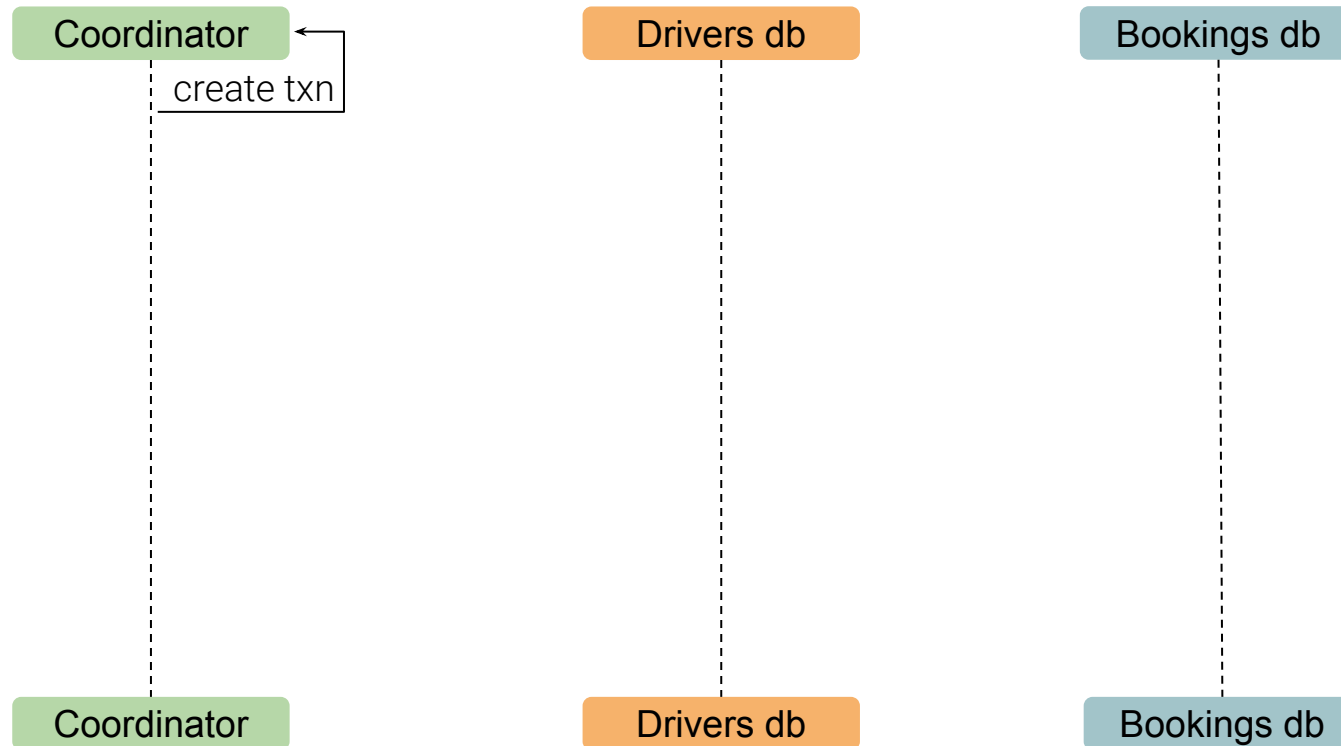


Sagas

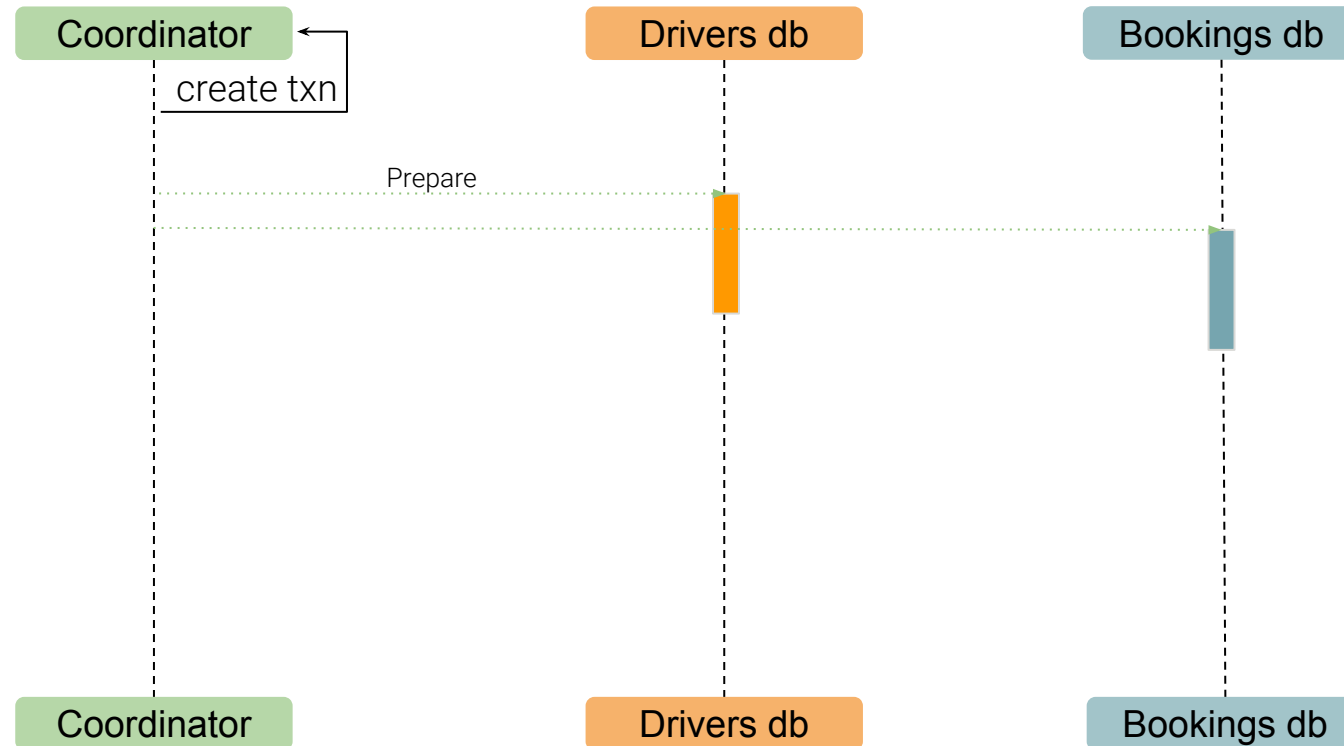
2-Phase Commit: XA Transaction



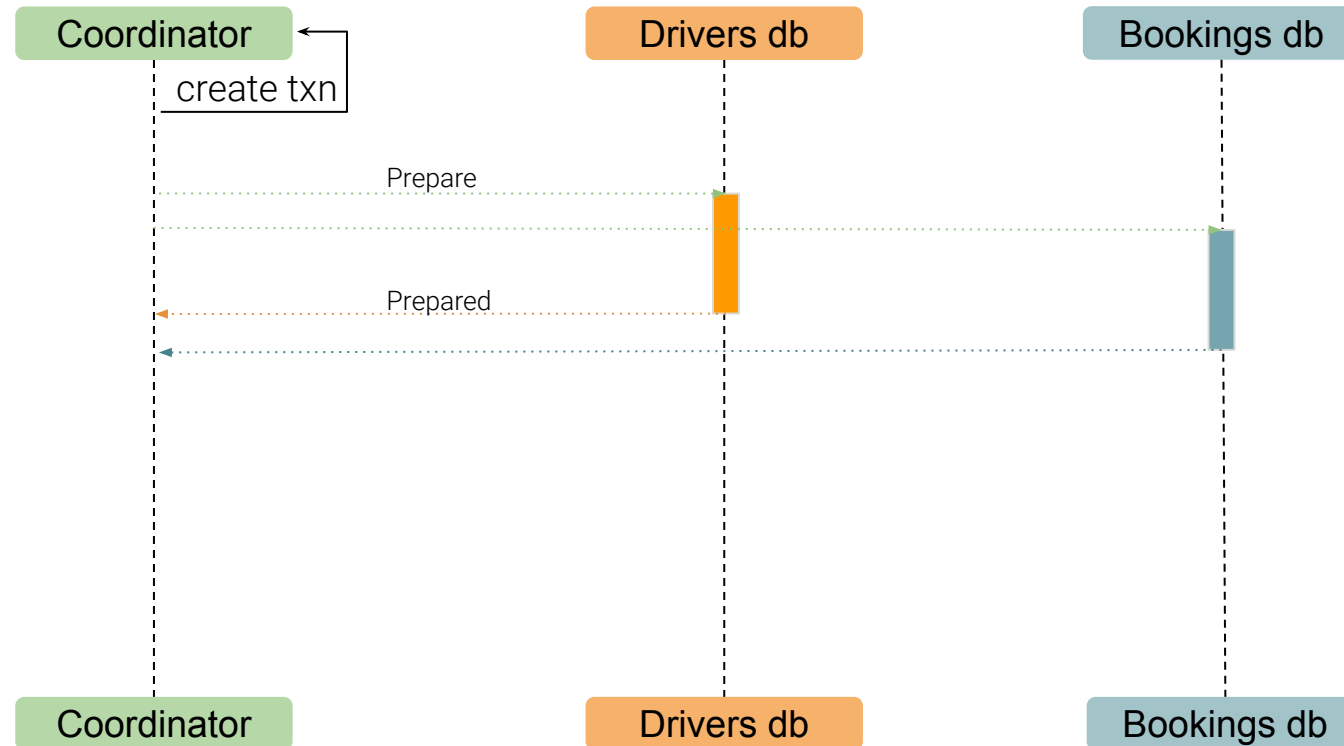
2-Phase Commit: XA Transaction



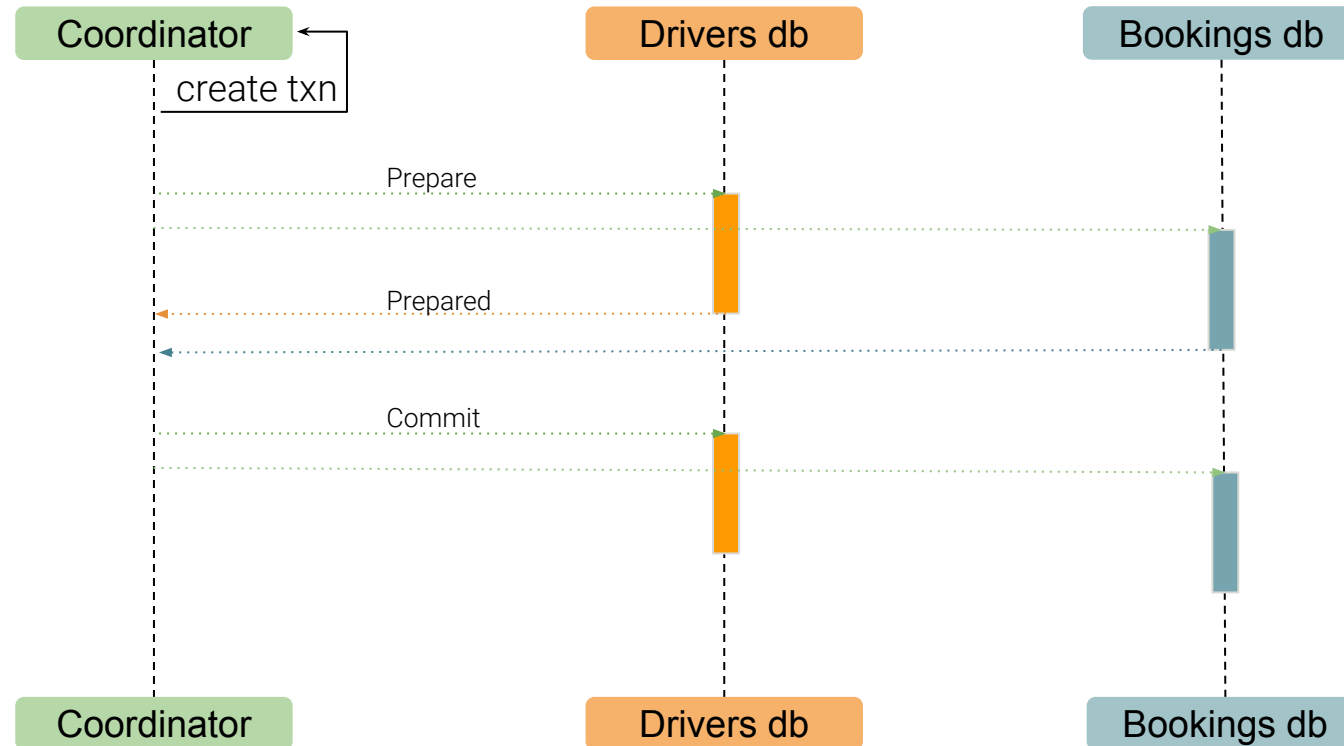
2-Phase Commit: XA Transaction



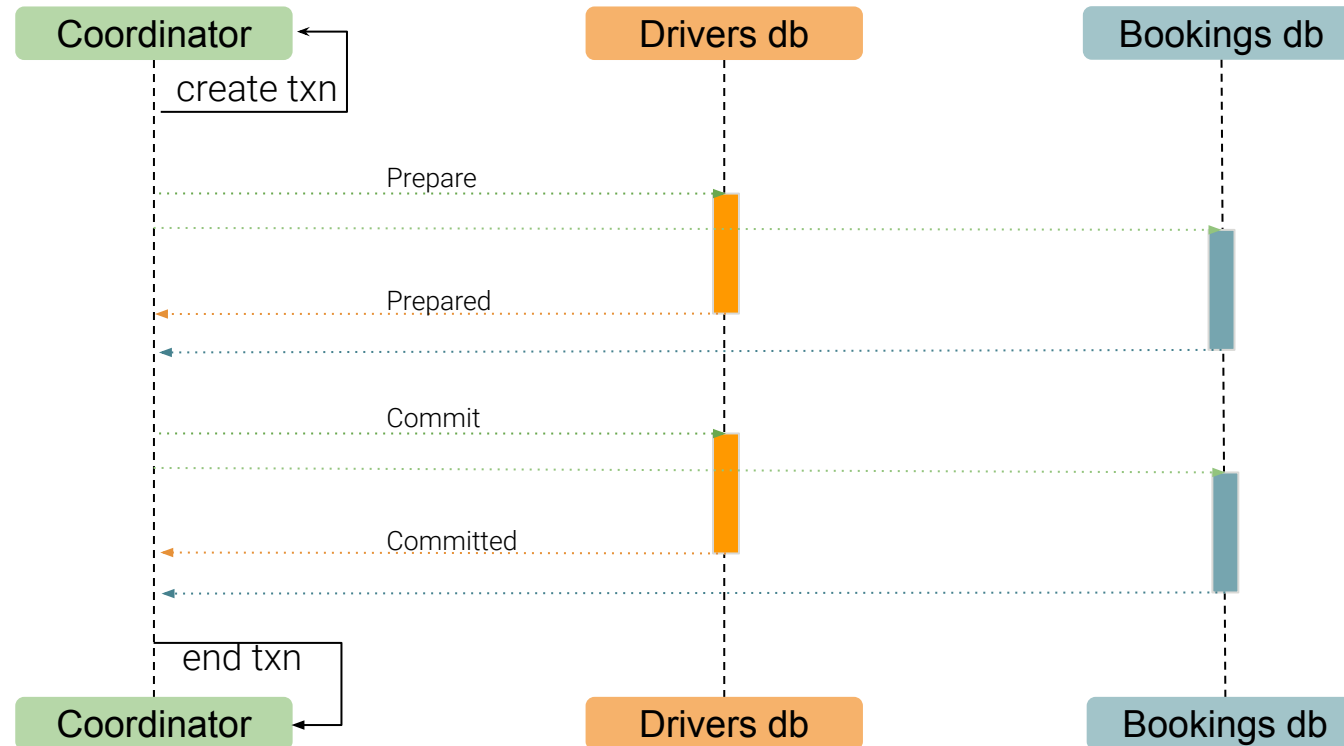
2-Phase Commit: XA Transaction



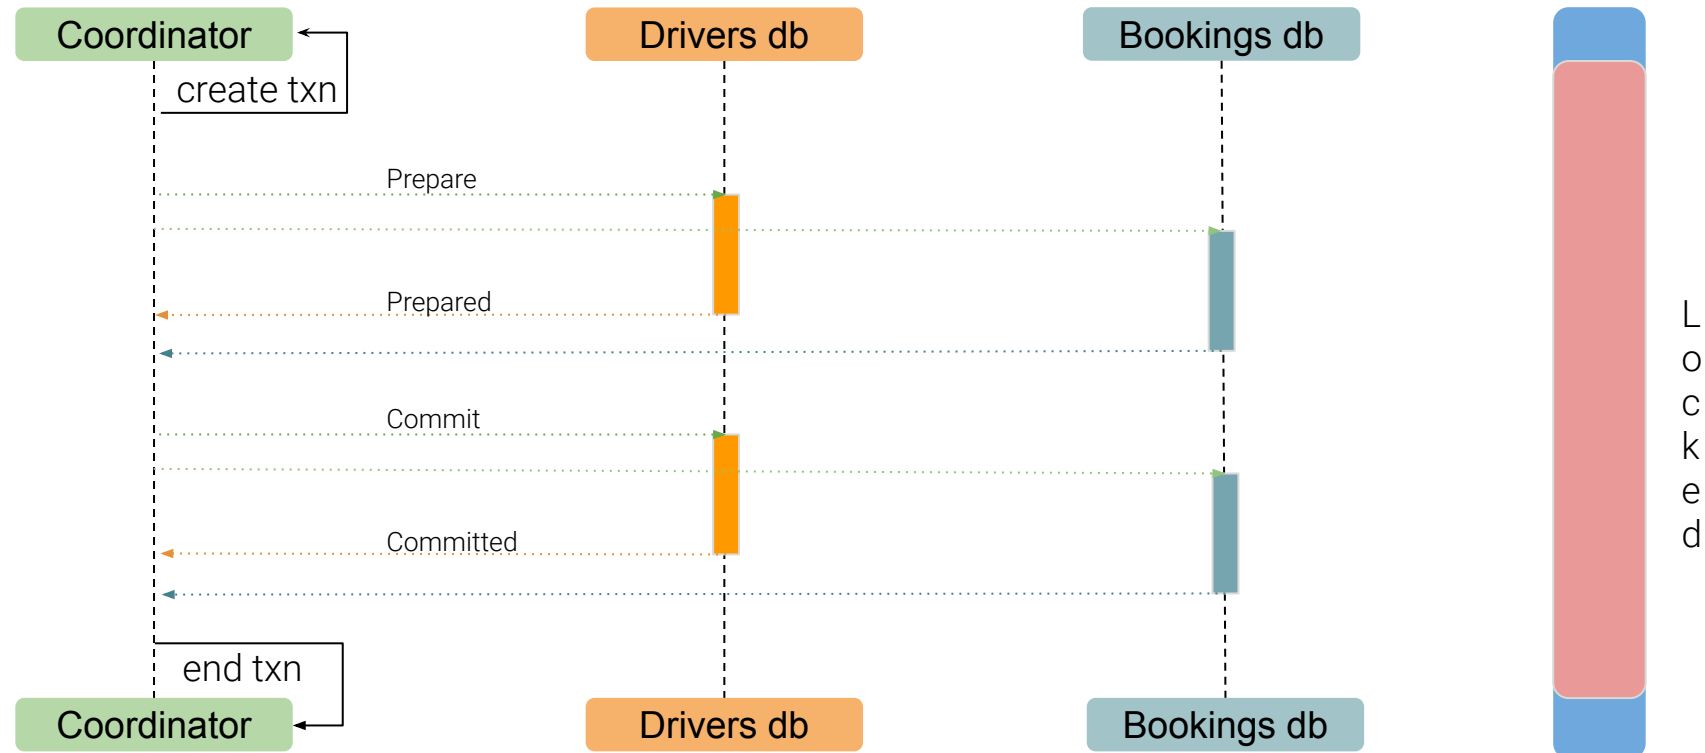
2-Phase Commit: XA Transaction



2-Phase Commit: XA Transaction



2-Phase Commit: XA Transaction

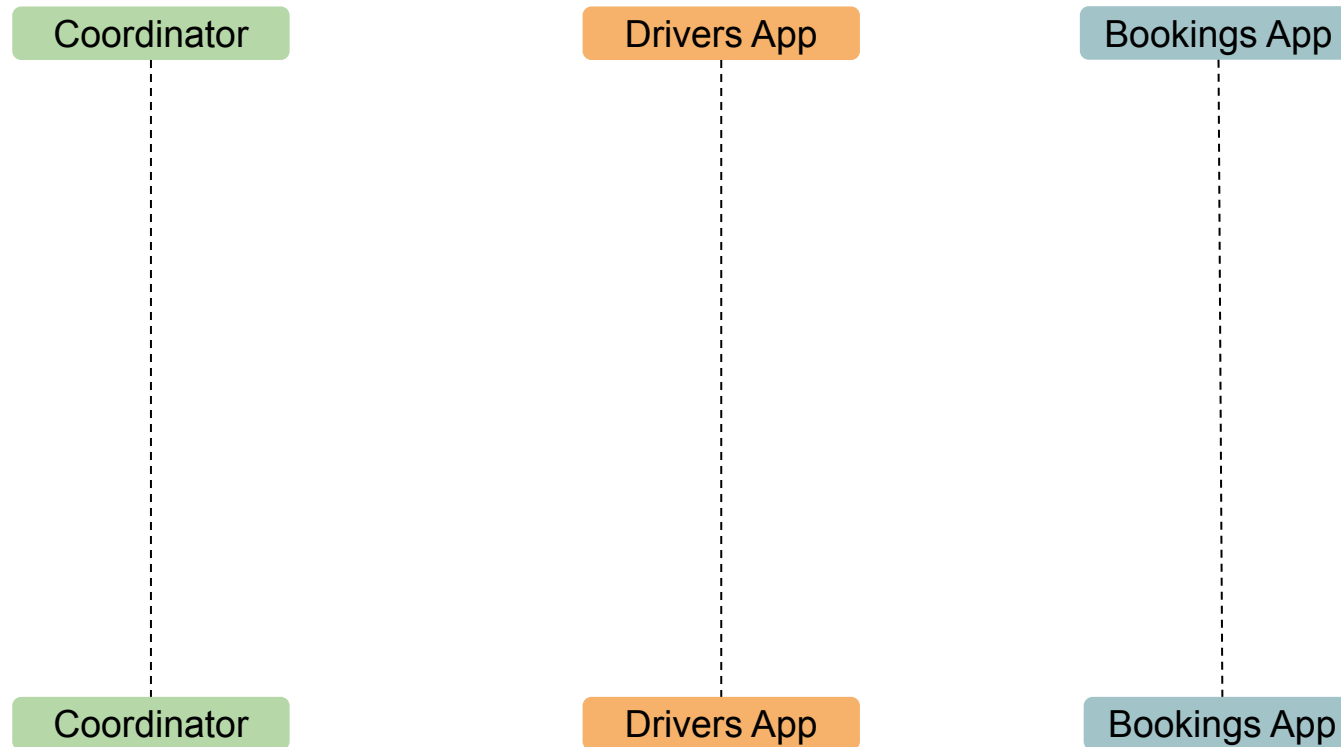




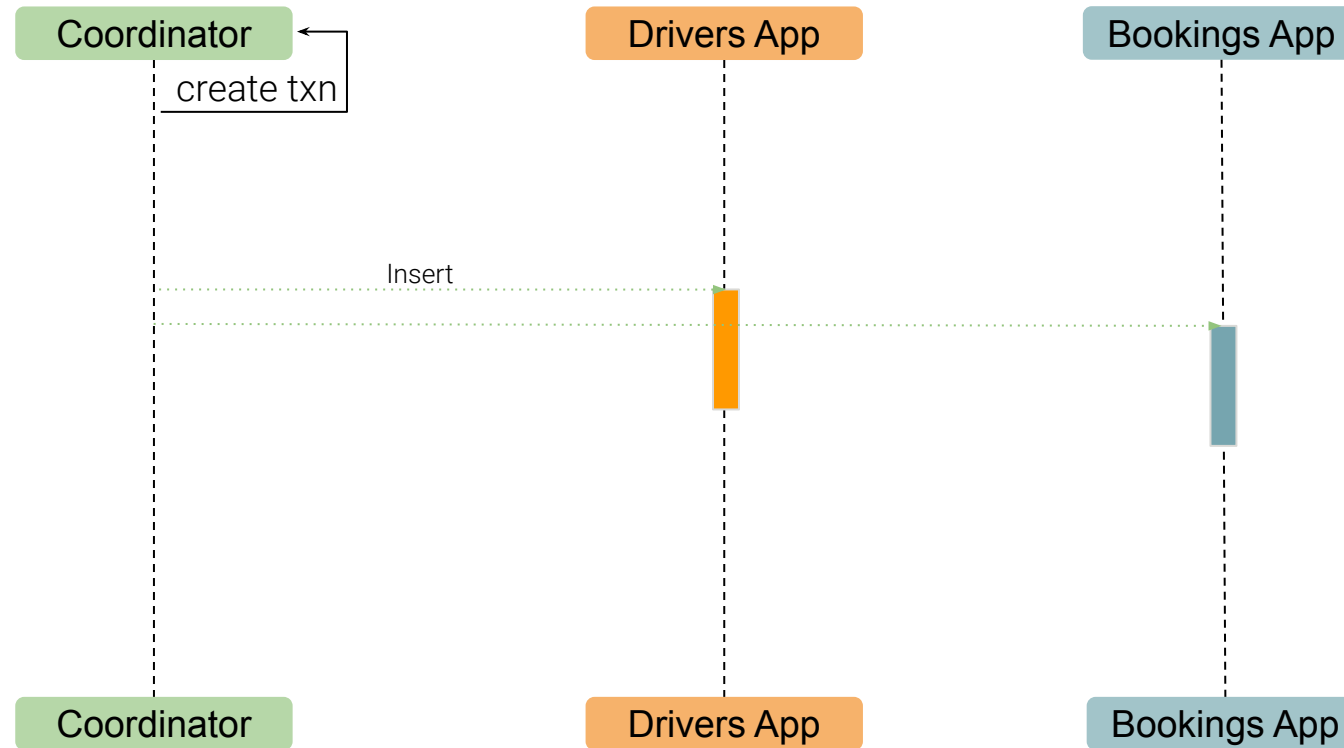
- Atomic Transaction
- Strong Consistency
- Read Write Isolation



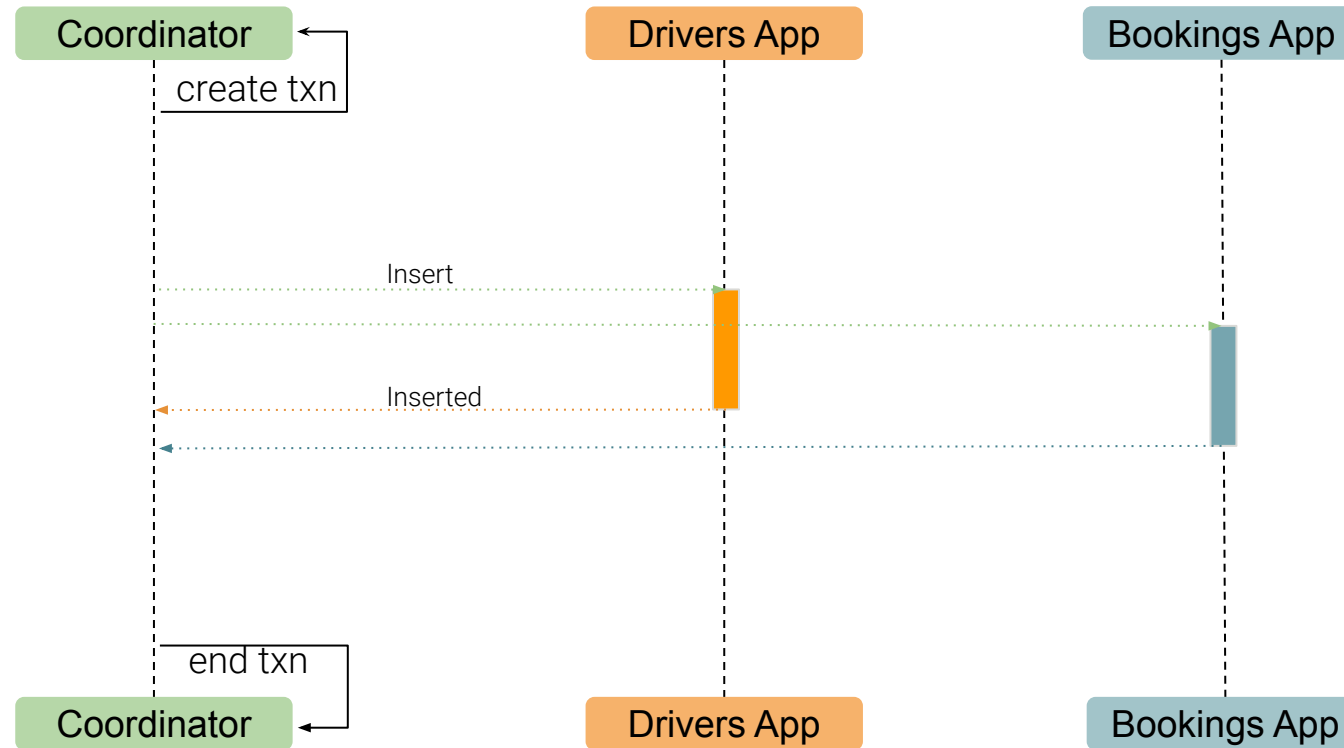
- Slow (due to Locks)
- Deadlock



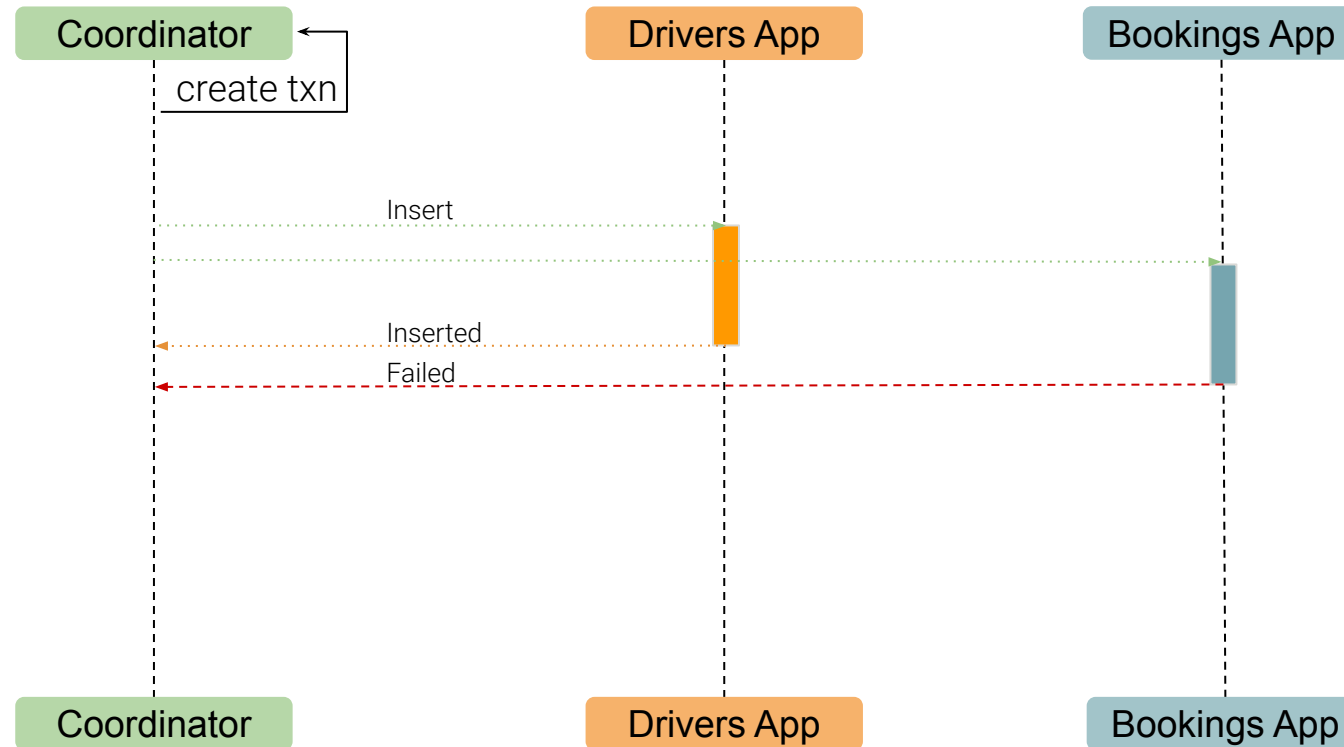
Sagas: Transaction



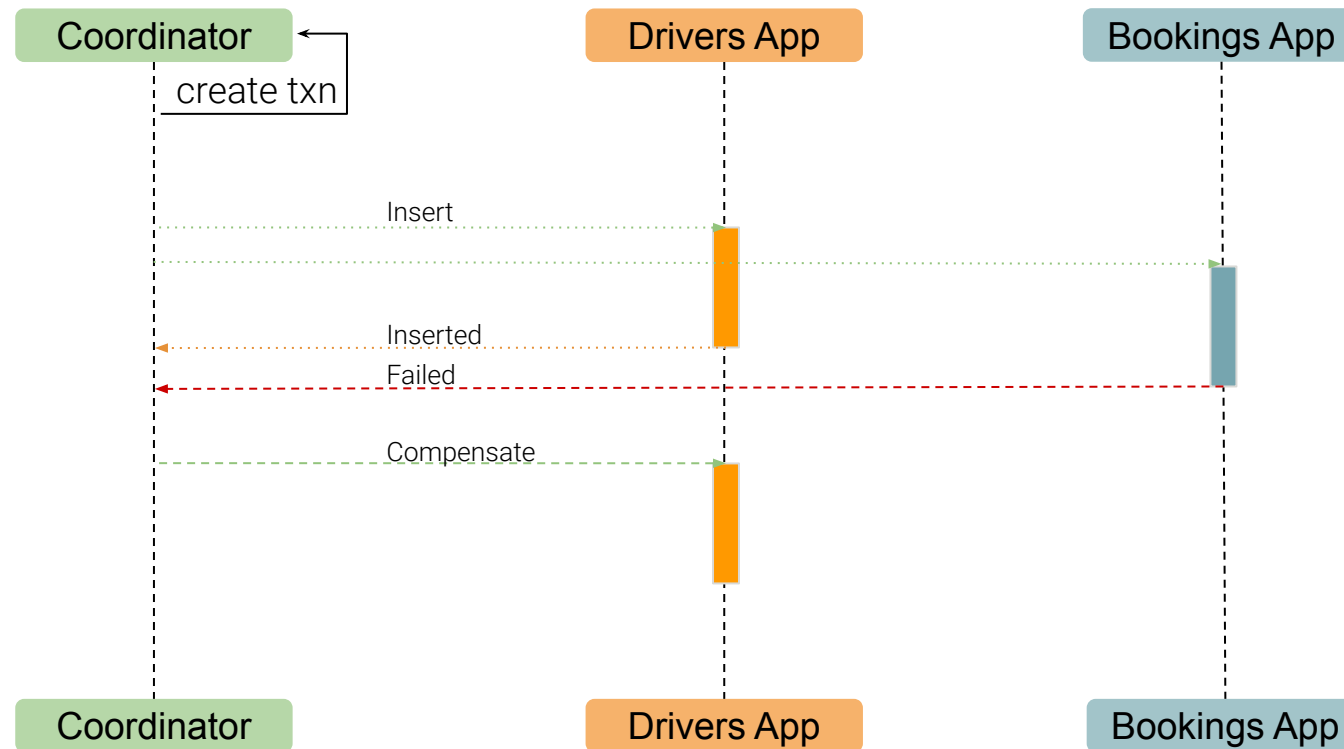
Sagas: Transaction



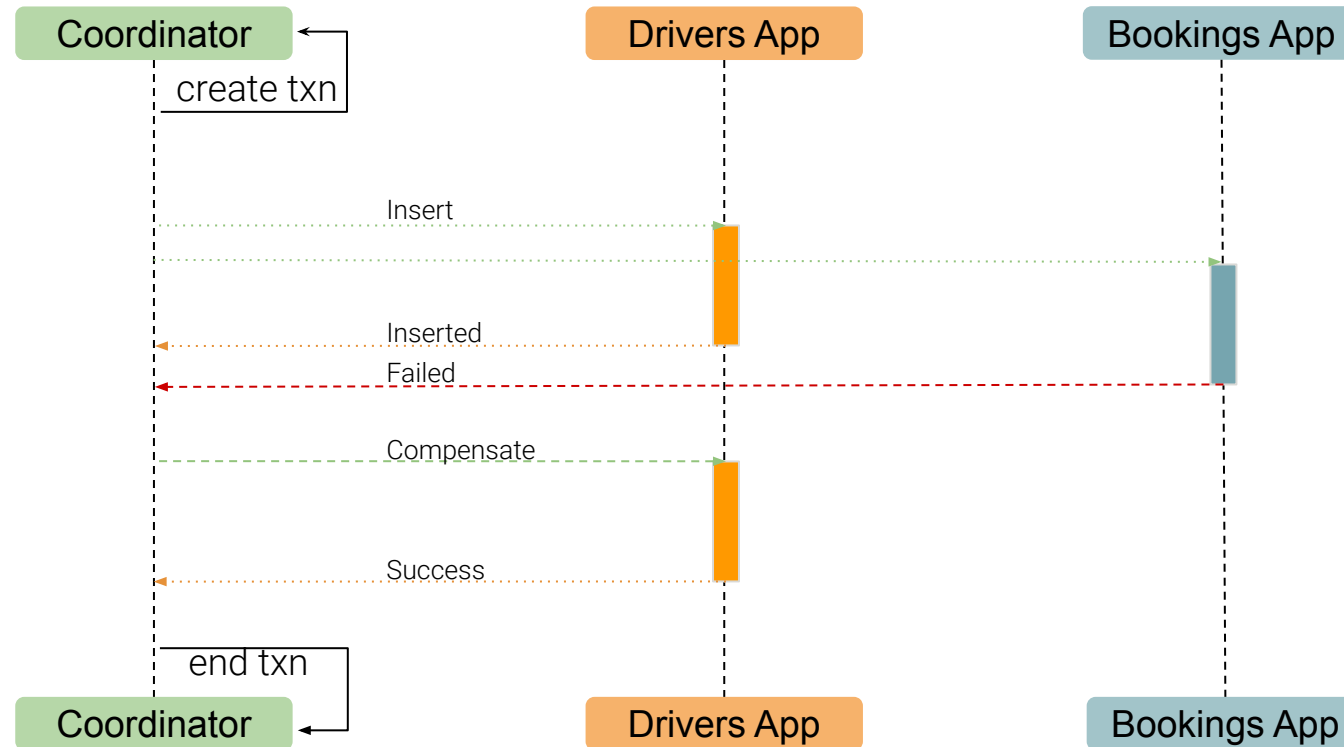
Sagas: Transaction



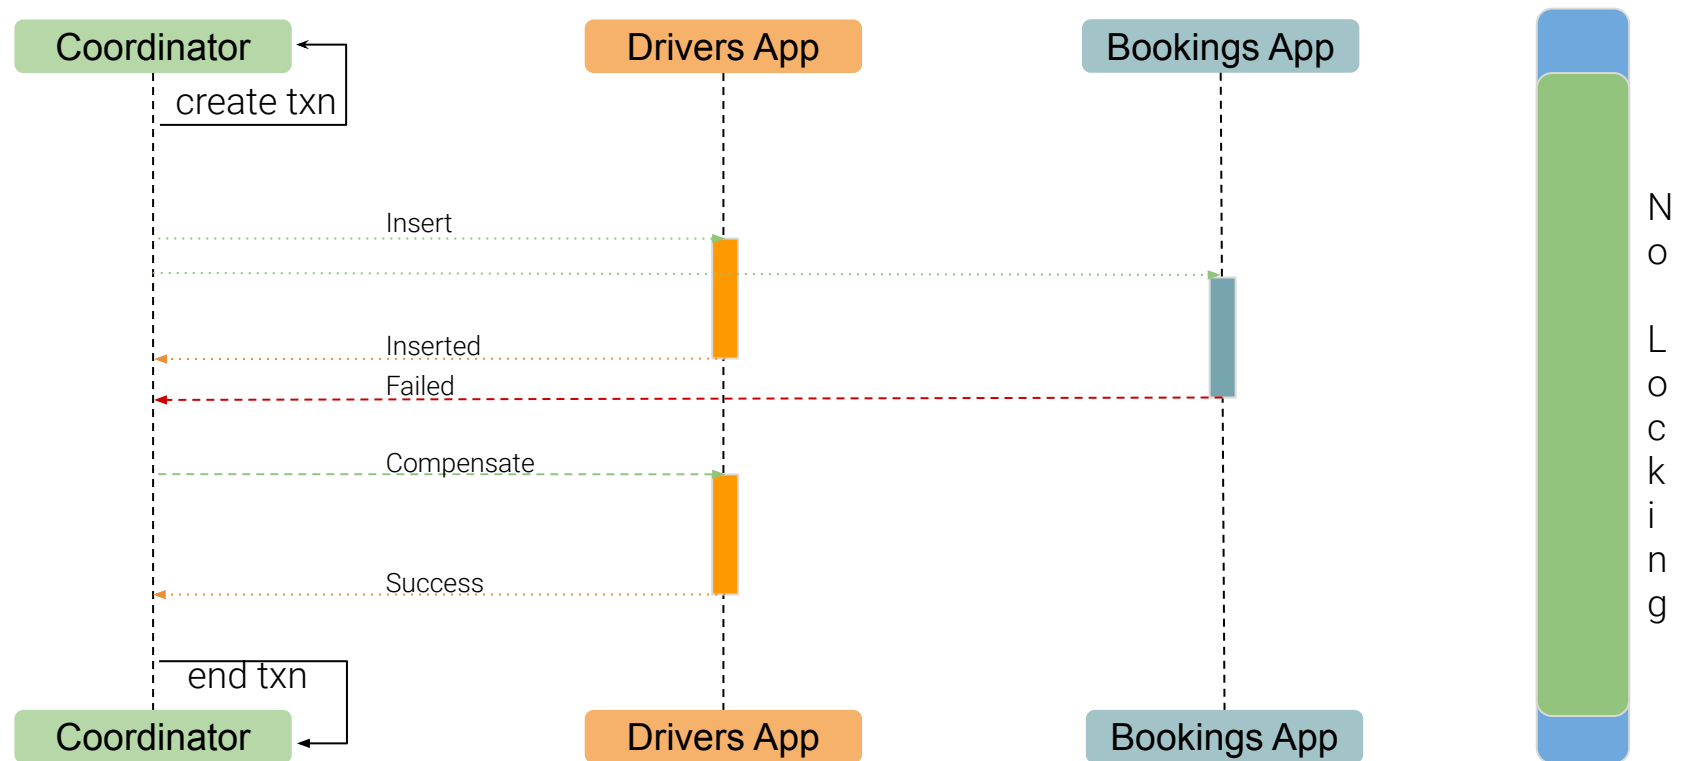
Sagas: Transaction



Sagas: Transaction



Sagas: Transaction





- No Locks
- Long Lived Transaction
- No tight coupling

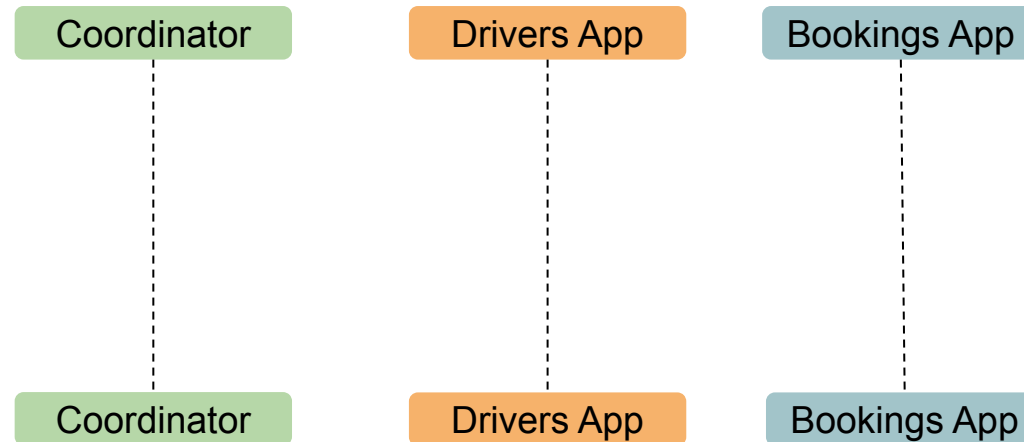


- Relaxed Isolation
- Difficult Compensations
- Bad User Experience

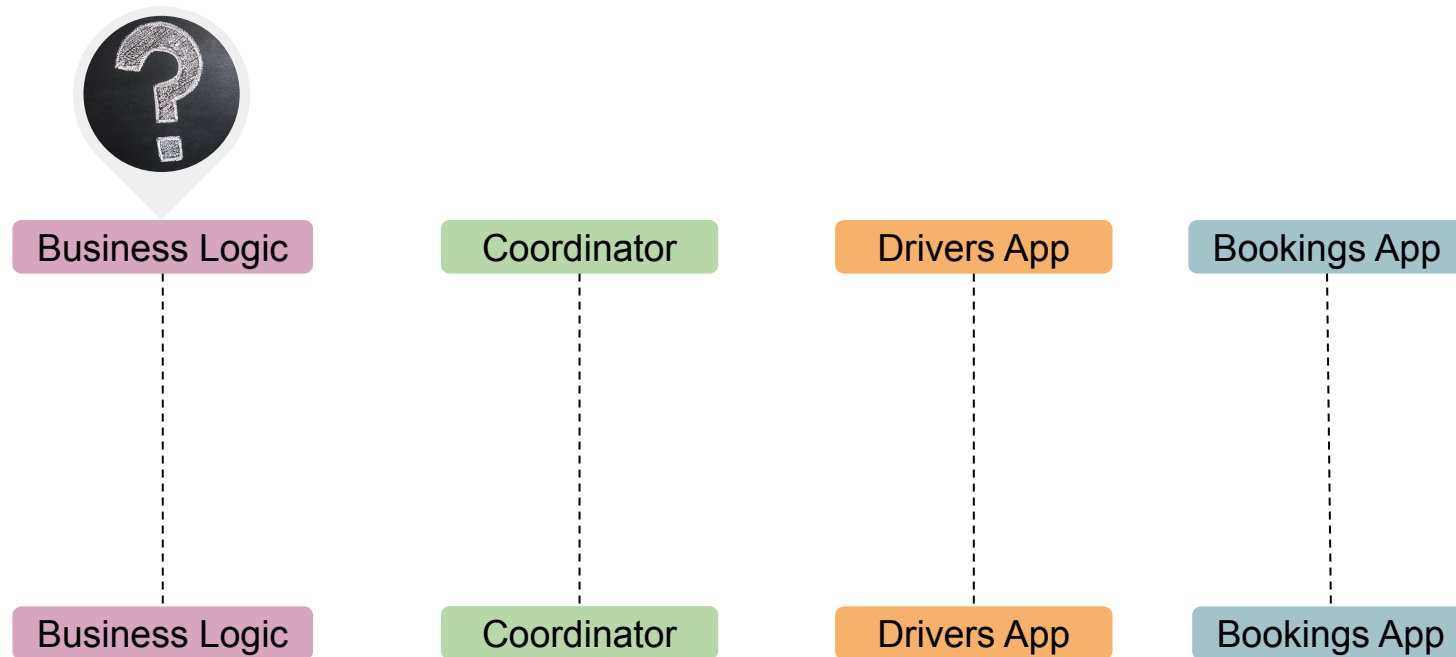
Is there any Better way?



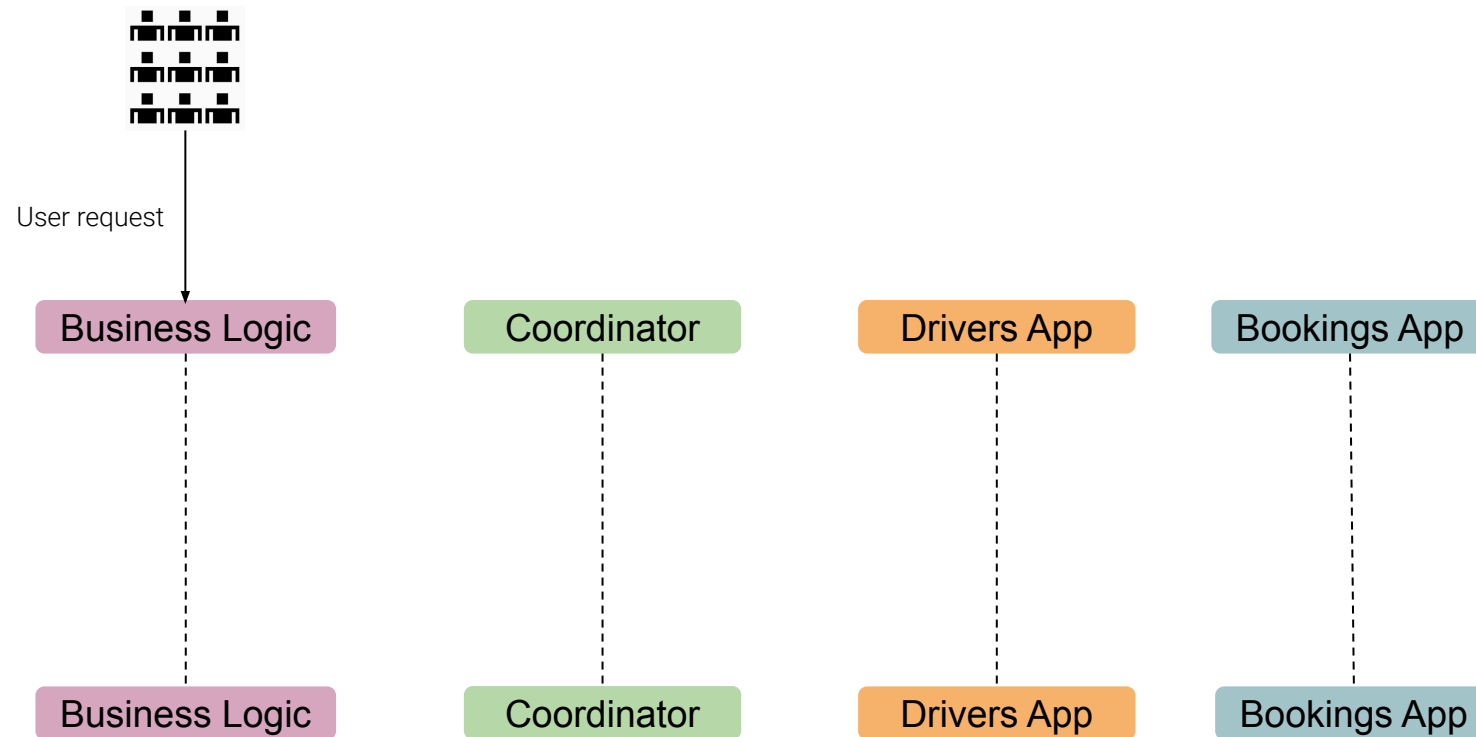
Shadows: A Distributed 2 Phase Commit



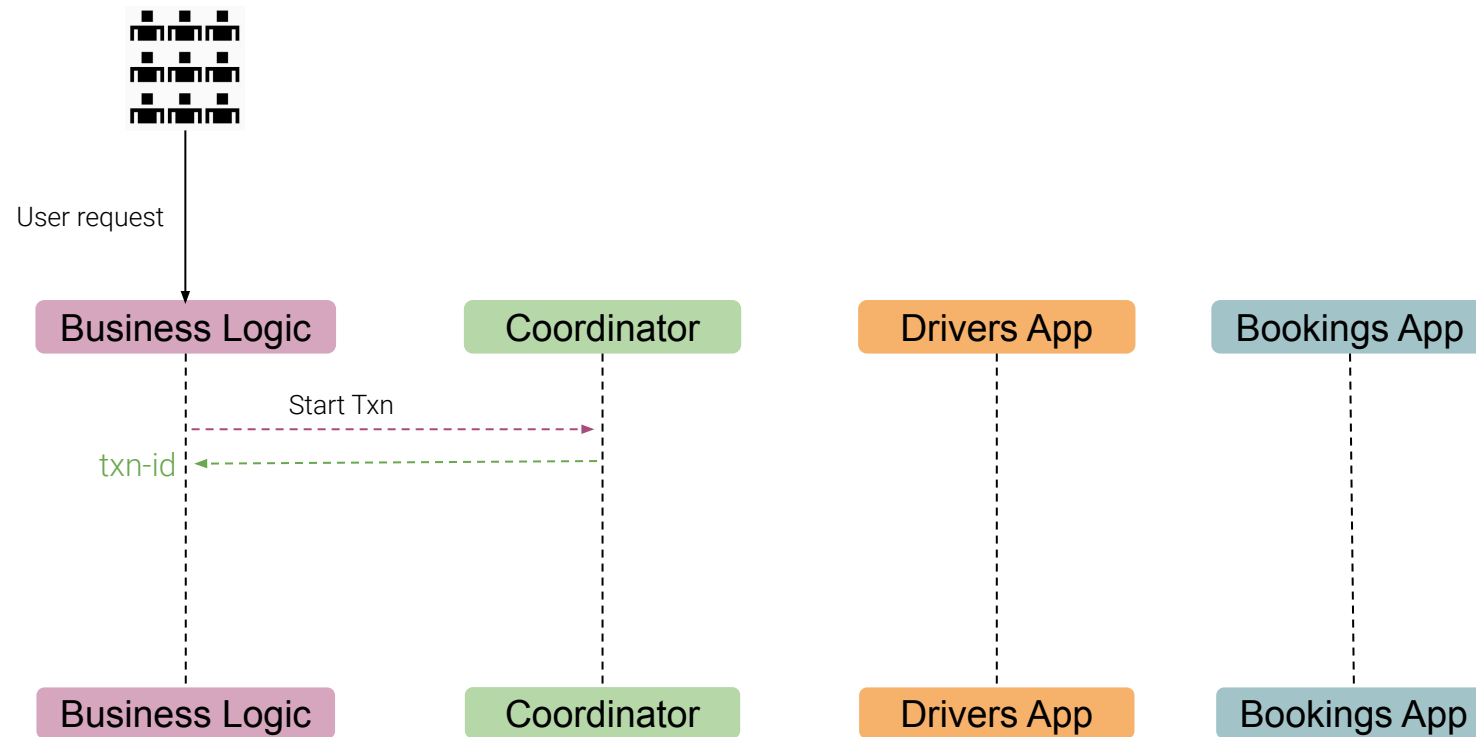
Shadows: Business Logic?



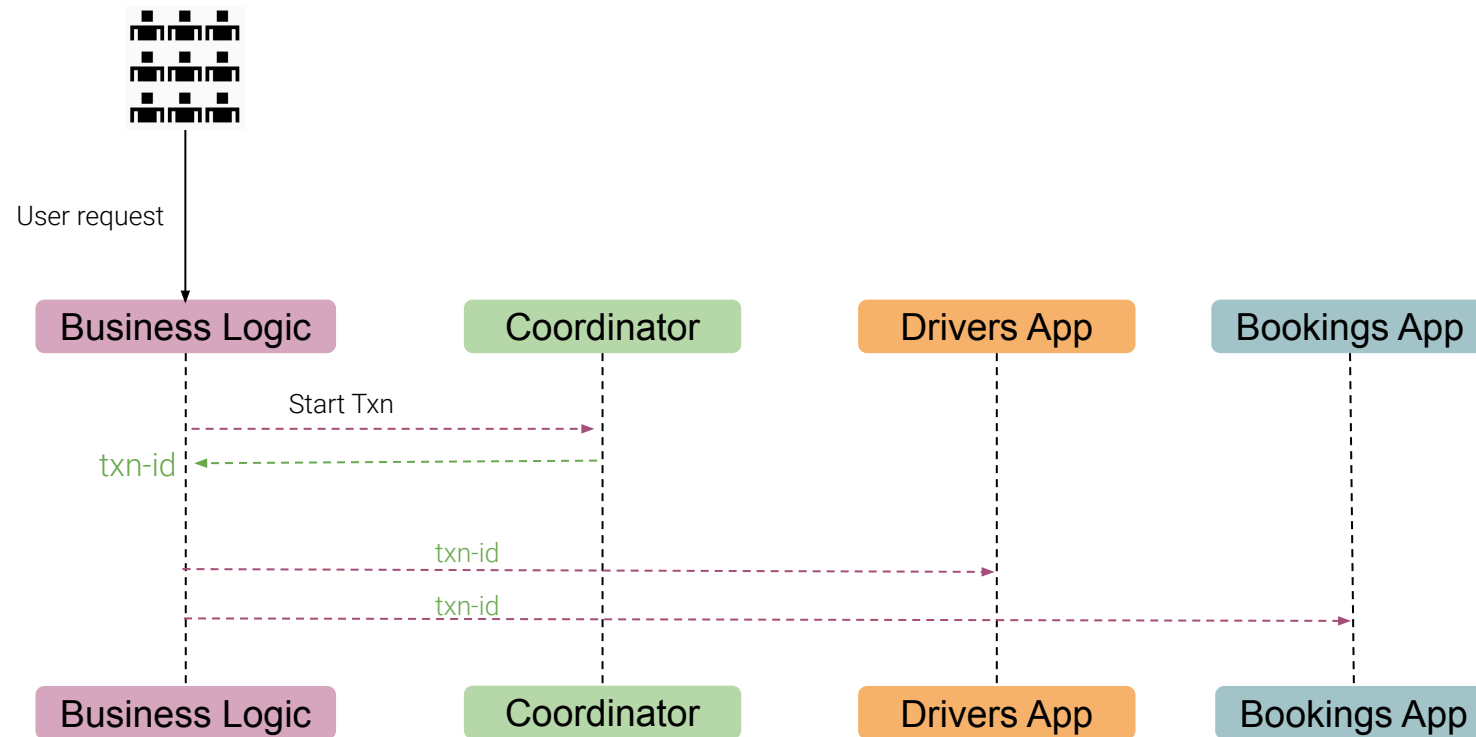
Shadows: A Distributed 2 Phase Commit



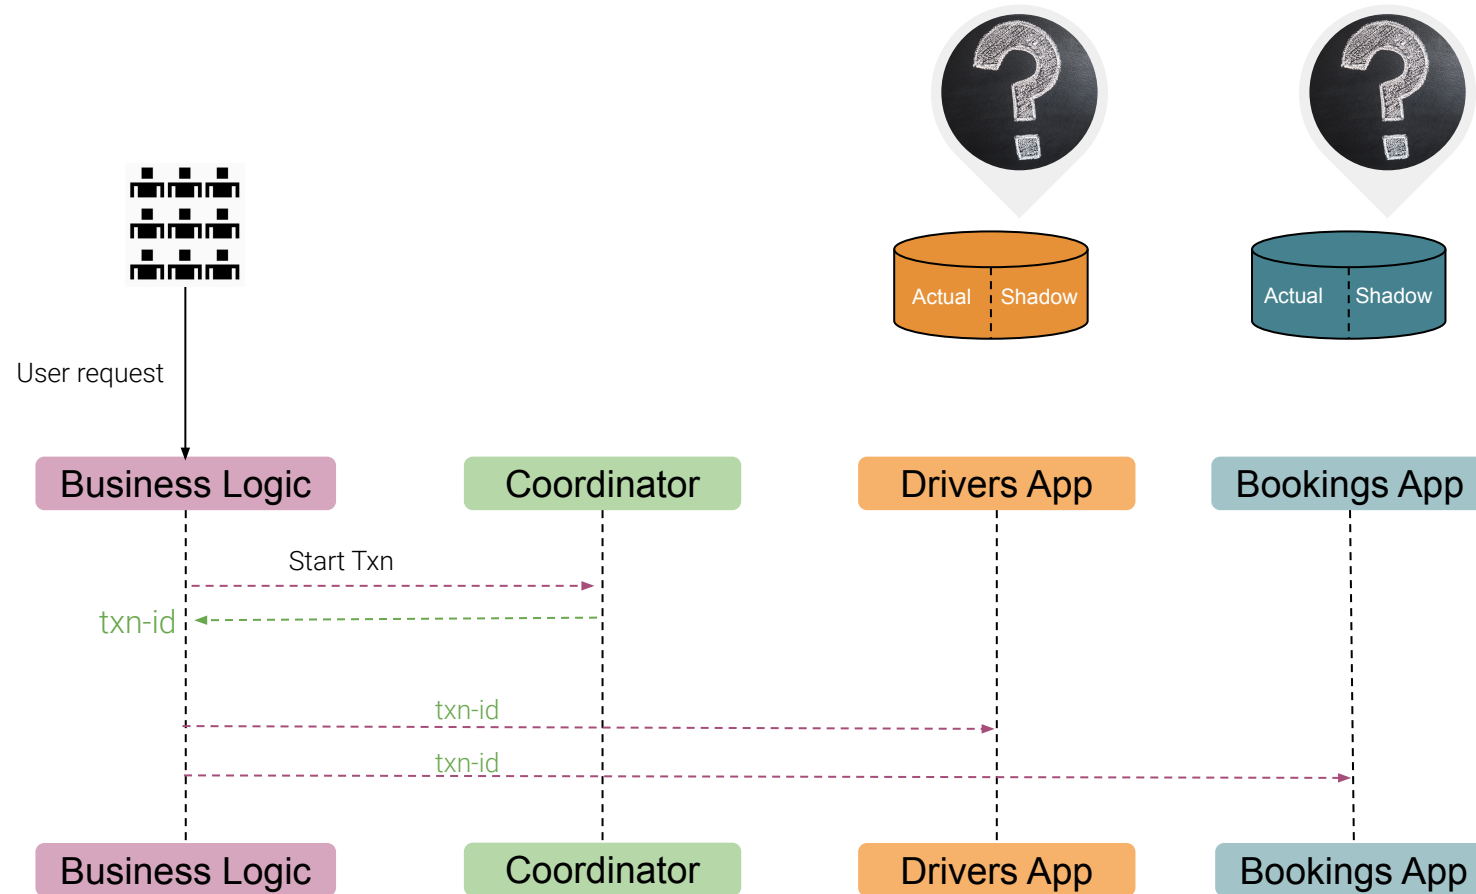
Shadows: A Distributed 2 Phase Commit



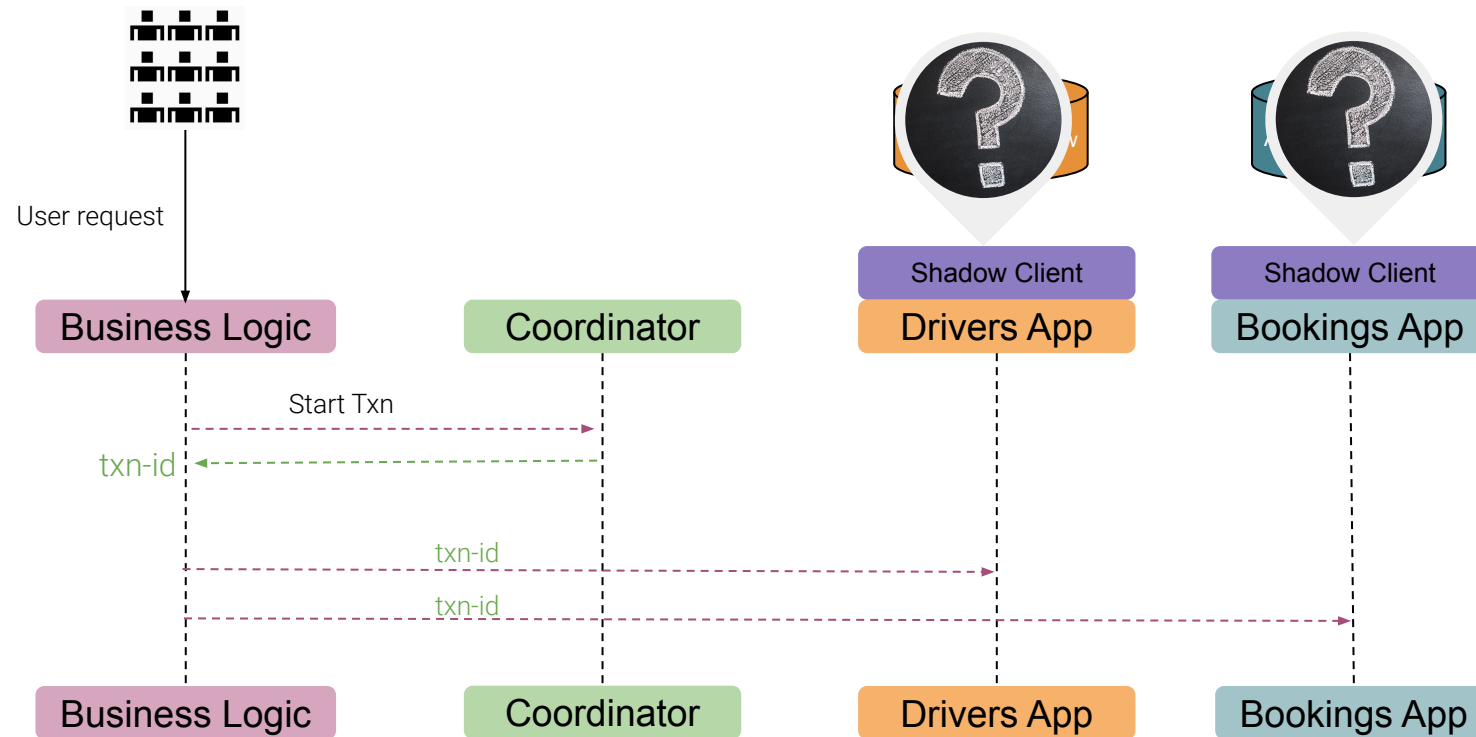
Shadows: A Distributed 2 Phase Commit



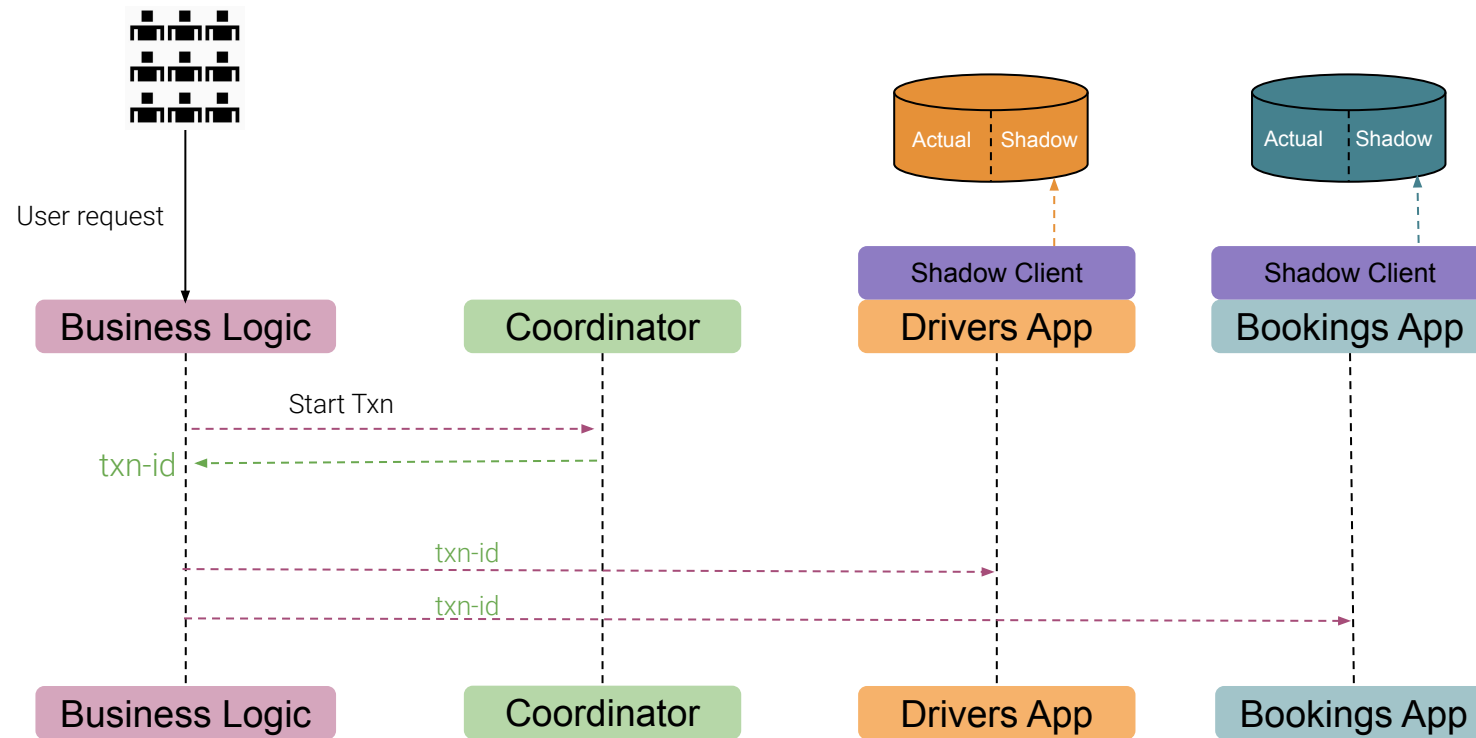
Shadows: Shadow tables?



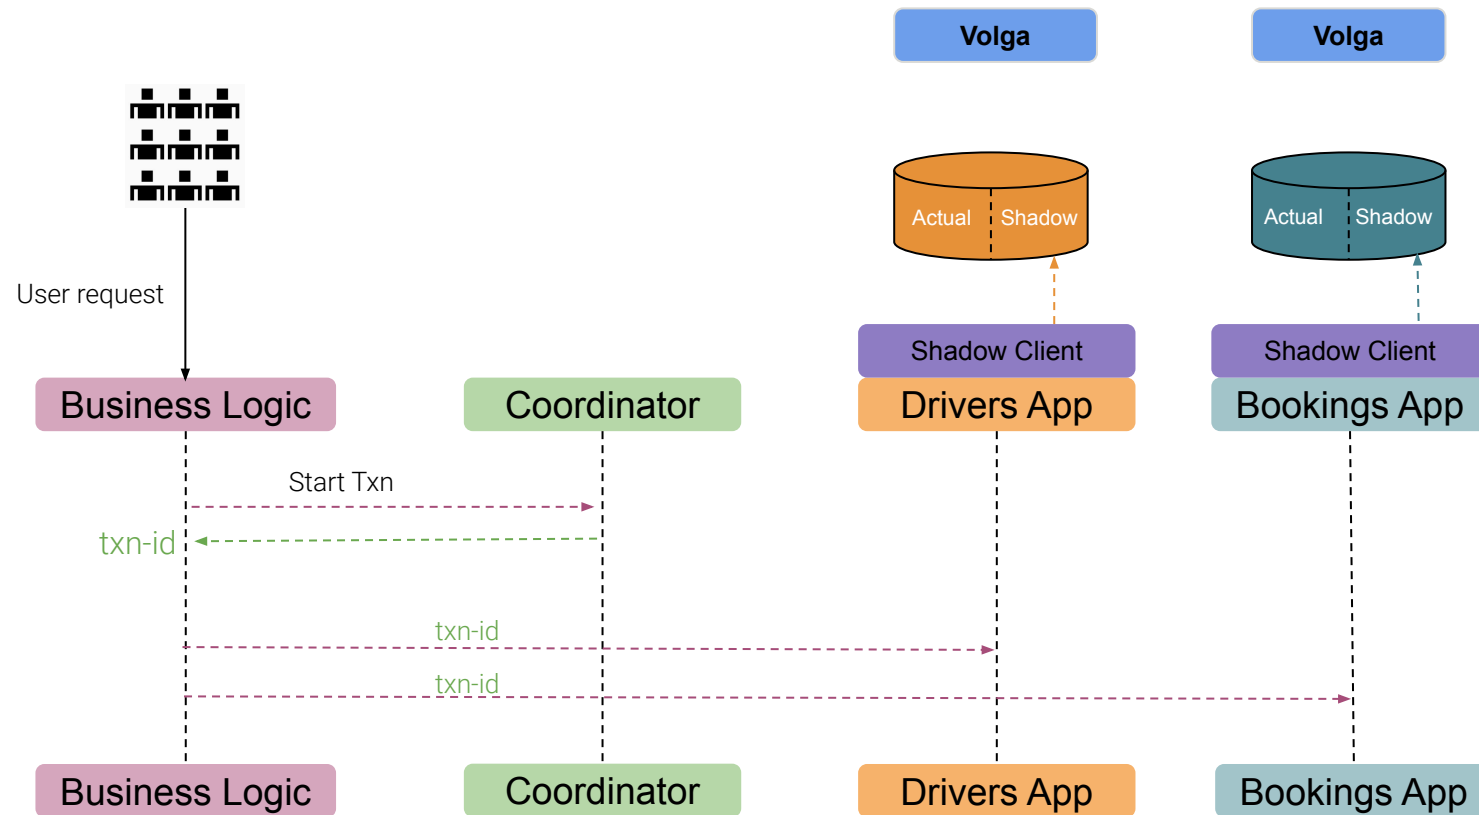
Shadows: Shadow Clients?



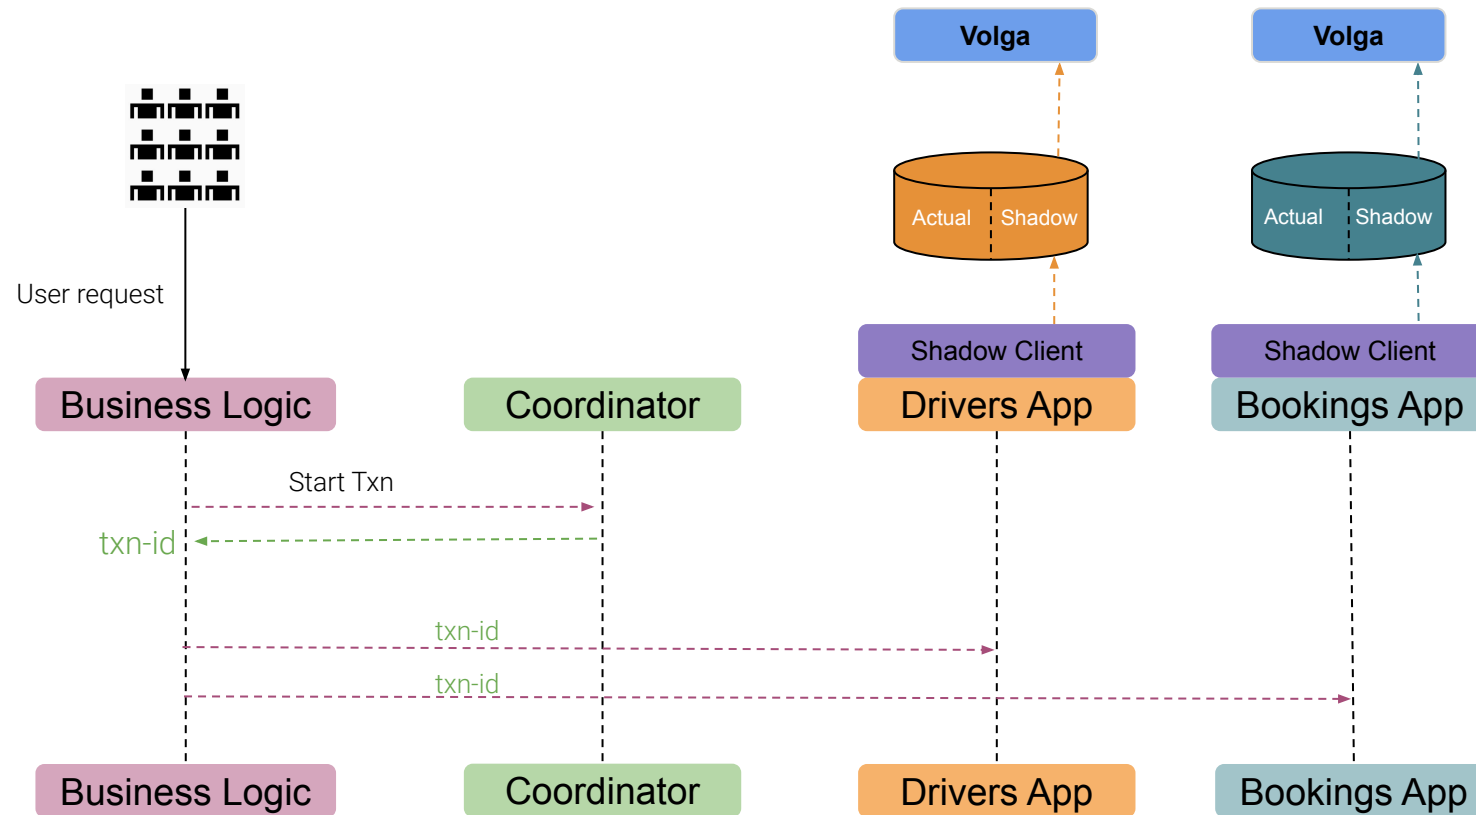
Shadows: A Distributed 2 Phase Commit



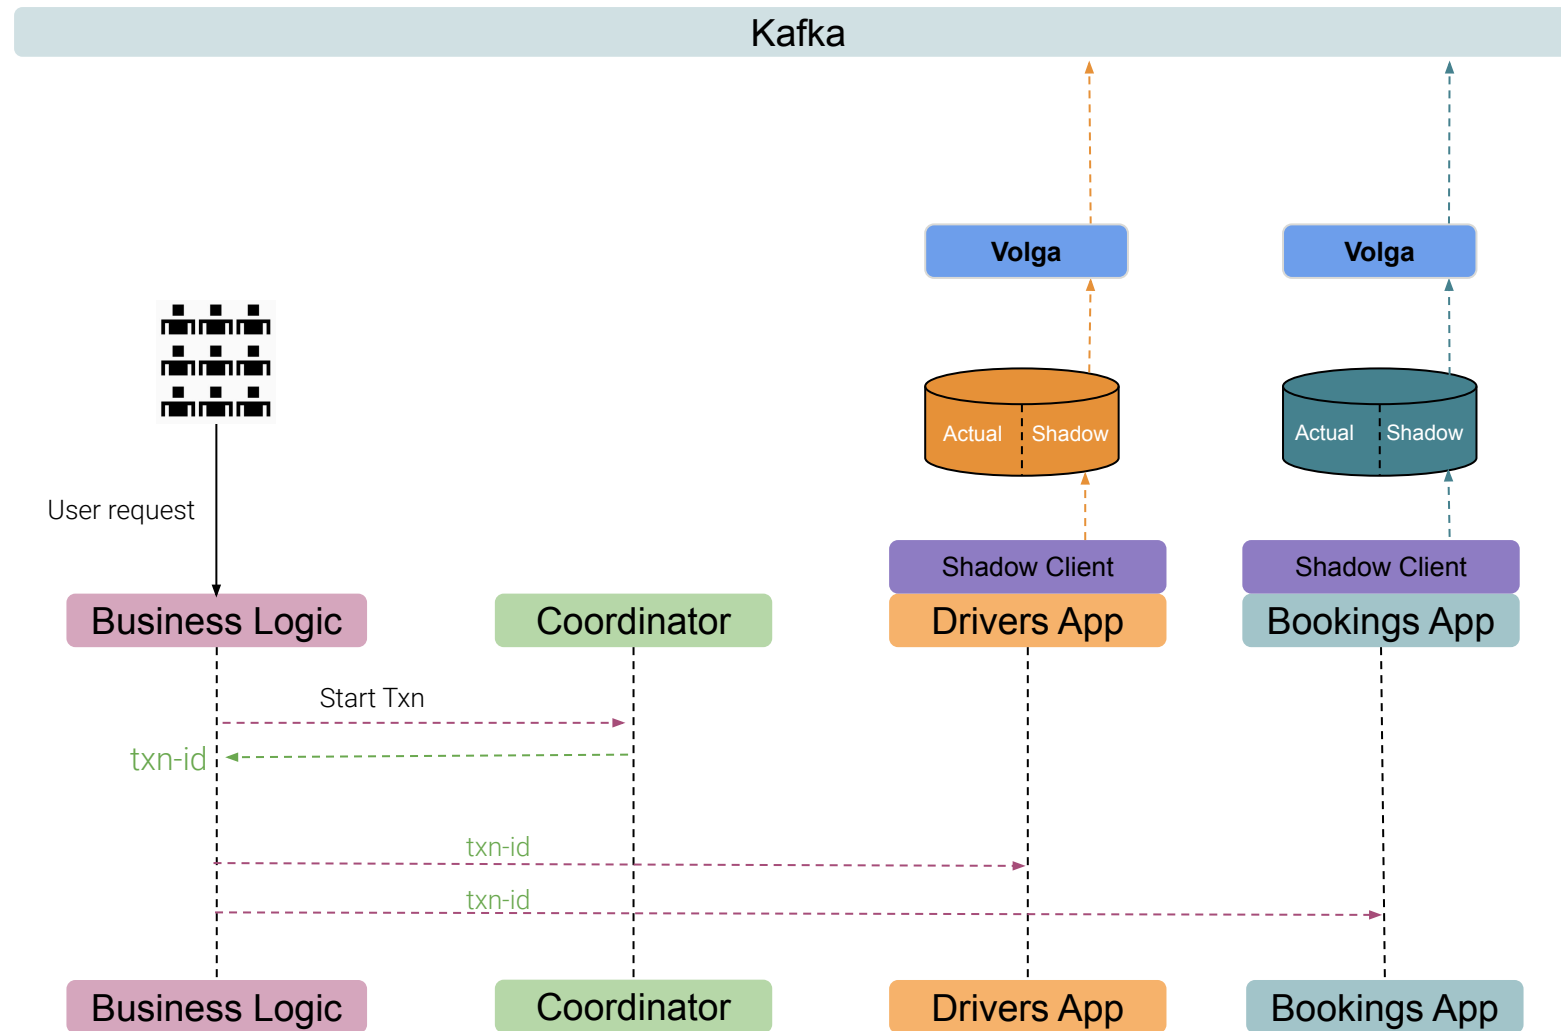
Shadows: A Distributed 2-Phase Commit



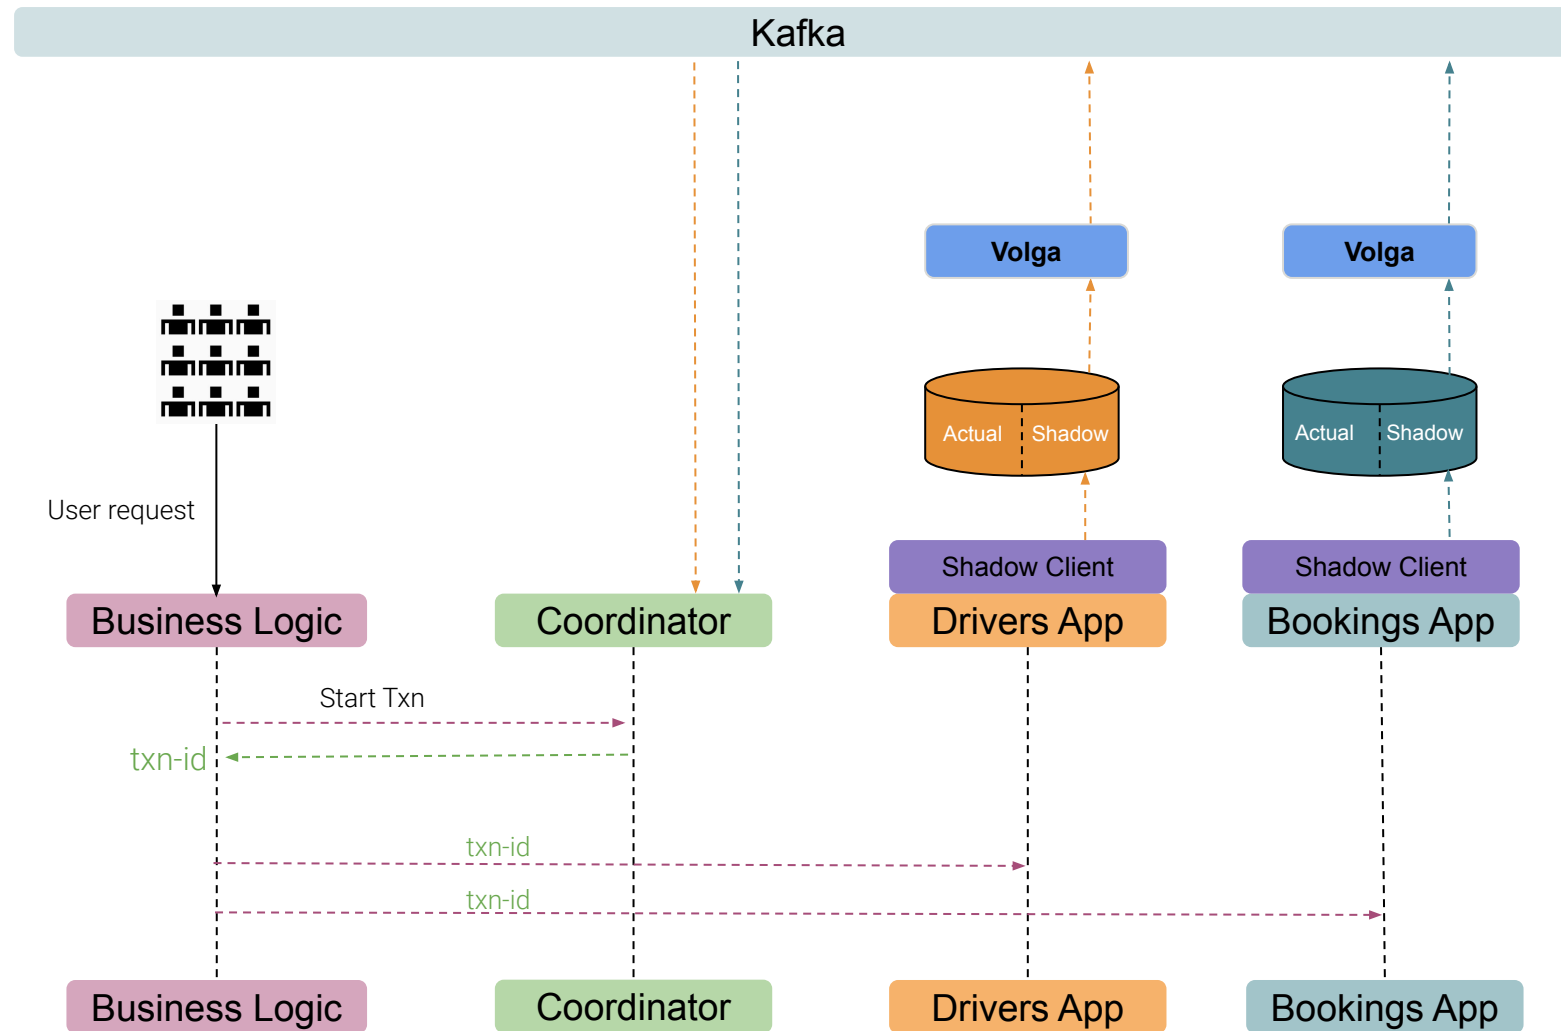
Shadows: A Distributed 2-Phase Commit



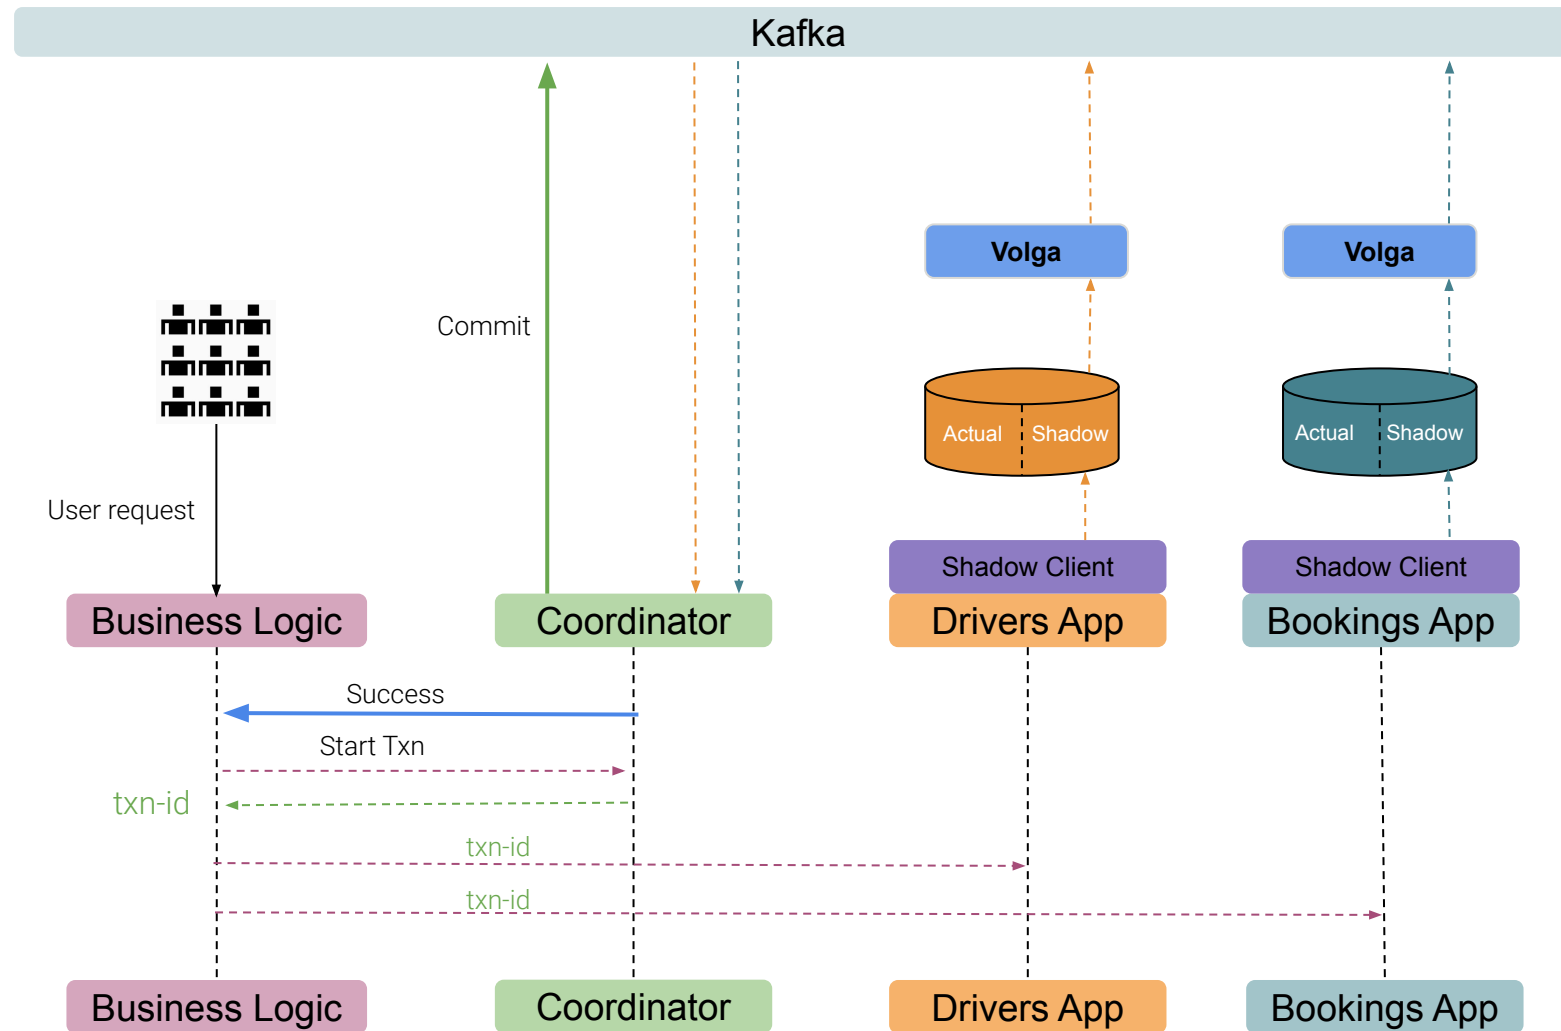
Shadows: A Distributed 2-Phase Commit



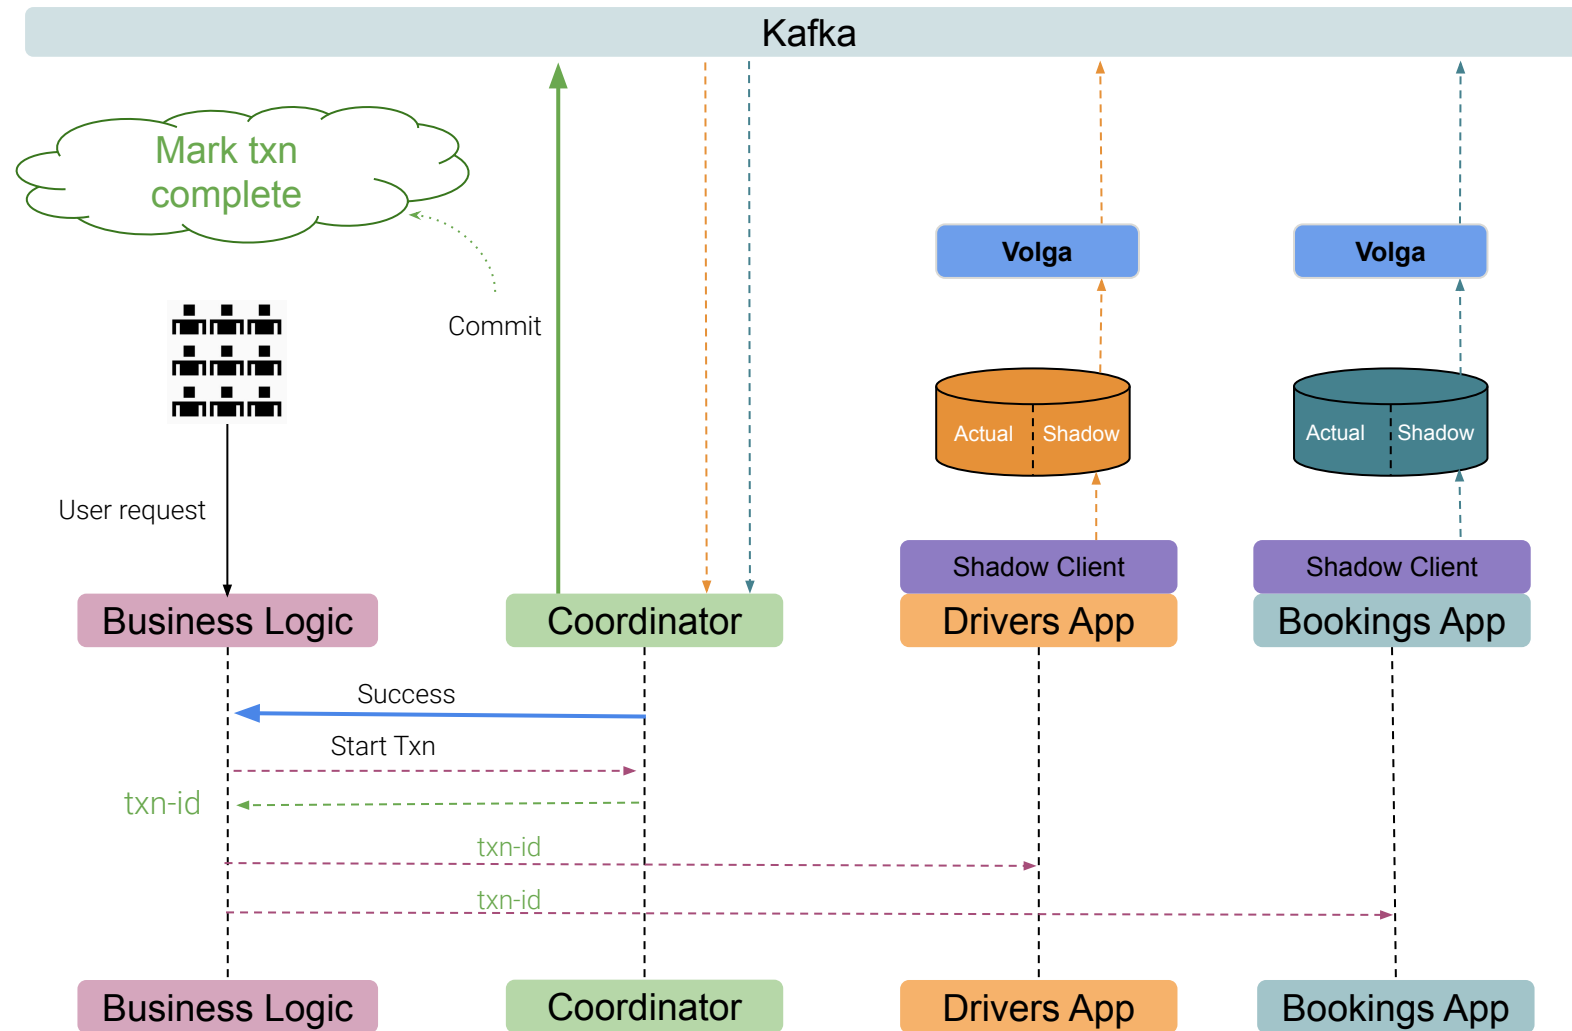
Shadows: A Distributed 2-Phase Commit



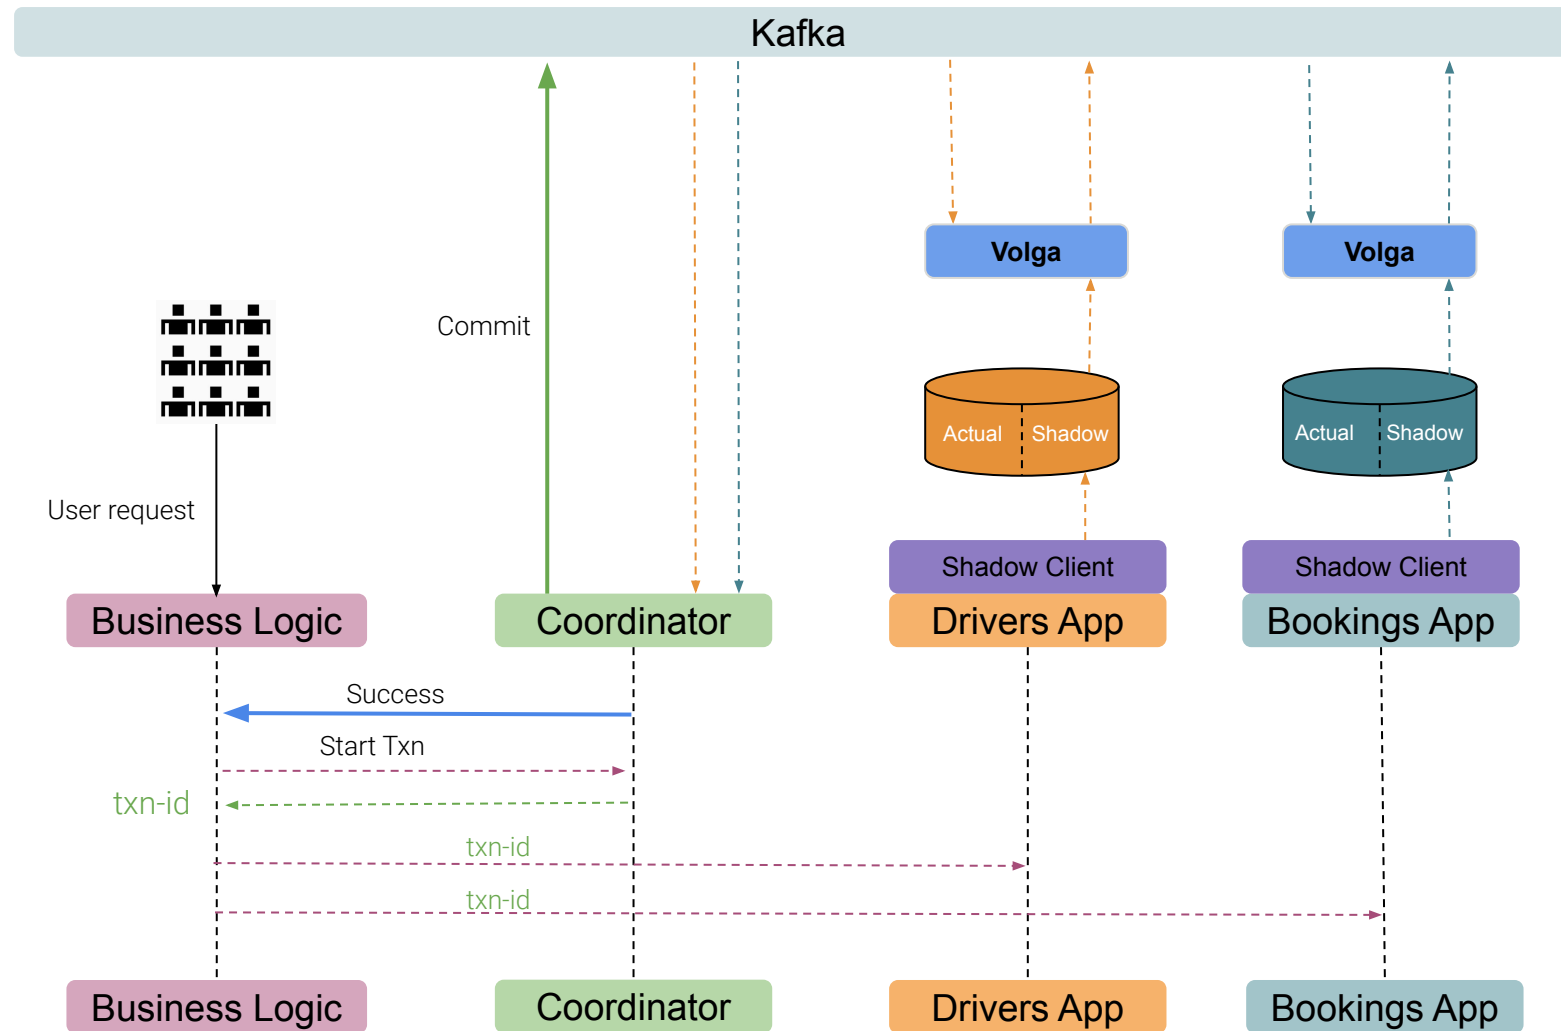
Shadows: A Distributed 2-Phase Commit



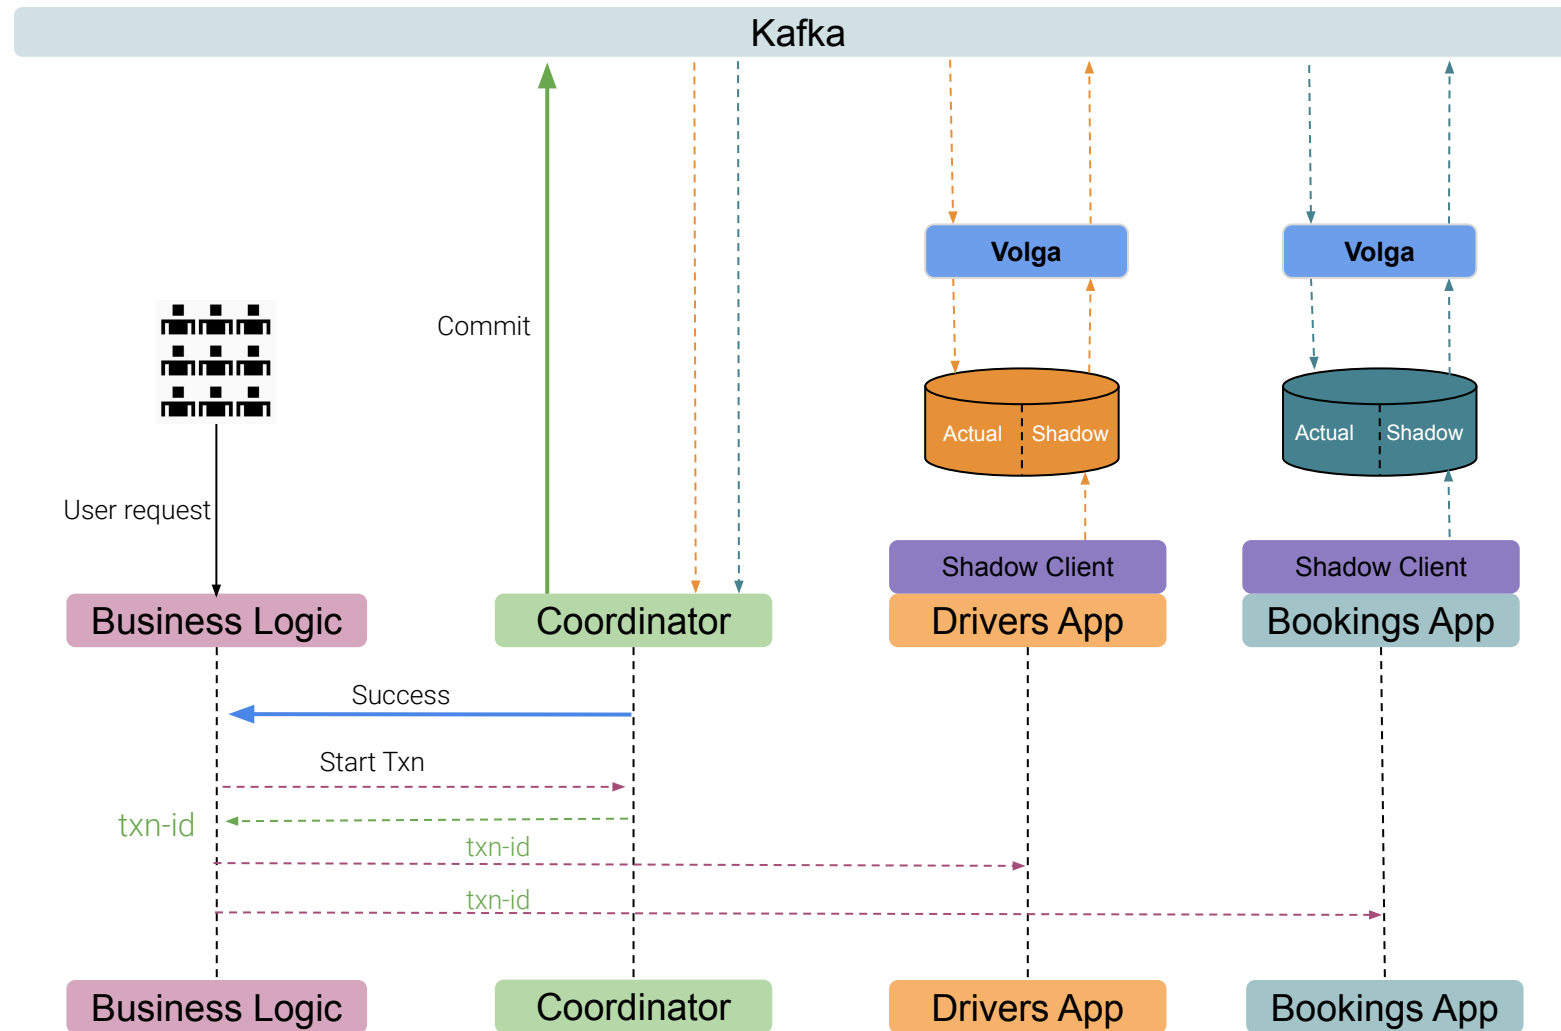
Shadows: A Distributed 2-Phase Commit



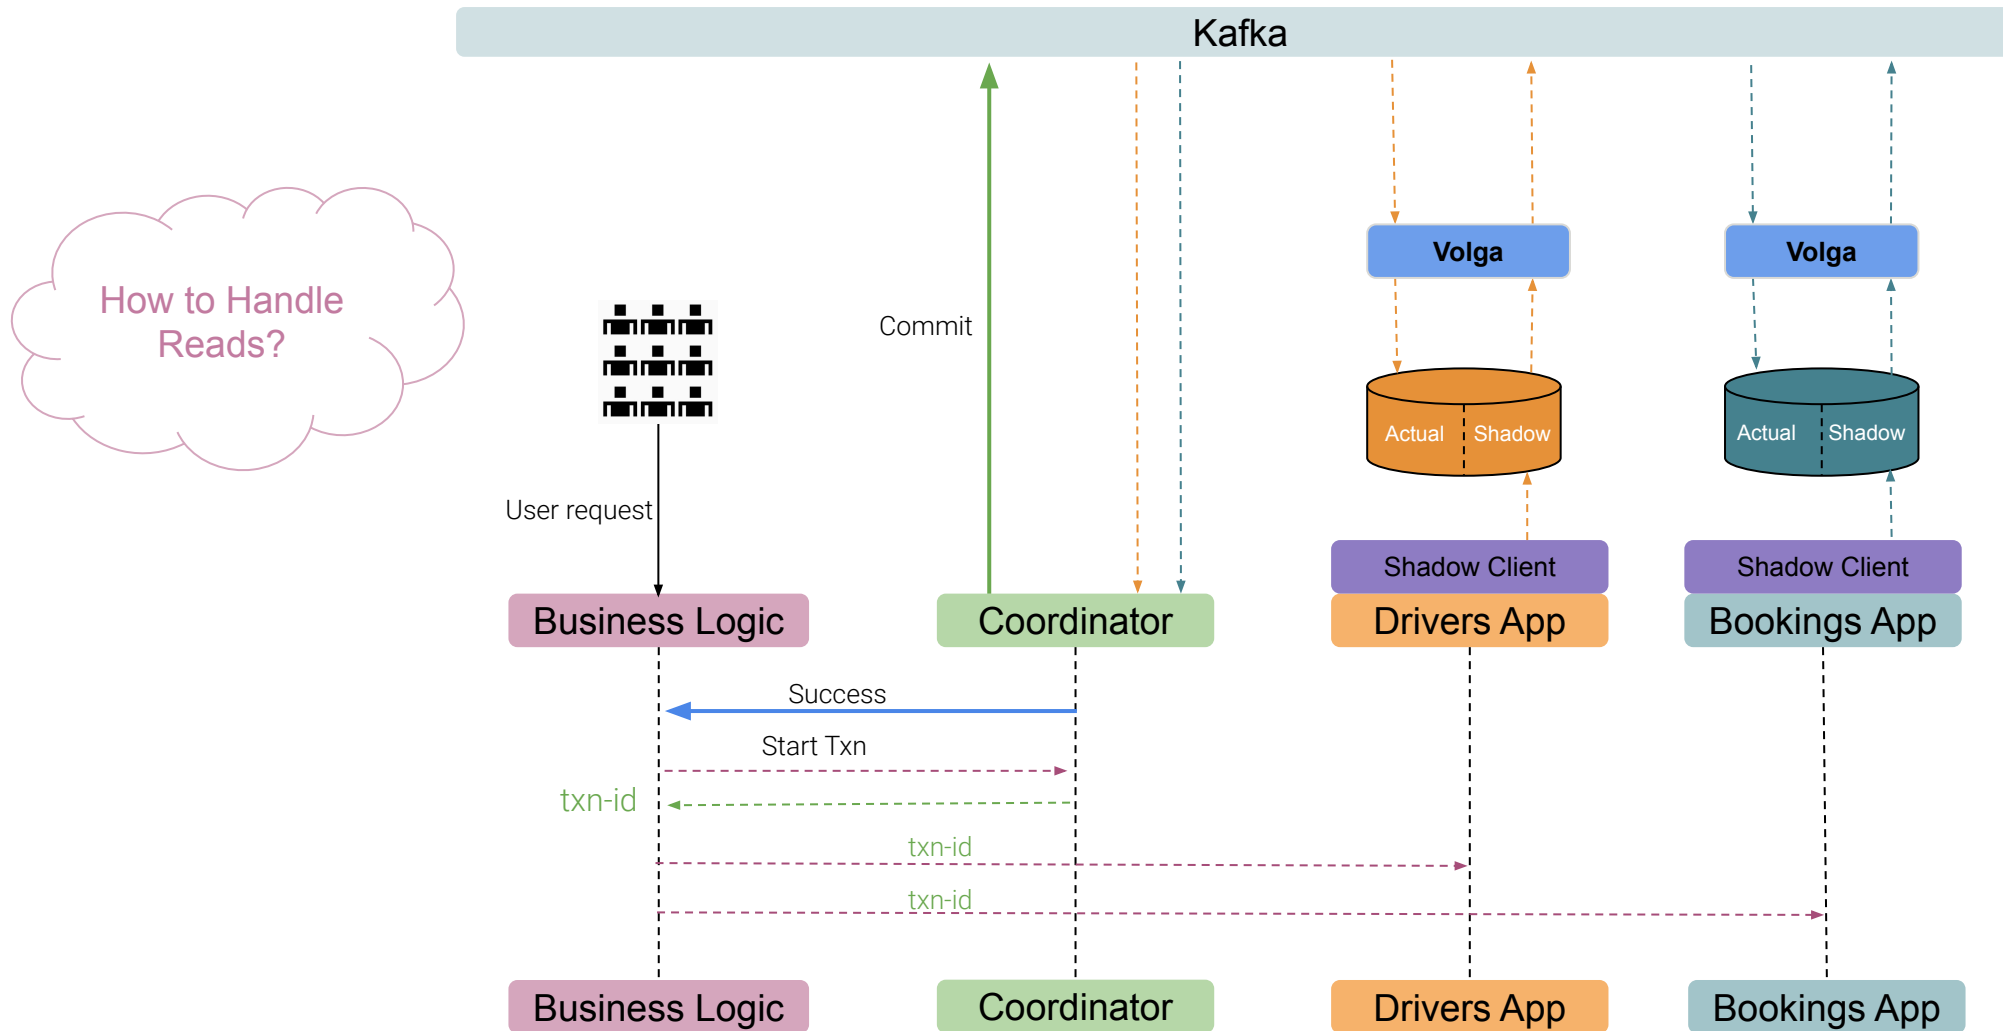
Shadows: A Distributed 2-Phase Commit



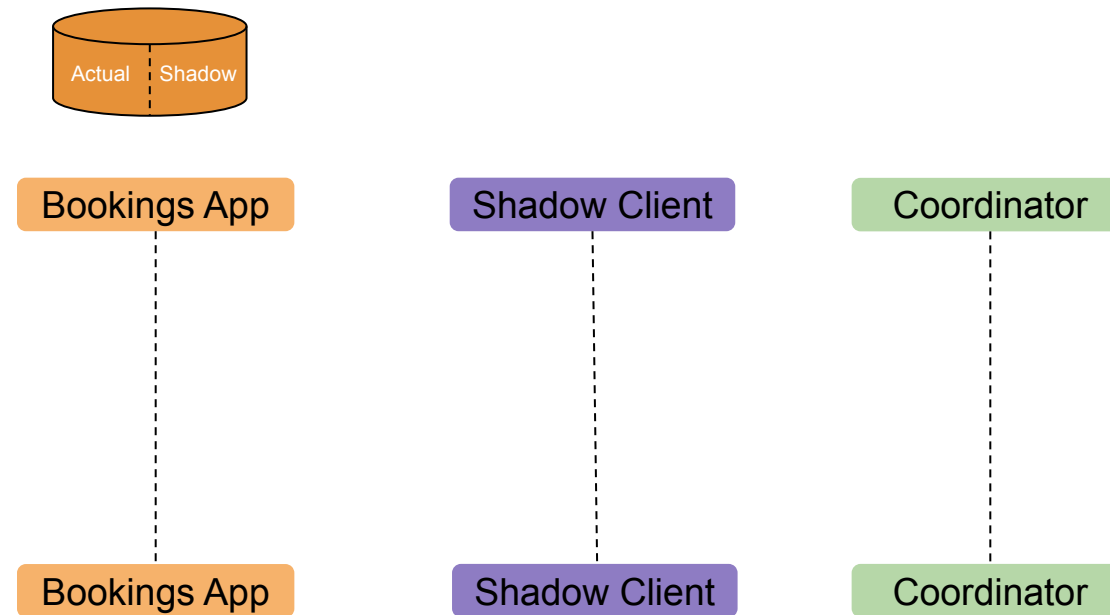
Shadows: A Distributed 2-Phase Commit



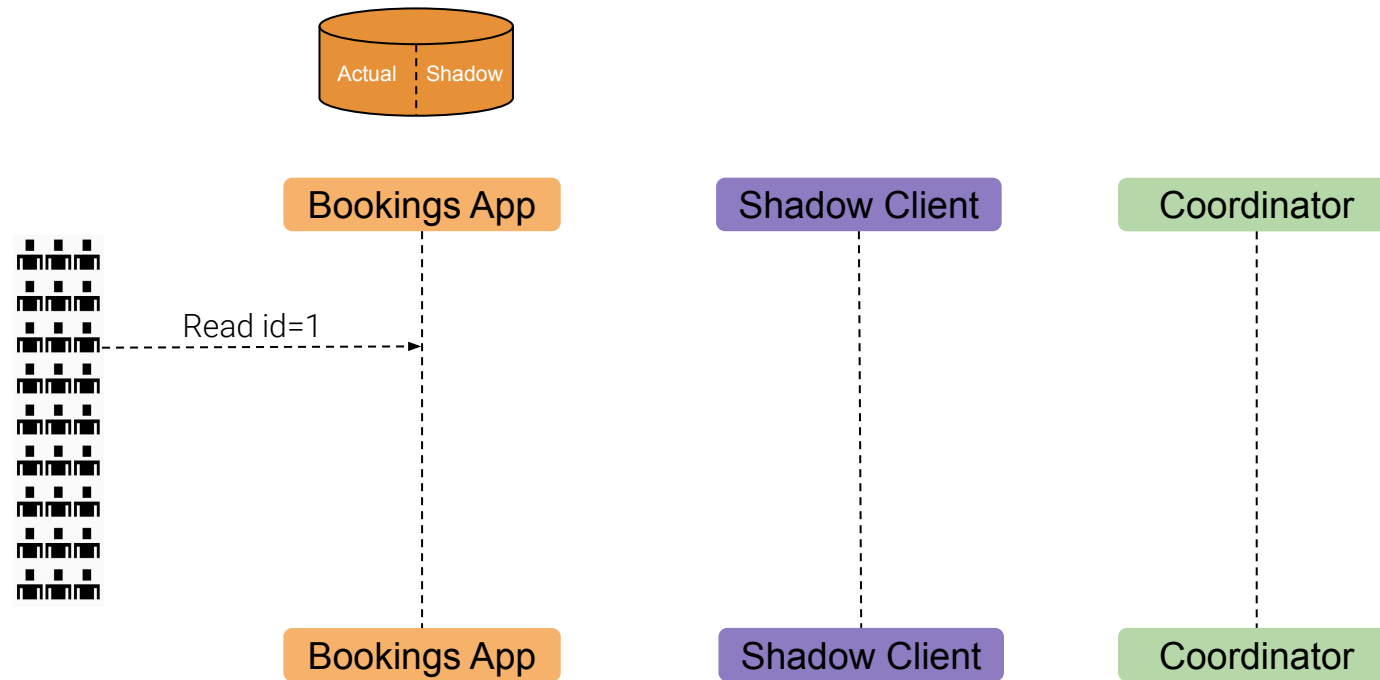
Shadows: A Distributed 2-Phase Commit



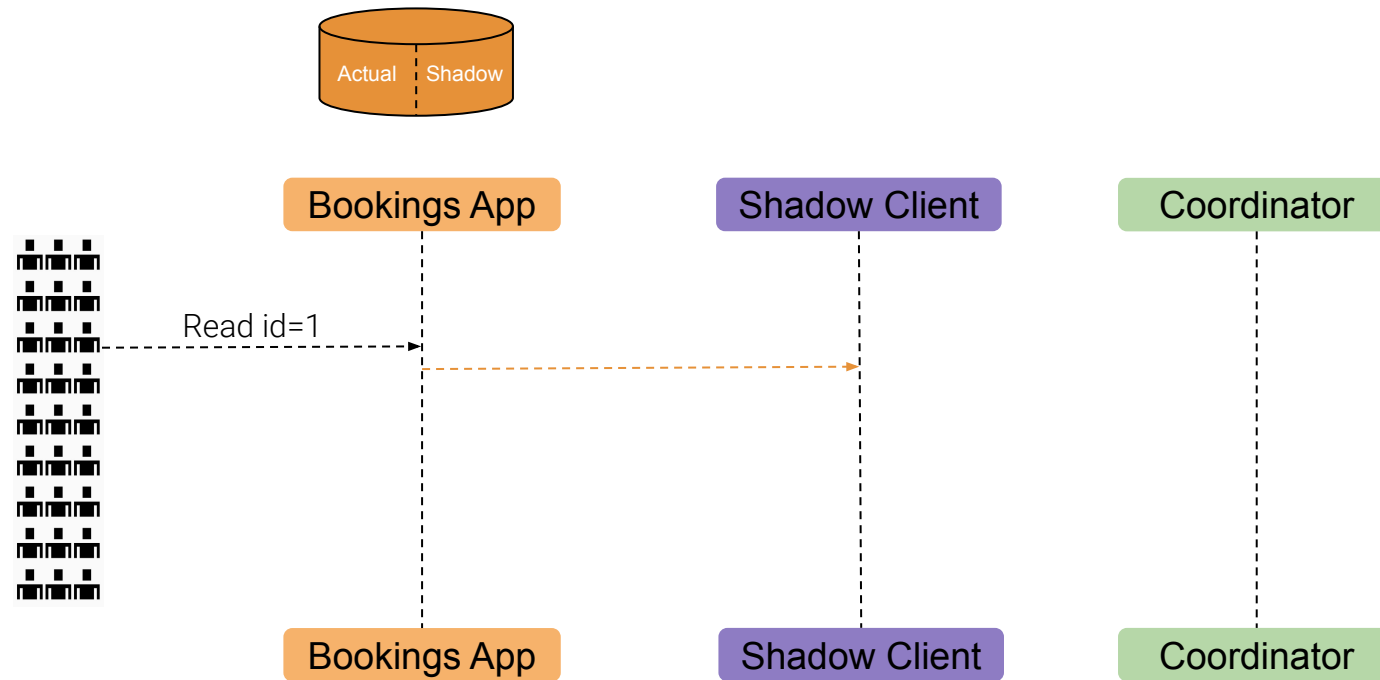
Consistent Reads



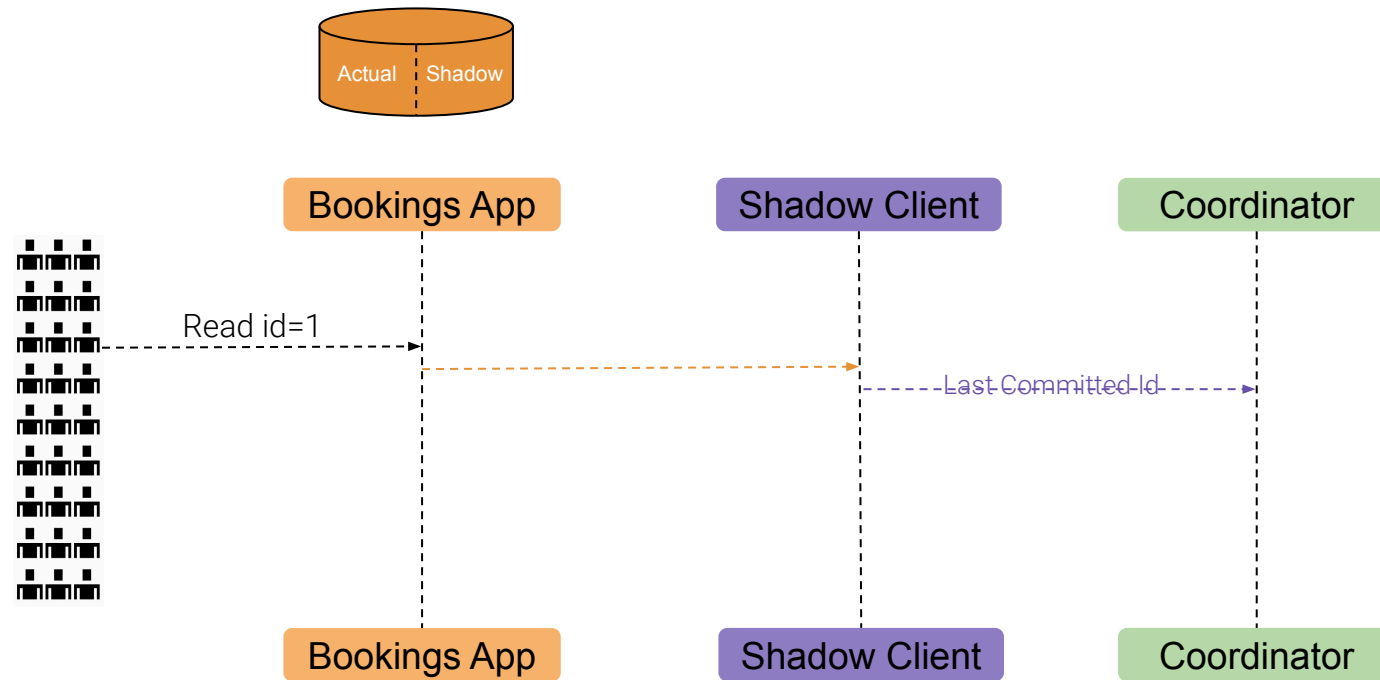
Consistent Reads



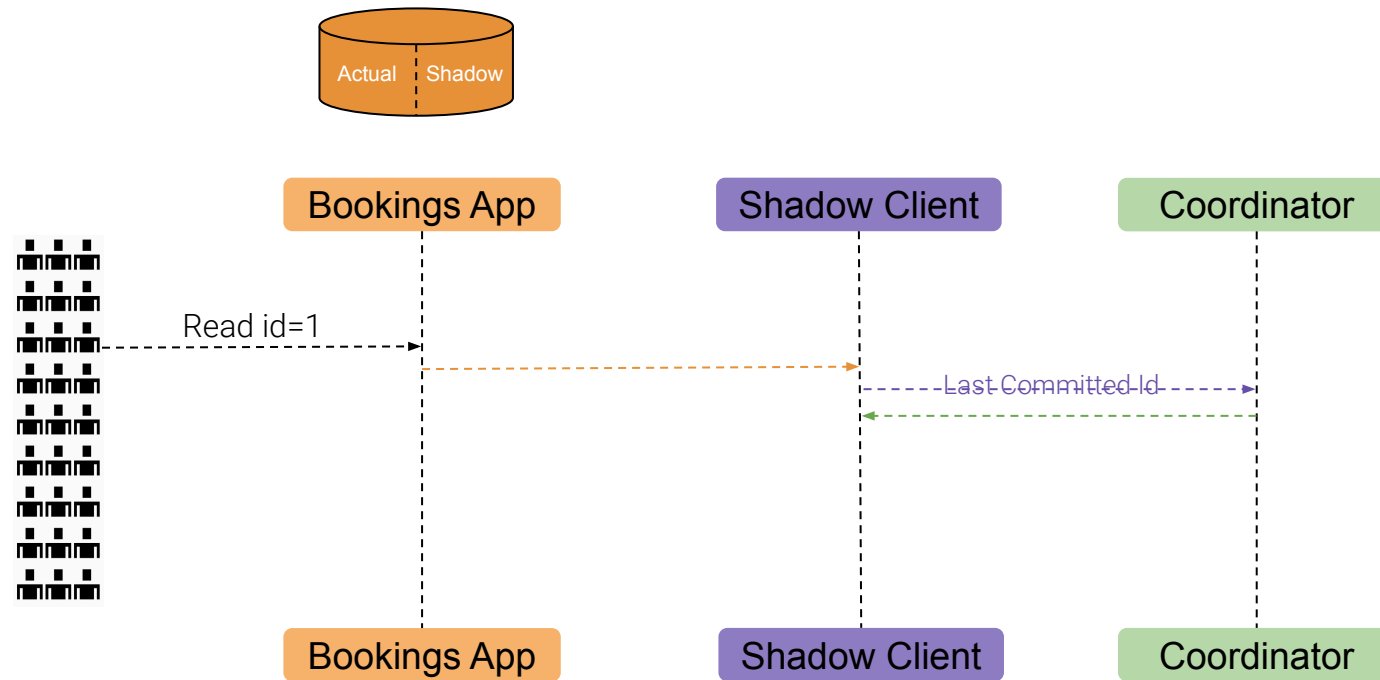
Consistent Reads



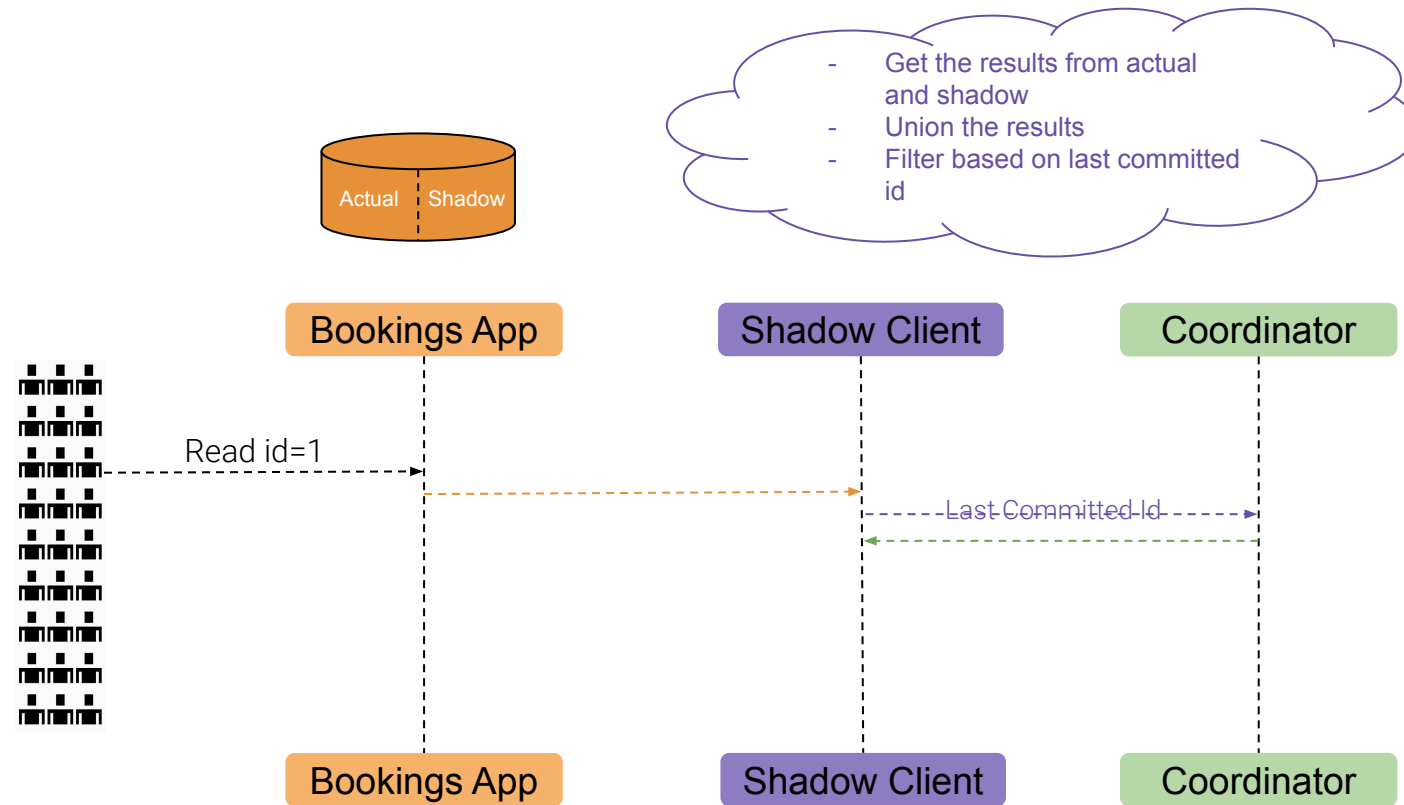
Consistent Reads



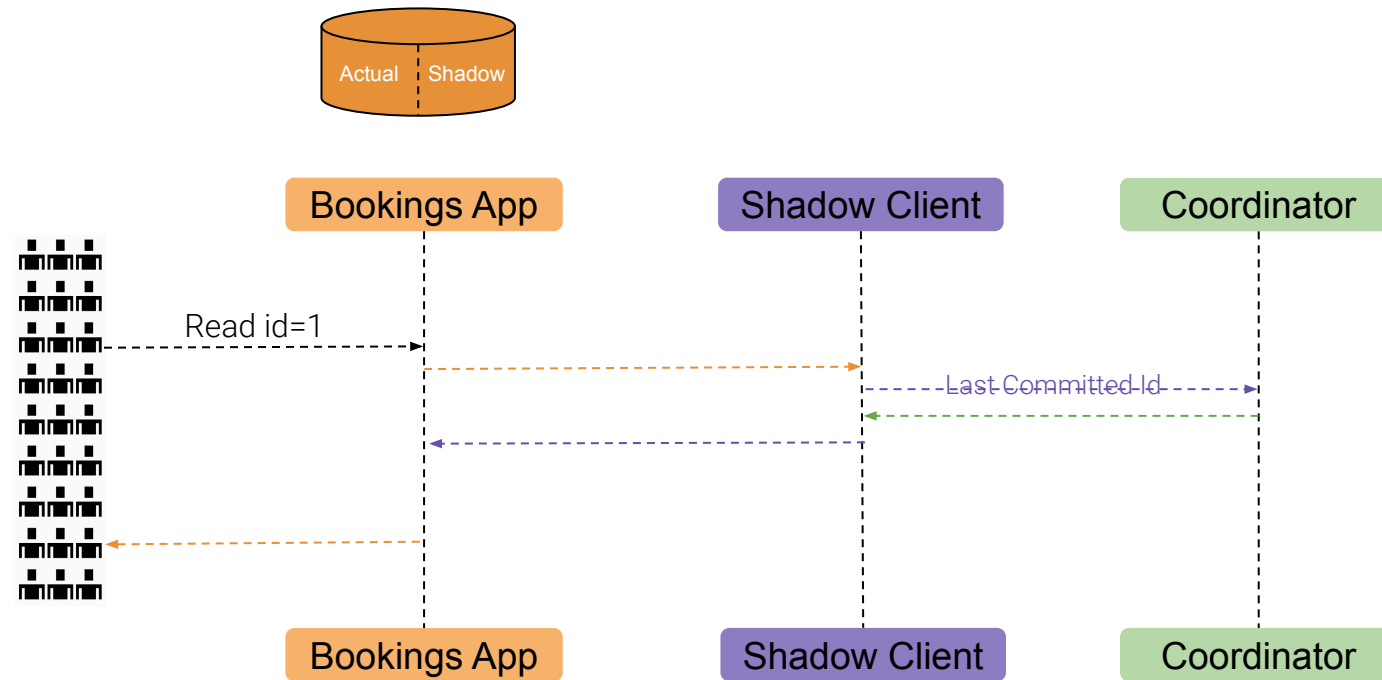
Consistent Reads



Consistent Reads



Consistent Reads





- No locking
- Read committed isolation
- Read committed consistency



- Slower Reads
- Complex to implement

Parameter	2 Phase Commit	Sagas	Shadows
Consistency	Linearizable	Eventual	Read Committed
Isolation	Serializable	Relaxed	Read Committed

So what to use?





Eventual

Shadows

High Availability

SAGAS

2PC

Isolation

Weak

Replication

Consistency

Partition Tolerance

Strong

XA Transaction

Read Committed

Latency

Relaxed

Throughput

Scalability



Eventual

Shadows

High Availability

SAGAS

2PC

Isolation

Weak

Replication

Consistency

Partition Tolerance

Strong

XA Transaction

Read Committed

Latency

Relaxed

Scalability

Throughput



Eventual

Shadows

High Availability

Weak

SAGAS

2PC

Isolation

Replication

Consistency

Partition Tolerance

Strong

XA Transaction

Read Committed

Latency

Relaxed

Scalability

Throughput



Eventual

Shadows

High Availability

SAGAS

2PC

Isolation

Week

Replication

Consistency

Partition Tolerance

Strong

XA Transaction

Read Committed

Relaxed

Latency

Scalability

Throughput



Eventual

Shadows

High Availability

SAGAS

2PC

Isolation

Week

Replication

Consistency

Partition Tolerance

Strong

XA Transaction

Read Committed

Latency

Relaxed

Scalability

Throughput



Eventual

Shadows

High Availability

SAGAS

2PC

Isolation

Week

Replication

Consistency

Partition Tolerance

Strong

XA Transaction

Read Committed

Latency

Relaxed

Scalability

Throughput



THANKYOU



QUESTIONS?



THANKYOU



Database Consistency: if two clients can see different states at the same point in time, we say that their view of the database is inconsistent

Database Isolation: the ability of a database to allow a transaction to execute as if there are no other concurrently running transactions

* Daniel Abadi in DBMS Musings (<http://dbmsmusings.blogspot.com>)



write -->1 ,, read -->2

<2

read -->1 write -->2

0-->

read -->1

result --0

write-->2

read-->3 (write --2)

read->4 write -->5

read == write --2

write--6 ,, read --7

read ...write --2,write.5 (edited)

read--8 ... write2,write5,write6

- According to the coordinator timestamp the transaction which started first or which has the lower timestamp have to finish first. This means that if a concurrent transaction happens and let's say first is read and second is write then it is okay as read will read all the data before this write but if write happens before read then the read will wait for the write to finish and then only return or wait until timeout.

if write write concurrent transaction happens then if first write is still going on then the second one will wait until timeout and if the first one finishes before the timeout then second will start or second will fail as soon as it times out.